

Python for Everybody¹

Exploring Data Using Python 3²

Dr. Charles R. Severance³

Credits 1

Editorial Support: Elliott Hauser, Sue Blumberg 2
Cover Design: Aimee Andrion

Printing History 3

- 2024-Jan-01 Update examples to Python 3.12, remove references to Twitter APIs, rewrite Databases chapter 4
- 2023-Jun-29 Many errata included, switch from Google APIs to OpenStreetMap APIs
- 2016-Jul-05 First Complete Python 3.0 version
- 2015-Dec-20 Initial Python 3.0 rough conversion

Copyright Details 5

Copyright 2009- Dr. Charles R. Severance 6

This work is licensed under a Creative Commons Attribution-NonCommercial-7 ShareAlike 3.0 Unported License. This license is available at

<http://creativecommons.org/licenses/by-nc-sa/3.0/> 8

You can see what the author considers commercial and non-commercial uses of this 9 material as well as license exemptions in the Appendix titled “Copyright Detail”.

Preface¹

Remixing an Open Book²

It is quite natural for academics who are continuously told to “publish or perish”³ to want to always create something from scratch that is their own fresh creation. This book is an experiment in not starting from scratch, but instead “remixing” the book titled *Think Python: How to Think Like a Computer Scientist* written by Allen B. Downey, Jeff Elkner, and others.

In December of 2009, I was preparing to teach *SI502 - Networked Programming* at the University of Michigan for the fifth semester in a row and decided it was time to write a Python textbook that focused on exploring data instead of understanding algorithms and abstractions. My goal in SI502 is to teach people lifelong data handling skills using Python. Few of my students were planning to be professional computer programmers. Instead, they planned to be librarians, managers, lawyers, biologists, economists, etc., who happened to want to skillfully use technology in their chosen field.

I never seemed to find the perfect data-oriented Python book for my course, so I⁵ set out to write just such a book. Luckily at a faculty meeting three weeks before I was about to start my new book from scratch over the holiday break, Dr. Atul Prakash showed me the *Think Python* book which he had used to teach his Python course that semester. It is a well-written Computer Science text with a focus on short, direct explanations and ease of learning.

The overall book structure has been changed to get to doing data analysis problems⁶ as quickly as possible and have a series of running examples and exercises about data analysis from the very beginning.

Chapters 2–10 are similar to the *Think Python* book, but there have been major⁷ changes. Number-oriented examples and exercises have been replaced with data-oriented exercises. Topics are presented in the order needed to build increasingly sophisticated data analysis solutions. Some topics like `try` and `except` are pulled forward and presented as part of the chapter on conditionals. Functions are given very light treatment until they are needed to handle program complexity rather than introduced as an early lesson in abstraction. Nearly all user-defined functions have been removed from the example code and exercises outside of Chapter 4. The word “recursion”¹ does not appear in the book at all.

In chapters 1 and 11–16, all of the material is brand new, focusing on real-world⁸ uses and simple examples of Python for data analysis including regular expressions for searching and parsing, automating tasks on your computer, retrieving data across the network, scraping web pages for data, object-oriented programming, using web services, parsing XML and JSON data, creating and using databases using Structured Query Language, and visualizing data.

The ultimate goal of all of these changes is to shift from a Computer Science to an⁹ Informatics focus and to only include topics into a first technology class that can be useful even if one chooses not to become a professional programmer.

¹Except, of course, for this line.

Students who find this book interesting and want to further explore should look ¹ at Allen B. Downey's *Think Python* book. Because there is a lot of overlap between the two books, students will quickly pick up skills in the additional areas of technical programming and algorithmic thinking that are covered in *Think Python*. And given that the books have a similar writing style, they should be able to move quickly through *Think Python* with a minimum of effort.

As the copyright holder of *Think Python*, Allen has given me permission to change the book's license on the material from his book that remains in this book from the GNU Free Documentation License to the more recent Creative Commons Attribution — Share Alike license. This follows a general shift in open documentation licenses moving from the GFDL to the CC-BY-SA (e.g., Wikipedia). Using the CC-BY-SA license maintains the book's strong copyleft tradition while making it even more straightforward for new authors to reuse this material as they see fit. ²

I feel that this book serves as an example of why open materials are so important ³ to the future of education, and I want to thank Allen B. Downey and Cambridge University Press for their forward-looking decision to make the book available under an open copyright. I hope they are pleased with the results of my efforts and I hope that you, the reader, are pleased with *our* collective efforts.

I would like to thank Allen B. Downey and Lauren Cowles for their help, patience, ⁴ and guidance in dealing with and resolving the copyright issues around this book.

Charles Severance ⁵

www.dr-chuck.com ⁶

Ann Arbor, MI, USA ⁷

September 9, 2013 ⁸

Charles Severance is a Clinical Associate Professor at the University of Michigan ⁹ School of Information.

Contents¹

1 Why should you learn to write programs?	12
1.1 Creativity and motivation	23
1.2 Computer hardware architecture	3
1.3 Understanding programming	4
1.4 Words and sentences	5
1.5 Conversing with Python	6
1.6 Terminology: Interpreter and compiler	8
1.7 Writing a program	10
1.8 What is a program?	10
1.9 The building blocks of programs	11
1.10 What could possibly go wrong?	12
1.11 Debugging	14
1.12 The learning journey	15
1.13 Glossary	15
1.14 Exercises	16
2 Variables, expressions, and statements	194
2.1 Values and types	195
2.2 Variables	20
2.3 Variable names and keywords	21
2.4 Statements	21
2.5 Operators and operands	22
2.6 Expressions	23
2.7 Order of operations	23
2.8 Modulus operator	24
2.9 String operations	24

2.10 Asking the user for input	25	1
2.11 Comments	26	
2.12 Choosing mnemonic variable names	27	
2.13 Debugging	28	
2.14 Glossary	29	
2.15 Exercises	30	

3 Conditional execution 312

3.1 Boolean expressions	31	3
3.2 Logical operators	32	
3.3 Conditional execution	32	
3.4 Alternative execution	34	
3.5 Chained conditionals	34	
3.6 Nested conditionals	35	
3.7 Catching exceptions using try and except	36	
3.8 Short-circuit evaluation of logical expressions	38	
3.9 Debugging	39	
3.10 Glossary	40	
3.11 Exercises	40	

4 Functions 434

4.1 Function calls	43	5
4.2 Built-in functions	43	
4.3 Type conversion functions	44	
4.4 Math functions	45	
4.5 Random numbers	46	
4.6 Adding new functions	47	
4.7 Definitions and uses	48	
4.8 Flow of execution	49	
4.9 Parameters and arguments	49	
4.10 Fruitful functions and void functions	51	
4.11 Why functions?	52	
4.12 Debugging	52	
4.13 Glossary	53	
4.14 Exercises	54	

5 Iteration	57
5.1 Updating variables	57
5.2 The <code>while</code> statement	57
5.3 Infinite loops	58
5.4 Finishing iterations with <code>continue</code>	59
5.5 Definite loops using <code>for</code>	60
5.6 Loop patterns	61
5.6.1 Counting and summing loops	61
5.6.2 Maximum and minimum loops	62
5.7 Debugging	64
5.8 Glossary	64
5.9 Exercises	64
6 Strings	67
6.1 A string is a sequence	67
6.2 Getting the length of a string using <code>len</code>	68
6.3 Traversal through a string with a loop	68
6.4 String slices	69
6.5 Strings are immutable	70
6.6 Looping and counting	70
6.7 The <code>in</code> operator	71
6.8 String comparison	71
6.9 String methods	71
6.10 Parsing strings	74
6.11 Formatted String Literals	74
6.12 Debugging	75
6.13 Glossary	76
6.14 Exercises	76
7 Files	79
7.1 Persistence	79
7.2 Opening files	79
7.3 Text files and lines	81
7.4 Reading files	82
7.5 Searching through a file	83

7.6	Letting the user choose the file name	85	1
7.7	Using <code>try</code> , <code>except</code> , and <code>open</code>	86	
7.8	Writing files	87	
7.9	Debugging	88	
7.10	Glossary	89	
7.11	Exercises	89	

8 Lists 91 2

8.1	A list is a sequence	91	3
8.2	Lists are mutable	92	
8.3	Traversing a list	92	
8.4	List operations	93	
8.5	List slices	94	
8.6	List methods	94	
8.7	Deleting elements	95	
8.8	Lists and functions	96	
8.9	Lists and strings	97	
8.10	Parsing lines	98	
8.11	Objects and values	99	
8.12	Aliasing	100	
8.13	List arguments	100	
8.14	Debugging	102	
8.15	Glossary	105	
8.16	Exercises	105	

9 Dictionaries 109 4

9.1	Dictionary as a set of counters	111	5
9.2	Dictionaries and files	112	
9.3	Looping and dictionaries	114	
9.4	Advanced text parsing	115	
9.5	Debugging	116	
9.6	Glossary	117	
9.7	Exercises	117	

10 Tuples	119¹
10.1 Tuples are immutable	119 ²
10.2 Comparing tuples	120
10.3 Tuple assignment	122
10.4 Dictionaries and tuples	123
10.5 Multiple assignment with dictionaries	124
10.6 The most common words	125
10.7 Using tuples as keys in dictionaries	126
10.8 Sequences: strings, lists, and tuples - Oh My!	126
10.9 List comprehension	127
10.10 Debugging	128
10.11 Glossary	128
10.12 Exercises	128
11 Regular expressions	131³
11.1 Character matching in regular expressions	132 ⁴
11.2 Extracting data using regular expressions	133
11.3 Combining searching and extracting	136
11.4 Escape character	140
11.5 Summary	140
11.6 Bonus section for Unix / Linux users	141
11.7 Debugging	142
11.8 Glossary	142
11.9 Exercises	143
12 Networked programs	145⁵
12.1 Hypertext Transfer Protocol - HTTP	145 ⁶
12.2 The world's simplest web browser	146
12.3 Retrieving an image over HTTP	148
12.4 Retrieving web pages with <code>urllib</code>	150
12.5 Reading binary files using <code>urllib</code>	151
12.6 Parsing HTML and scraping the web	152
12.7 Parsing HTML using regular expressions	152
12.8 Parsing HTML using BeautifulSoup	154
12.9 Bonus section for Unix / Linux users	157
12.10 Glossary	157
12.11 Exercises	158

13 Using Web Services	159	1
13.1 eXtensible Markup Language - XML	159	2
13.2 Parsing XML	160	
13.3 Looping through nodes	161	
13.4 JavaScript Object Notation - JSON	162	
13.5 Parsing JSON	163	
13.6 Application Programming Interfaces	164	
13.7 Security and API usage	165	
13.8 Glossary	166	
14 Object-oriented programming	167	3
14.1 Managing larger programs	167	4
14.2 Getting started	168	
14.3 Using objects	168	
14.4 Starting with programs	169	
14.5 Subdividing a problem	171	
14.6 Our first Python object	171	
14.7 Classes as types	174	
14.8 Object lifecycle	175	
14.9 Multiple instances	176	
14.10 Inheritance	177	
14.11 Summary	178	
14.12 Glossary	179	
15 Using Databases and SQL	181	5
15.1 What is a database?	181	6
15.2 Database concepts	181	
15.3 Database Browser for SQLite	182	
15.4 Creating a database table	182	
15.5 Structured Query Language summary	185	
15.6 Multiple tables and basic data modeling	187	
15.7 Data model diagrams	189	
15.8 Automatically creating primary keys	190	
15.9 Logical keys for fast lookup	191	
15.10 Adding constraints to the database	192	

15.11 Sample multi-table application	193	1
15.12 Many to many relationships in databases	196	
15.13 Modeling data at the many-to-many connection	200	
15.14 Summary	201	
15.15 Debugging	202	
15.16 Glossary	202	
16 Visualizing data	205	2
16.1 Building a OpenStreetMap from geocoded data	205	3
16.2 Visualizing networks and interconnections	207	
16.3 Visualizing mail data	210	
A Contributions	217	4
A.1 Translations	217	5
A.2 Contributor List for Python for Everybody	217	
A.3 Contributor List for Python for Informatics	218	
A.4 Preface for “Think Python”	218	
A.4.1 The strange history of “Think Python”	218	6
A.4.2 Acknowledgements for “Think Python”	219	
A.5 Contributor List for “Think Python”	220	
B Copyright Detail	221	7

Chapter 1¹

Why should you learn to² write programs?

Writing programs (or programming) is a very creative and rewarding activity. You³ can write programs for many reasons, ranging from making your living to solving a difficult data analysis problem to having fun to helping someone else solve a problem. This book assumes that *everyone* needs to know how to program, and that once you know how to program you will figure out what you want to do with your newfound skills.

We are surrounded in our daily lives with computers ranging from laptops to cell phones. We can think of these computers as our “personal assistants” who can take care of many things on our behalf. The hardware in our current-day computers is essentially built to continuously ask us the question, “What would you like me to do next?”

Programmers add an operating system and a set of applications to the hardware⁵ and we end up with a Personal Digital Assistant that is quite helpful and capable of helping us do many different things.

Our computers are fast and have vast amounts of memory and could be very helpful⁶ to us if we only knew the language to speak to explain to the computer what we would like it to “do next”. If we knew this language, we could tell the computer to do tasks on our behalf that were repetitive. Interestingly, the kinds of things computers can do best are often the kinds of things that we humans find boring and mind-numbing.



Figure 1.1: Personal Digital Assistant⁸



Figure 1.2: Programmers Talking to You 2

For example, look at the first three paragraphs of this chapter and tell me the 3 most commonly used word and how many times the word is used. While you were able to read and understand the words in a few seconds, counting them is almost painful because it is not the kind of problem that human minds are designed to solve. For a computer, the opposite is true, reading and understanding text from a piece of paper is hard for a computer to do but counting the words and telling you how many times the most used word was used is very easy for the computer:

```
python words.py 4
Enter file: words.txt
to 16
```

Our “personal information analysis assistant” quickly told us that the word “to” 5 was used sixteen times in the first three paragraphs of this chapter.


This very fact that computers are good at things that humans are not is why you 6 need to become skilled at talking “computer language”. Once you learn this new language, you can delegate mundane tasks to your partner (the computer), leaving more time for you to do the things that you are uniquely suited for. You bring creativity, intuition, and inventiveness to this partnership.

1.1 Creativity and motivation 7

While this book is not intended for professional programmers, professional programming can be a very rewarding job both financially and personally. Building useful, elegant, and clever programs for others to use is a very creative activity. Your computer or Personal Digital Assistant (PDA) usually contains many different programs from many different groups of programmers, each competing for your attention and interest. They try their best to meet your needs and give you a great user experience in the process. In some situations, when you choose a piece of software, the programmers are directly compensated because of your choice.

If we think of programs as the creative output of groups of programmers, perhaps 9 the following figure is a more sensible version of our PDA:

For now, our primary motivation is not to make money or please end users, but instead for us to be more productive in handling the data and information that we will encounter in our lives. When you first start, you will be both the programmer and the end user of your programs. As you gain skill as a programmer and programming feels more creative to you, your thoughts may turn toward developing programs for others.



1

Figure 1.3: Hardware Architecture 2


1.2 Computer hardware architecture 3

Before we start learning the language we speak to give instructions to computers to develop software, we need to learn a small amount about how computers are built. If you were to take apart your computer or cell phone and look deep inside, you would find the following parts:

The high-level definitions of these parts are as follows: 5

- The *Central Processing Unit* (or CPU) is the part of the computer that is 6 built to be obsessed with “what is next?” If your computer is rated at 3.0 Gigahertz, it means that the CPU will ask “What next?” three billion times per second. You are going to have to learn how to talk fast to keep up with the CPU.
- The *Main Memory* is used to store information that the CPU needs in a hurry. The main memory is nearly as fast as the CPU. But the information stored in the main memory vanishes when the computer is turned off.
- The *Secondary Memory* is also used to store information, but it is much slower than the main memory. The advantage of the secondary memory is that it can store information even when there is no power to the computer. Examples of secondary memory are disk drives or flash memory (typically found in USB sticks and portable music players).
- The *Input and Output Devices* are simply our screen, keyboard, mouse, microphone, speaker, touchpad, etc. They are all of the ways we interact with the computer.
- These days, most computers also have a *Network Connection* to retrieve information over a network. We can think of the network as a very slow place to store and retrieve data that might not always be “up”. So in a sense, the network is a slower and at times unreliable form of *Secondary Memory*.

While most of the detail of how these components work is best left to computer 7 builders, it helps to have some terminology so we can talk about these different parts as we write our programs.



1

Figure 1.4: Where Are You? 2

As a programmer, your job is to use and orchestrate each of these resources to 3 solve the problem that you need to solve and analyze the data you get from the solution. As a programmer you will mostly be “talking” to the CPU and telling it what to do next. Sometimes you will tell the CPU to use the main memory, secondary memory, network, or the input/output devices.

You need to be the person who answers the CPU’s “What next?” question. But it 4 would be very uncomfortable to shrink you down to 5mm tall and insert you into the computer just so you could issue a command three billion times per second. So instead, you must write down your instructions in advance. We call these stored instructions a *program* and the act of writing these instructions down and getting the instructions to be correct *programming*.

1.3 Understanding programming⁵

In the rest of this book, we will try to turn you into a person who is skilled in the art 6 of programming. In the end you will be a *programmer* - perhaps not a professional programmer, but at least you will have the skills to look at a data/information analysis problem and develop a program to solve the problem.

In a sense, you need two skills to be a programmer:⁷

- First, you need to know the programming language (Python) - you need to 8 know the vocabulary and the grammar. You need to be able to spell the words in this new language properly and know how to construct well-formed “sentences” in this new language.
- Second, you need to “tell a story”. In writing a story, you combine words and sentences to convey an idea to the reader. There is a skill and art in constructing the story, and skill in story writing is improved by doing some writing and getting some feedback. In programming, our program is the “story” and the problem you are trying to solve is the “idea”.

Once you learn one programming language such as Python, you will find it much 9 easier to learn a second programming language such as JavaScript or C++. The

new programming language has very different vocabulary and grammar but the problem-solving skills will be the same across all programming languages.

You will learn the “vocabulary” and “sentences” of Python pretty quickly. It will take longer for you to be able to write a coherent program to solve a brand-new problem. We teach programming much like we teach writing. We start reading and explaining programs, then we write simple programs, and then we write increasingly complex programs over time. At some point you “get your muse” and see the patterns on your own and can see more naturally how to take a problem and write a program that solves that problem. And once you get to that point, programming becomes a very pleasant and creative process.

We start with the vocabulary and structure of Python programs. Be patient as the simple examples remind you of when you started reading for the first time.

1.4 Words and sentences⁴

Unlike human languages, the Python vocabulary is actually pretty small. We call this “vocabulary” the “reserved words” or “keywords”. These are words that have very special meaning to Python. When Python sees these words in a Python program, they have one and only one meaning to Python. Later as you write programs you will make up your own words that have meaning to you called *variables*. You will have great latitude in choosing your names for your variables, but you cannot use any of Python’s reserved words as a name for a variable.

When we train a dog, we use special words like “sit”, “stay”, and “fetch”. When you talk to a dog and don’t use any of the reserved words, they just look at you with a quizzical look on their face until you say a reserved word. For example, if you say, “I wish more people would walk to improve their overall health”, what most dogs likely hear is, “blah blah blah *walk* blah blah blah blah.” That is because “walk” is a reserved word in dog language. Many might suggest that the language between humans and cats has no reserved words¹.

The reserved words in the language where humans talk to Python include the following:

False	await	else	import	pass	8
None	break	except	in	raise	
True	class	finally	is	return	
and	continue	for	lambda	try	
as	def	from	nonlocal	while	
assert	del	global	not	with	
async	elif	if	or	yield	

That is it, and unlike a dog, Python is already completely trained. When you say “try”, Python will try every time you say it without fail.

We will learn these reserved words and how they are used in good time, but for now we will focus on the Python equivalent of “speak” (in human-to-dog language). The nice thing about telling Python to speak is that we can even tell it what to say by giving it a message in quotes:

¹<http://xkcd.com/231/>

```
print('Hello world!') 1
```

And we have even written our first syntactically correct Python sentence. Our 2 sentence starts with the function *print* followed by a string of text of our choosing enclosed in single quotes. The strings in the print statements are enclosed in quotes. Single quotes and double quotes do the same thing; most people use single quotes except in cases where a single quote (which is also an apostrophe) appears in the string.

1.5 Conversing with Python³

Now that we have a word and a simple sentence that we know in Python, we need 4 to know how to start a conversation with Python to test our new language skills.

Before you can converse with Python, you must first install the Python software 5 on your computer and learn how to start Python on your computer. That is too much detail for this chapter so I suggest that you consult www.py4e.com where I have detailed instructions and screencasts of setting up and starting Python on Macintosh and Windows systems. At some point, you will be in a terminal or command window and you will type *python* and the Python interpreter will start executing in interactive mode and appear somewhat as follows:

```
Python 3.11.6 (main, Nov 2 2023, 04:39:43) 6
[Clang 14.0.3 (clang-1403.0.22.14.1)] on darwin
Type "help", "copyright", "credits" or "license" for more
information.
>>>
```

The **>>>** prompt is the Python interpreter's way of asking you, “What do you want 7 me to do next?” Python is ready to have a conversation with you. All you have to know is how to speak the Python language.

Let's say for example that you did not know even the simplest Python language 8 words or sentences. You might want to use the standard line that astronauts use when they land on a faraway planet and try to speak with the inhabitants of the planet:

```
>>> I come in peace, please take me to your leader 9
File "<stdin>", line 1
      I come in peace take me tou your leader
          ^
>>>
```

SyntaxError: invalid syntax 10

This is not going so well. Unless you think of something quickly, the inhabitants 11 of the planet are likely to stab you with their spears, put you on a spit, roast you over a fire, and eat you for dinner.

Luckily you brought a copy of this book on your travels, and you thumb to this 12 very page and try again:

```
>>> print('Hello world!')  
Hello world!
```

This is looking much better, so you try to communicate some more:²

```
>>> print('You must be the legendary god that comes from the sky')  
You must be the legendary god that comes from the sky  
>>> print('We have been waiting for you for a long time')  
We have been waiting for you for a long time  
>>> print('Our legend says you will be very tasty with mustard')  
Our legend says you will be very tasty with mustard  
>>> print 'We will have a feast tonight unless you say  
    File "<stdin>", line 1  
        print 'We will have a feast tonight unless you say  
                ^
```

```
SyntaxError: unterminated string literal (detected at line 1)  
>>>
```

The conversation was going so well for a while and then you made the tiniest⁵ mistake using the Python language and Python brought the spears back out.

At this point, you should also realize that while Python is amazingly complex and⁶ powerful and very picky about the syntax you use to communicate with it, Python is *not* intelligent. You are really just having a conversation with yourself, but using proper syntax.

In a sense, when you use a program written by someone else the conversation is⁷ between you and those other programmers with Python acting as an intermediary. Python is a way for the creators of programs to express how the conversation is supposed to proceed. And in just a few more chapters, you will be one of those programmers using Python to talk to the users of your program.

Before we leave our first conversation with the Python interpreter, you should⁸ probably know the proper way to say “good-bye” when interacting with the inhabitants of Planet Python:

```
>>> good-bye  
Traceback (most recent call last):  
  File "<stdin>", line 1, in <module>  
NameError: name 'good' is not defined  
>>> if you don't mind, I need to leave  
    File "<stdin>", line 1  
        if you don't mind, I need to leave  
                ^
```

```
SyntaxError: unterminated string literal (detected at line 1)  
>>> quit()
```

You will notice that the error is different for the first two incorrect attempts. The second error is different because *if* is a reserved word and Python saw the reserved word and thought we were trying to say something but got the syntax of the sentence wrong.¹¹

The proper way to say “good-bye” to Python is to enter `quit()` at the interactive chevron `>>>` prompt. It would have probably taken you quite a while to guess that one, so having a book handy probably will turn out to be helpful.¹

1.6 Terminology: Interpreter and compiler²

Python is a *high-level* language intended to be relatively straightforward for humans to read and write and for computers to read and process. Other high-level languages include Java, C++, PHP, Ruby, Basic, Perl, JavaScript, and many more. The actual hardware inside the Central Processing Unit (CPU) does not understand any of these high-level languages.³

The CPU understands a language we call *machine language*. Machine language is very simple and frankly very tiresome to write because it is represented all in zeros and ones:⁴

```
001010001101001001010100000011115  
1110011000001110100101011011016  
...7
```

Machine language seems quite simple on the surface, given that there are only zeros and ones, but its syntax is even more complex and far more intricate than Python. So very few programmers ever write machine language. Instead we build various translators to allow programmers to write in high-level languages like Python or JavaScript and these translators convert the programs to machine language for actual execution by the CPU.⁸

Since machine language is tied to the computer hardware, machine language is not *portable* across different types of hardware. Programs written in high-level languages can be moved between different computers by using a different interpreter on the new machine or recompiling the code to create a machine language version of the program for the new machine.⁹

These programming language translators fall into two general categories: (1) interpreters and (2) compilers.¹⁰

An *interpreter* reads the source code of the program as written by the programmer, parses the source code, and interprets the instructions on the fly. Python is an interpreter and when we are running Python interactively, we can type a line of Python (a sentence) and Python processes it immediately and is ready for us to type another line of Python.¹¹

Some of the lines of Python tell Python that you want it to remember some value for later. We need to pick a name for that value to be remembered and we can use that symbolic name to retrieve the value later. We use the term *variable* to refer to the labels we use to refer to this stored data.¹²

```
>>> x = 613  
>>> print(x)14  
615  
>>> y = x * 716
```

```
>>> print(y) 1  
42  
>>>
```

In this example, we ask Python to remember the value six and use the label `x` so we can retrieve the value later. We verify that Python has actually remembered the value using `print`. Then we ask Python to retrieve `x` and multiply it by seven and put the newly computed value in `y`. Then we ask Python to print out the value currently in `y`.

Even though we are typing these commands into Python one line at a time, Python is treating them as an ordered sequence of statements with later statements able to retrieve data created in earlier statements. We are writing our first simple paragraph with four sentences in a logical and meaningful order.

It is the nature of an *interpreter* to be able to have an interactive conversation as shown above. A *compiler* needs to be handed the entire program in a file, and then it runs a process to translate the high-level source code into machine language and then the compiler puts the resulting machine language into a file for later execution.

If you have a Windows system, often these executable machine language programs have a suffix of “.exe” or “.dll” which stand for “executable” and “dynamic link library” respectively. In Linux and Macintosh, there is no suffix that uniquely marks a file as executable.

If you were to open an executable file in a text editor, it would look completely crazy and be unreadable:

It is not easy to read or write machine language, so it is nice that we have *interpreters* and *compilers* that allow us to write in high-level languages like Python or C.

Now at this point in our discussion of compilers and interpreters, you should be 9
wondering a bit about the Python interpreter itself. What language is it written
in? Is it written in a compiled language? When we type “python”, what exactly
is happening?

The Python interpreter is written in a high-level language called “C”. You can look at the actual source code for the Python interpreter by going to www.python.org and working your way to their source code. So Python is a program itself and it is compiled into machine code. When you installed Python on your computer (or the vendor installed it), you copied a machine-code copy of the translated Python program onto your system. In Windows, the executable machine code for Python itself is likely in a file with a name like:

That is more than you really need to know to be a Python programmer, but 1 sometimes it pays to answer those little nagging questions right at the beginning.

1.7 Writing a program²

Typing commands into the Python interpreter is a great way to experiment with 3 Python's features, but it is not recommended for solving more complex problems.

When we want to write a program, we use a text editor to write the Python 4 instructions into a file, which is called a *script*. By convention, Python scripts have names that end with .py.

To execute the script, you have to tell the Python interpreter the name of the file. 5 In a command window, you would type `python hello.py` as follows:

```
$ cat hello.py      6
print('Hello world!')
$ python hello.py
Hello world!
```

The “\$” is the operating system prompt, and the “cat hello.py” is showing us that 7 the file “hello.py” has a one-line Python program to print a string.

We call the Python interpreter and tell it to read its source code from the file 8 “hello.py” instead of prompting us for lines of Python code interactively.

You will notice that there was no need to have `quit()` at the end of the Python 9 program in the file. When Python is reading your source code from a file, it knows to stop when it reaches the end of the file.

1.8 What is a program?¹⁰

The definition of a *program* at its most basic is a sequence of Python statements 11 that have been crafted to do something. Even our simple *hello.py* script is a program. It is a one-line program and is not particularly useful, but in the strictest definition, it is a Python program.

It might be easiest to understand what a program is by thinking about a problem 12 that a program might be built to solve, and then looking at a program that would solve that problem.

Lets say you are doing Social Computing research on Facebook posts and you are 13 interested in the most frequently used word in a series of posts. You could print out the stream of Facebook posts and pore over the text looking for the most common word, but that would take a long time and be very mistake prone. You would be smart to write a Python program to handle the task quickly and accurately so you can spend the weekend doing something fun.

For example, look at the following text about a clown and a car. Look at the text 14 and figure out the most common word and how many times it occurs.

```
the clown ran after the car and the car ran into the tent 1
and the tent fell down on the clown and the car
```

Then imagine that you are doing this task looking at millions of lines of text.² Frankly it would be quicker for you to learn Python and write a Python program to count the words than it would be to manually scan the words.

The even better news is that I already came up with a simple program to find the most common word in a text file. I wrote it, tested it, and now I am giving it to you to use so you can save some time.³

```
name = input('Enter file: ')
handle = open(name, 'r')
counts = dict()

for line in handle:
    words = line.split()
    for word in words:
        counts[word] = counts.get(word, 0) + 1

bigcount = None
bigword = None
for word, count in list(counts.items()):
    if bigcount is None or count > bigcount:
        bigword = word
        bigcount = count

print(bigword, bigcount)
```

Code: <https://www.py4e.com/code3/words.py> ⁵

You don't even need to know Python to use this program. You will need to get through Chapter 10 of this book to fully understand the awesome Python techniques that were used to make the program. You are the end user, you simply use the program and marvel at its cleverness and how it saved you so much manual effort. You simply type the code into a file called *words.py* and run it or you download the source code from <http://www.py4e.com/code3/> and run it.⁶

This is a good example of how Python and the Python language are acting as an intermediary between you (the end user) and me (the programmer). Python is a way for us to exchange useful instruction sequences (i.e., programs) in a common language that can be used by anyone who installs Python on their computer. So neither of us are talking *to Python*, instead we are communicating with each other *through Python*.⁷

1.9 The building blocks of programs⁸

In the next few chapters, we will learn more about the vocabulary, sentence structure, paragraph structure, and story structure of Python. We will learn about the⁹

powerful capabilities of Python and how to compose those capabilities together to **1** create useful programs.

There are some low-level conceptual patterns that we use to construct programs. **2** These constructs are not just for Python programs, they are part of every programming language from machine language up to the high-level languages.

input Get data from the “outside world”. This might be reading data from a **3** file, or even some kind of sensor like a microphone or GPS. In our initial programs, our input will come from the user typing data on the keyboard.

output Display the results of the program on a screen or store them in a file or **4** perhaps write them to a device like a speaker to play music or speak text.

sequential execution Perform statements one after another in the order they **5** are encountered in the script.

conditional execution Check for certain conditions and then execute or skip a **6** sequence of statements.

repeated execution Perform some set of statements repeatedly, usually with **7** some variation.

reuse Write a set of instructions once and give them a name and then reuse those **8** instructions as needed throughout your program.

It sounds almost too simple to be true, and of course it is never so simple. It is like **9** saying that walking is simply “putting one foot in front of the other”. The “art” of writing a program is composing and weaving these basic elements together many times over to produce something that is useful to its users.

The word counting program above directly uses all of these patterns except for **10** one.

1.10 What could possibly go wrong?¹¹

As we saw in our earliest conversations with Python, we must communicate very **12** precisely when we write Python code. The smallest deviation or mistake will cause Python to give up looking at your program.

Beginning programmers often take the fact that Python leaves no room for errors **13** as evidence that Python is mean, hateful, and cruel. While Python seems to like everyone else, Python knows them personally and holds a grudge against them. Because of this grudge, Python takes our perfectly written programs and rejects them as “unfit” just to torment us.

```
>>> print 'Hello world!' 14
File "<stdin>", line 1
    print 'Hello world!
    ^^^^^^^^^^'
```

SyntaxError: invalid syntax **15**

```
>>> print ('Hello world') 16
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
```

```
NameError: name 'primit' is not defined. Did you mean: 'print'? 1
```

```
>>> I hate you Python! 2
File "<stdin>", line 1
    I hate you Python!
    ^^^^
```

```
SyntaxError: invalid syntax 3
```

```
>>> if you come out of there, I would teach you a lesson 4
File "<stdin>", line 1
    if you come out of there, I would teach you a lesson
    ^^^^
```

```
SyntaxError: invalid syntax 5
```

```
>>>
```

There is little to be gained by arguing with Python. It is just a tool. It has no 6 emotions and it is happy and ready to serve you whenever you need it. Its error messages sound harsh, but they are just Python's call for help. It has looked at what you typed, and it simply cannot understand what you have entered.

Python is much more like a dog, loving you unconditionally, having a few key words 7 that it understands, looking at you with a sweet look on its face (>>>), and waiting for you to say something it understands. When Python says "SyntaxError: invalid syntax", it is simply wagging its tail and saying, "You seemed to say something but I just don't understand what you meant, but please keep talking to me (>>>)."

As your programs become increasingly sophisticated, you will encounter three general types of errors: 8

Syntax errors These are the first errors you will make and the easiest to fix. A 9 syntax error means that you have violated the "grammar" rules of Python. Python does its best to point right at the line and character where it noticed it was confused. The only tricky bit of syntax errors is that sometimes the mistake that needs fixing is actually earlier in the program than where Python noticed it was confused. So the line and character that Python indicates in a syntax error may just be a starting point for your investigation.

Logic errors A logic error is when your program has good syntax but there is 10 a mistake in the order of the statements or perhaps a mistake in how the statements relate to one another. A good example of a logic error might be, "take a drink from your water bottle, put it in your backpack, walk to the library, and then put the top back on the bottle."

Semantic errors A semantic error is when your description of the steps to take is 11 syntactically perfect and in the right order, but there is simply a mistake in the program. The program is perfectly correct but it does not do what you intended for it to do. A simple example would be if you were giving a person directions to a restaurant and said, "... when you reach the intersection with the gas station, turn left and go one mile and the restaurant is a red building on your left." Your friend is very late and calls you to tell you that they are on a farm and walking around behind a barn, with no sign of a restaurant. Then you say "did you turn left or right at the gas station?" and they say, "I followed your directions perfectly, I have them written down, it says turn left and go one mile at the gas station." Then you say, "I am very sorry, because

while my instructions were syntactically correct, they sadly contained a small 1
but undetected semantic error.”.

Again in all three types of errors, Python is merely trying its hardest to do exactly 2
what you have asked.

1.11 Debugging 3

When Python spits out an error or even when it gives you a result that is different 4
from what you had intended, then begins the hunt for the cause of the error.
Debugging is the process of finding the cause of the error in your code. When you
are debugging a program, and especially if you are working on a hard bug, there
are four things to try:

reading Examine your code, read it back to yourself, and check that it says what 5
you meant to say.

running Experiment by making changes and running different versions. Often 6
if you display the right thing at the right place in the program, the prob-
lem becomes obvious, but sometimes you have to spend some time to build
scaffolding.

ruminating Take some time to think! What kind of error is it: syntax, runtime, 7
semantic? What information can you get from the error messages, or from
the output of the program? What kind of error could cause the problem
you’re seeing? What did you change last, before the problem appeared?

retreating At some point, the best thing to do is back off, undoing recent changes, 8
until you get back to a program that works and that you understand. Then
you can start rebuilding.

Beginning programmers sometimes get stuck on one of these activities and forget 9
the others. Finding a hard bug requires reading, running, ruminating, and some-
times retreating. If you get stuck on one of these activities, try the others. Each
activity comes with its own failure mode.

For example, reading your code might help if the problem is a typographical error, 10
but not if the problem is a conceptual misunderstanding. If you don’t understand
what your program does, you can read it 100 times and never see the error, because
the error is in your head.

Running experiments can help, especially if you run small, simple tests. But if 11
you run experiments without thinking or reading your code, you might fall into
a pattern I call “random walk programming”, which is the process of making
random changes until the program does the right thing. Needless to say, random
walk programming can take a long time.

You have to take time to think. Debugging is like an experimental science. You 12
should have at least one hypothesis about what the problem is. If there are two or
more possibilities, try to think of a test that would eliminate one of them.

Taking a break helps with the thinking. So does talking. If you explain the problem 13
to someone else (or even to yourself), you will sometimes find the answer before
you finish asking the question.

But even the best debugging techniques will fail if there are too many errors, or if the code you are trying to fix is too big and complicated. Sometimes the best option is to retreat, simplifying the program until you get to something that works and that you understand.

Beginning programmers are often reluctant to retreat because they can't stand to delete a line of code (even if it's wrong). If it makes you feel better, copy your program into another file before you start stripping it down. Then you can paste the pieces back in a little bit at a time.

1.12 The learning journey³

As you progress through the rest of the book, don't be afraid if the concepts don't seem to fit together well the first time. When you were learning to speak, it was not a problem for your first few years that you just made cute gurgling noises. And it was OK if it took six months for you to move from simple vocabulary to simple sentences and took 5-6 more years to move from sentences to paragraphs, and a few more years to be able to write an interesting complete short story on your own.

We want you to learn Python much more rapidly, so we teach it all at the same time over the next few chapters. But it is like learning a new language that takes time to absorb and understand before it feels natural. That leads to some confusion as we visit and revisit topics to try to get you to see the big picture while we are defining the tiny fragments that make up that big picture. While the book is written linearly, and if you are taking a course it will progress in a linear fashion, don't hesitate to be very nonlinear in how you approach the material. Look forwards and backwards and read with a light touch. By skimming more advanced material without fully understanding the details, you can get a better understanding of the "why?" of programming. By reviewing previous material and even redoing earlier exercises, you will realize that you actually learned a lot of material even if the material you are currently staring at seems a bit impenetrable.

Usually when you are learning your first programming language, there are a few wonderful "Ah Hah!" moments where you can look up from pounding away at some rock with a hammer and chisel and step away and see that you are indeed building a beautiful sculpture.

If something seems particularly hard, there is usually no value in staying up all night and staring at it. Take a break, take a nap, have a snack, explain what you are having a problem with to someone (or perhaps your dog), and then come back to it with fresh eyes. I assure you that once you learn the programming concepts in the book you will look back and see that it was all really easy and elegant and it simply took you a bit of time to absorb it.

1.13 Glossary⁸

bug An error in a program.

central processing unit The heart of any computer. It is what runs the software that we write; also called "CPU" or "the processor".

compile	To translate a program written in a high-level language into a low-level language all at once, in preparation for later execution.	1
high-level language	A programming language like Python that is designed to be easy for humans to read and write.	2
interactive mode	A way of using the Python interpreter by typing commands and expressions at the prompt.	3
interpret	To execute a program in a high-level language by translating it one line at a time.	4
low-level language	A programming language that is designed to be easy for a computer to execute; also called “machine code” or “assembly language”.	5
machine code	The lowest-level language for software, which is the language that is directly executed by the central processing unit (CPU).	6
main memory	Stores programs and data. Main memory loses its information when the power is turned off.	7
parse	To examine a program and analyze the syntactic structure.	8
portability	A property of a program that can run on more than one kind of computer.	9
print function	An instruction that causes the Python interpreter to display a value on the screen.	10
problem solving	The process of formulating a problem, finding a solution, and expressing the solution.	11
program	A set of instructions that specifies a computation.	12
prompt	When a program displays a message and pauses for the user to type some input to the program.	13
secondary memory	Stores programs and data and retains its information even when the power is turned off. Generally slower than main memory. Examples of secondary memory include disk drives and flash memory in USB sticks.	14
semantics	The meaning of a program.	15
semantic error	An error in a program that makes it do something other than what the programmer intended.	16
source code	A program in a high-level language.	17

1.14 Exercises¹⁸

Exercise 1: What is the function of the secondary memory in a computer?¹⁹

- a) Execute all of the computation and logic of the program
- b) Retrieve web pages over the Internet
- c) Store information for the long term, even beyond a power cycle
- d) Take input from the user

Exercise 2: What is a program?²¹

Exercise 3: What is the difference between a compiler and an interpreter?²²

Exercise 4: Which of the following contains “machine code”?²³

- a) The Python interpreter
- b) The keyboard
- c) Python source file
- d) A word processing document

Exercise 5: What is wrong with the following code?²⁵

```
>>> print 'Hello world!' 1
File "<stdin>", line 1
    print 'Hello world!'
           ^
SyntaxError: invalid syntax
>>>
```

Exercise 6: Where in the computer is a variable such as “x” stored after the following Python line finishes? 2

x = 123 3

- a) Central processing unit 4
- b) Main Memory
- c) Secondary Memory
- d) Input Devices
- e) Output Devices

Exercise 7: What will the following program print out: 5

```
x = 43 6
x = x - 1 7
print(x) 8
```

- a) 43 9
- b) 42
- c) x - 1
- d) Error because $x = x - 1$ is not possible mathematically

Exercise 8: Explain each of the following using an example of a human capability: 10
(1) Central processing unit, (2) Main Memory, (3) Secondary Memory, (4) Input Device, and (5) Output Device. For example, “What is the human equivalent to a Central Processing Unit”?

Chapter 2¹

Variables, expressions, and² statements

2.1 Values and types³

A *value* is one of the basic things a program works with, like a letter or a number.⁴
The values we have seen so far are 1, 2, and “Hello, World!”

These values belong to different *types*: 2 is an integer, and “Hello, World!” is a ⁵
string, so called because it contains a “string” of letters. You (and the interpreter)
can identify strings because they are enclosed in quotation marks.

The `print` statement also works for integers. We use the `python` command to start ⁶
the interpreter.

```
python      7  
>>> print(4)  
4
```

If you are not sure what type a value has, the interpreter can tell you.⁸

```
>>> type('Hello, World!') 9  
<class 'str'>  
>>> type(17)  
<class 'int'>
```

Not surprisingly, strings belong to the type `str` and integers belong to the type `int`.¹⁰
Less obviously, numbers with a decimal point belong to a type called `float`,
because these numbers are represented in a format called *floating point*.

```
>>> type(3.2) 11  
<class 'float'>
```

What about values like “17” and “3.2”? They look like numbers, but they are in ¹²
quotation marks like strings.

```
>>> type('17') 1
<class 'str'>
>>> type('3.2')
<class 'str'>
```

They're strings.²

When you type a large integer, you might be tempted to use commas between groups of three digits, as in 1,000,000. This is not a legal integer in Python, but it is legal:³

```
>>> print(1,000,000)4
1 0 0
```

Well, that's not what we expected at all! Python interprets 1,000,000 as a comma-separated sequence of integers, which it prints with spaces between.⁵

This is the first example we have seen of a semantic error: the code runs without producing an error message, but it doesn't do the “right” thing.⁶

2.2 Variables⁷

One of the most powerful features of a programming language is the ability to manipulate *variables*. A variable is a name that refers to a value.⁸

An *assignment statement* creates new variables and gives them values:⁹

```
>>> message = 'And now for something completely different'10
>>> n = 17
>>> pi = 3.1415926535897931
```

This example makes three assignments. The first assigns a string to a new variable named `message`; the second assigns the integer 17 to `n`; the third assigns the (approximate) value of π to `pi`.¹¹

To display the value of a variable, you can use a `print` statement:¹²

```
>>> print(n) 13
17
>>> print(pi)
3.141592653589793
```

The type of a variable is the type of the value it refers to.¹⁴

```
>>> type(message)15
<class 'str'>
>>> type(n)
<class 'int'>
>>> type(pi)
<class 'float'>
```

2.3 Variable names and keywords¹

Programmers generally choose names for their variables that are meaningful and document what the variable is used for.²

Variable names can be arbitrarily long. They can contain both letters and numbers,³ but they cannot start with a number. It is legal to use uppercase letters, but it is a good idea to begin variable names with a lowercase letter (you'll see why later).

The underscore character (`_`) can appear in a name. It is often used in names with multiple words, such as `my_name` or `airspeed_of_unladen_swallow`. Variable names can start with an underscore character, but we generally avoid doing this unless we are writing library code for others to use.⁴

If you give a variable an illegal name, you get a syntax error:⁵

```
>>> 76trombones = 'big parade'                                6
SyntaxError: invalid syntax
>>> more@ = 1000000
SyntaxError: invalid syntax
>>> class = 'Advanced Theoretical Zymurgy'
SyntaxError: invalid syntax
```

`76trombones` is illegal because it begins with a number. `more@` is illegal because it contains an illegal character, `@`. But what's wrong with `class`?⁷

It turns out that `class` is one of Python's *keywords*. The interpreter uses keywords to recognize the structure of the program, and they cannot be used as variable names.⁸

Python reserves 35 keywords:⁹

<code>False</code>	<code>await</code>	<code>else</code>	<code>import</code>	<code>pass</code>	¹⁰
<code>None</code>	<code>break</code>	<code>except</code>	<code>in</code>	<code>raise</code>	
<code>True</code>	<code>class</code>	<code>finally</code>	<code>is</code>	<code>return</code>	
<code>and</code>	<code>continue</code>	<code>for</code>	<code>lambda</code>	<code>try</code>	
<code>as</code>	<code>def</code>	<code>from</code>	<code>nonlocal</code>	<code>while</code>	
<code>assert</code>	<code>del</code>	<code>global</code>	<code>not</code>	<code>with</code>	
<code>async</code>	<code>elif</code>	<code>if</code>	<code>or</code>	<code>yield</code>	

You might want to keep this list handy. If the interpreter complains about one of your variable names and you don't know why, see if it is on this list.¹¹

2.4 Statements¹²

A *statement* is a unit of code that the Python interpreter can execute. We have seen two kinds of statements: `print` being an expression statement and assignment.¹³

When you type a statement in interactive mode, the interpreter executes it and displays the result, if there is one.¹⁴

A script usually contains a sequence of statements. If there is more than one **1** statement, the results appear one at a time as the statements execute.

For example, the script **2**

```
print(1) 3  
x = 2 4  
print(x) 5
```

produces the output **6**

```
1 7  
2
```

The assignment statement produces no output. **8**

2.5 Operators and operands **9**

Operators are special symbols that represent computations like addition and multiplication. The values the operator is applied to are called *operands*. **10**

The operators `+`, `-`, `*`, `/`, and `**` perform addition, subtraction, multiplication, **11** division, and exponentiation, as in the following examples:

```
20+32 12  
hour-1 13  
hour*60+minute 14  
minute/60 15  
5**2 16  
(5+9)*(15-7) 17
```

There has been a change in the division operator between Python 2 and Python 3. **18**
In Python 3, the result of this division is a floating point result:

```
>>> minute = 59 19  
>>> minute/60 20  
0.9833333333333333 21
```

The division operator in Python 2 would divide two integers and truncate the **22** result to an integer:

```
>>> minute = 59 23  
>>> minute/60 24  
0 25
```

To obtain the same answer in Python 3 use floored (`// integer`) division. **26**

```
>>> minute = 591
>>> minute//60
0
```

In Python 3 integer division functions much more as you would expect if you ² entered the expression on a calculator.

2.6 Expressions ³

An *expression* is a combination of values, variables, and operators. A value all by ⁴ itself is considered an expression, and so is a variable, so the following are all legal expressions (assuming that the variable `x` has been assigned a value):

```
17      5
x
x + 17
```

If you type an expression in interactive mode, the interpreter *evaluates* it and ⁶ displays the result:

```
>>> 1 + 17
2
```

But in a script, an expression all by itself doesn't do anything! This is a common ⁸ source of confusion for beginners.

Exercise 1: Type the following statements in the Python interpreter to see what ⁹ they do:

```
5      10
x = 5
x + 1
```

2.7 Order of operations ¹¹

When more than one operator appears in an expression, the order of evaluation ¹² depends on the *rules of precedence*. For mathematical operators, Python follows mathematical convention. The acronym *PEMDAS* is a useful way to remember the rules:

- Parentheses have the highest precedence and can be used to force an expression to evaluate in the order you want. Since expressions in parentheses are evaluated first, `2 * (3-1)` is 4, and `(1+1)**(5-2)` is 8. You can also use parentheses to make an expression easier to read, as in `(minute * 100) / 60`, even if it doesn't change the result. ¹³

- Exponentiation has the next highest precedence, so $2**1+1$ is 3, not 4, and $3*1**3$ is 3, not 27.
- Multiplication and Division have the same precedence, which is higher than Addition and Subtraction, which also have the same precedence. So $2*3-1$ is 5, not 4, and $6+4/2$ is 8, not 5.
- Operators with the same precedence are evaluated from left to right. So the expression $5-3-1$ is 1, not 3, because the $5-3$ happens first and then 1 is subtracted from 2.

When in doubt, always put parentheses in your expressions to make sure the computations are performed in the order you intend. 2

2.8 Modulus operator³

The *modulus operator* works on integers and yields the remainder when the first operand is divided by the second. In Python, the modulus operator is a percent sign (%). The syntax is the same as for other operators: 4

```
>>> quotient = 7 // 3 5
>>> print(quotient)
2
>>> remainder = 7 % 3
>>> print(remainder)
1
```

So 7 divided by 3 is 2 with 1 left over. 6

The modulus operator turns out to be surprisingly useful. For example, you can check whether one number is divisible by another: if $x \% y$ is zero, then x is divisible by y . 7

You can also extract the right-most digit or digits from a number. For example, $x \% 10$ yields the right-most digit of x (in base 10). Similarly, $x \% 100$ yields the last two digits. 8

2.9 String operations⁹

The + operator works with strings, but it is not addition in the mathematical sense. Instead it performs *concatenation*, which means joining the strings by linking them end to end. For example: 10

```
>>> first = 10 11
>>> second = 15
>>> print(first+second)
25
>>> first = '100'
```

```
>>> second = '150'      1
>>> print(first + second)
100150
```

The * operator also works with strings by multiplying the content of a string by 2 an integer. For example:

```
>>> first = 'Test '      3
>>> second = 3
>>> print(first * second)
Test Test Test
```

2.10 Asking the user for input⁴

Sometimes we would like to take the value for a variable from the user via their 5 keyboard. Python provides a built-in function called `input` that gets input from the keyboard¹. When this function is called, the program stops and waits for the user to type something. When the user presses Return or Enter, the program resumes and `input` returns what the user typed as a string.

```
>>> inp = input() 6
Some silly stuff
>>> print(inp)
Some silly stuff
```

Before getting input from the user, it is a good idea to print a prompt telling the 7 user what to input. You can pass a string to `input` to be displayed to the user before pausing for input:

```
>>> name = input('What is your name?\n') 8
What is your name?
Chuck
>>> print(name)
Chuck
```

The sequence \n at the end of the prompt represents a *newline*, which is a special 9 character that causes a line break. That's why the user's input appears below the prompt.

If you expect the user to type an integer, you can try to convert the return value 10 to `int` using the `int()` function:

```
>>> prompt = 'What...is the airspeed velocity of an unladen swallow?\n' 11
>>> speed = int(input(prompt))
What...is the airspeed velocity of an unladen swallow?
17
```

¹In Python 2, this function was named `raw_input`.

```
>>> int(speed)    1
17
>>> int(speed) + 5
22
```

But if the user types something other than a string of digits, you get an error:²

```
>>> speed = input(prompt)
What... is the airspeed velocity of an unladen swallow? 3
What do you mean, an African or a European swallow?
>>> int(speed)
ValueError: invalid literal for int() with
base 10: 'What do you mean, an African or a European swallow?'
```

We will see how to handle this kind of error later.⁴

2.11 Comments⁵

As programs get bigger and more complicated, they get more difficult to read.⁶
Formal languages are dense, and it is often difficult to look at a piece of code and figure out what it is doing, or why.

For this reason, it is a good idea to add notes to your programs to explain in natural language what the program is doing. These notes are called *comments*, and in Python they start with the # symbol:⁷

```
# compute the percentage of the hour that has elapsed 8
percentage = (minute * 100) / 60
```

In this case, the comment appears on a line by itself. You can also put comments at the end of a line:⁹

```
percentage = (minute * 100) / 60      # percentage of an hour 10
```

Everything from the # to the end of the line is ignored; it has no effect on the program.¹¹

Comments are most useful when they document non-obvious features of the code.¹² It is reasonable to assume that the reader can figure out *what* the code does; it is much more useful to explain *why*.

This comment is redundant with the code and useless:¹³

```
v = 5      # assign 5 to v 14
```

This comment contains useful information that is not in the code:¹⁵

```
v = 5      # velocity in meters/second. 16
```

Good variable names can reduce the need for comments, but long names can make complex expressions hard to read, so there is a trade-off.¹⁷

2.12 Choosing mnemonic variable names¹

As long as you follow the simple rules of variable naming, and avoid reserved words, you have a lot of choice when you name your variables. In the beginning, this choice can be confusing both when you read a program and when you write your own programs. For example, the following three programs are identical in terms of what they accomplish, but very different when you read them and try to understand them.

```
a = 35.0 3
```

```
b = 12.50 4
```

```
c = a * b 5
```

```
print(c) 6
```

```
hours = 35.0 7
```

```
rate = 12.50 8
```

```
pay = hours * rate 9
```

```
print(pay) 10
```

```
x1q3z9ahd = 35.0 11
```

```
x1q3z9afd = 12.50 12
```

```
x1q3p9afd = x1q3z9ahd * x1q3z9afd 13
```

```
print(x1q3p9afd) 14
```

The Python interpreter sees all three of these programs as *exactly the same* but humans see and understand these programs quite differently. Humans will most quickly understand the *intent* of the second program because the programmer has chosen variable names that reflect their intent regarding what data will be stored in each variable.

We call these wisely chosen variable names “mnemonic variable names”. The word *mnemonic*² means “memory aid”. We choose mnemonic variable names to help us remember why we created the variable in the first place.

While this all sounds great, and it is a very good idea to use mnemonic variable names, mnemonic variable names can get in the way of a beginning programmer’s ability to parse and understand code. This is because beginning programmers have not yet memorized the reserved words (there are only 35 of them) and sometimes variables with names that are too descriptive start to look like part of the language and not just well-chosen variable names.

Take a quick look at the following Python sample code which loops through some data. We will cover loops soon, but for now try to just puzzle through what this means:

```
for word in words: 19
    print(word)
```

²See <https://en.wikipedia.org/wiki/Mnemonic> for an extended description of the word “mnemonic”.

What is happening here? Which of the tokens (for, word, in, etc.) are reserved words and which are just variable names? Does Python understand at a fundamental level the notion of words? Beginning programmers have trouble separating what parts of the code *must* be the same as this example and what parts of the code are simply choices made by the programmer.

The following code is equivalent to the above code:²

```
for slice in pizza:3
    print(slice)
```

It is easier for the beginning programmer to look at this code and know which parts are reserved words defined by Python and which parts are simply variable names chosen by the programmer. It is pretty clear that Python has no fundamental understanding of pizza and slices and the fact that a pizza consists of a set of one or more slices.

But if our program is truly about reading data and looking for words in the data,⁵ **pizza** and **slice** are very un-mnemonic variable names. Choosing them as variable names distracts from the meaning of the program.

After a pretty short period of time, you will know the most common reserved words⁶ and you will start to see the reserved words jumping out at you:

The parts of the code that are defined by Python (**for**, **in**, **print**, and **:**) are in bold and the programmer-chosen variables (**word** and **words**) are not in bold. Many text editors are aware of Python syntax and will color reserved words differently to give you clues to keep your variables and reserved words separate. After a while you will begin to read Python and quickly determine what is a variable and what is a reserved word.

2.13 Debugging⁸

At this point, the syntax error you are most likely to make is an illegal variable name, like **class** and **yield**, which are keywords, or **odd-job** and **US\$**, which contain illegal characters.

If you put a space in a variable name, Python thinks it is two operands without an operator:¹⁰

```
>>> bad name = 511
SyntaxError: invalid syntax12
```

For syntax errors, the error messages don't help much. The most common message¹³ is **SyntaxError: invalid syntax** which is not very informative.

The runtime error you are most likely to make is a "use before def;" that is, trying¹⁴ to use a variable before you have assigned a value. This can happen if you spell a variable name wrong:

```
>>> principal = 327.68
>>> interest = principle * rate
NameError: name 'principle' is not defined
```

1

Variables names are case sensitive, so `\LaTeX` is not the same as `latex`.²

At this point, the most likely cause of a semantic error is the order of operations.³
For example, to evaluate $1/2\pi$, you might be tempted to write

```
>>> 1.0 / 2.0 * pi
```

4

But the division happens first, so you would get $\pi/2$, which is not the same thing!⁵
There is no way for Python to know what you meant to write, so in this case you don't get an error message; you just get the wrong answer.

2.14 Glossary⁶

assignment A statement that assigns a value to a variable.⁷

concatenate To join two operands end to end.⁸

comment Information in a program that is meant for other programmers (or anyone reading the source code) and has no effect on the execution of the program.

evaluate To simplify an expression by performing the operations in order to yield¹⁰ a single value.

expression A combination of variables, operators, and values that represents a¹¹ single result value.

floating point A type that represents numbers with fractional parts.¹²

integer A type that represents whole numbers.¹³

keyword A reserved word that is used by the compiler to parse a program; you¹⁴ cannot use keywords like `if`, `def`, and `while` as variable names.

mnemonic A memory aid. We often give variables mnemonic names to help us¹⁵ remember what is stored in the variable.

modulus operator An operator, denoted with a percent sign (%), that works on¹⁶ integers and yields the remainder when one number is divided by another.

operand One of the values on which an operator operates.¹⁷

operator A special symbol that represents a simple computation like addition,¹⁸ multiplication, or string concatenation.

rules of precedence The set of rules governing the order in which expressions¹⁹ involving multiple operators and operands are evaluated.

statement A section of code that represents a command or action. So far, the²⁰ statements we have seen are assignments and print expression statement.

string A type that represents sequences of characters.²¹

type A category of values. The types we have seen so far are integers (type `int`),²² floating-point numbers (type `float`), and strings (type `str`).

value One of the basic units of data, like a number or string, that a program²³ manipulates.

variable A name that refers to a value.²⁴

2.15 Exercises¹

Exercise 2: Write a program that uses `input` to prompt a user for their name² and then welcomes them.

```
Enter your name: Chuck3
```

```
Hello Chuck4
```

Exercise 3: Write a program to prompt the user for hours and rate per hour to⁵ compute gross pay.

```
Enter Hours: 356
```

```
Enter Rate: 2.757
```

```
Pay: 96.258
```

We won't worry about making sure our pay has exactly two digits after the decimal⁹ place for now. If you want, you can play with the built-in Python `round` function to properly round the resulting pay to two decimal places.

Exercise 4: Assume that we execute the following assignment statements:¹⁰

```
width = 1711
```

```
height = 12.012
```

For each of the following expressions, write the value of the expression and the¹³ type (of the value of the expression).

```
1. width//214
```

```
2. width/2.0
```

```
3. height/3
```

```
4. 1 + 2 * 5
```

Use the Python interpreter to check your answers.¹⁵

Exercise 5: Write a program which prompts the user for a Celsius temperature,¹⁶ convert the temperature to Fahrenheit, and print out the converted temperature.

Chapter 3¹

Conditional execution²

3.1 Boolean expressions³

A *boolean expression* is an expression that is either true or false. The following ⁴ examples use the operator `==`, which compares two operands and produces `True` if they are equal and `False` otherwise:

```
>>> 5 == 5 5
True 6
>>> 5 == 6 7
False 8
```

`True` and `False` are special values that belong to the class `bool`; they are not ⁹ strings:

```
>>> type(True) 10
<class 'bool'> 11
>>> type(False) 12
<class 'bool'> 13
```

The `==` operator is one of the *comparison operators*; the others are: ¹⁴

```
x != y          # x is not equal to y 15
x > y           # x is greater than y 16
x < y           # x is less than y 17
x >= y          # x is greater than or equal to y 18
x <= y          # x is less than or equal to y 19
x is y          # x is the same as y 20
x is not y      # x is not the same as y 21
```

Although these operations are probably familiar to you, the Python symbols are ²² different from the mathematical symbols for the same operations. A common error is to use a single equal sign (`=`) instead of a double equal sign (`==`). Remember that `=` is an assignment operator and `==` is a comparison operator. There is no such thing as `=<` or `=>`.

3.2 Logical operators¹

There are three *logical operators*: `and`, `or`, and `not`. The semantics (meaning) of ² these operators is similar to their meaning in English. For example,

```
x > 0 and x < 103
```

is true only if `x` is greater than 0 *and* less than 10.⁴

`n%2 == 0 or n%3 == 0` is true if *either* of the conditions is true, that is, if the ⁵ number is divisible by 2 *or* 3.

Finally, the `not` operator negates a boolean expression, so `not (x > y)` is true if ⁶ `x > y` is false.

```
>>> x = 17
>>> y = 2
>>> x > y
False
>>> not (x > y)
True
```

Strictly speaking, the operands of the logical operators should be boolean expressions, but Python is not very strict. Any nonzero number is interpreted as “true.”⁸

```
>>> 17 and True9
True
```

This flexibility can be useful in some situations, but there are some subtleties to it ¹⁰ that might be confusing. You might want to avoid it until you are sure you know what you are doing.

3.3 Conditional execution¹¹

In order to write useful programs, we almost always need the ability to check conditions and change the behavior of the program accordingly. *Conditional statements* give us this ability. The simplest form is the `if` statement:¹²

```
if x > 0 :13
    print('x is positive')
```

The boolean expression after the `if` statement is called the *condition*. We end the ¹⁴ `if` statement with a colon character (`:`) and the line(s) after the `if` statement are indented.

If the logical condition is true, then the indented statement gets executed. If the ¹⁵ logical condition is false, the indented statement is skipped.

`if` statements have the same structure as function definitions or `for` loops¹. The ¹⁶ statement consists of a header line that ends with the colon character (`:`) followed

¹We will learn about functions in Chapter 4 and loops in Chapter 5.




Figure 3.1: If Logic 2

by an indented block. Statements like this are called *compound statements* because 3 they stretch across more than one line.

```
if x > y: 4
    print(x)
    print(y)
```

There is no limit on the number of statements that can appear in the body, but 5 there must be at least one. Occasionally, it is useful to have a body with no statements (usually as a place holder for code you haven't written yet). In that case, you can use the `pass` statement to pass the Python interpreter check, which does nothing.

```
if x < 0 :6
    pass  # need to handle negative values, do nothing for now.7
```

If you enter an `if` statement in the Python interpreter, the prompt will change 8 from three chevrons (`>>>`) to three dots (...) to indicate you are in the middle of a block of statements, as shown below:

```
>>> x = 3 9
>>> if x < 10:
...     print('Small')
...
Small
>>>
```

When using the Python interpreter, you must leave a blank line at the end of a 10 block, otherwise Python will return an error:

```
>>> x = 3 11
>>> if x < 10:
...     print('Small')
... print('Done')
File "<stdin>", line 3
    print('Done')
    ^
SyntaxError: invalid syntax
```




Figure 3.2: If-Then-Else Logic 2

A blank line at the end of a block of statements is not necessary when writing and 3 executing a script, but it may improve readability of your code.

3.4 Alternative execution⁴

A second form of the `if` statement is *alternative execution*, in which there are two 5 possibilities and the condition determines which one gets executed. The syntax looks like this:

```

if x % 2 == 0:          6
    print('x is even')
else:
    print('x is odd')
    
```

If the remainder when `x` is divided by 2 is 0, then we know that `x` is even, and the 7 program displays a message to that effect. If the condition is false, the second set of statements is executed.

Since the condition must either be true or false, exactly one of the alternatives will 8 be executed. The alternatives are called *branches*, because they are branches in the flow of execution.

3.5 Chained conditionals⁹

Sometimes there are more than two possibilities and we need more than two 10 branches. One way to express a computation like that is a *chained conditional*:

```

if x < y:                11
    print('x is less than y')
elif x > y:
    print('x is greater than y')
else:
    print('x and y are equal')
    
```




Figure 3.3: If-Then-ElseIf Logic 2

`elif` is an abbreviation of “else if.” Again, exactly one branch will be executed.³

There is no limit on the number of `elif` statements. If there is an `else` clause, it⁴ has to be at the end, but there doesn’t have to be one.

```

if choice == 'a':
    print('Bad guess')
elif choice == 'b':
    print('Good guess')
elif choice == 'c':
    print('Close, but not correct')

```

Each condition is checked in order. If the first is false, the next is checked, and so on. If one of them is true, the corresponding branch executes, and the statement ends. Even if more than one condition is true, only the first true branch executes.⁵

3.6 Nested conditionals⁷


One conditional can also be nested within another. We could have written the⁸ three-branch example like this:

```

if x == y:
    print('x and y are equal')
else:
    if x < y:
        print('x is less than y')
    else:
        print('x is greater than y')

```

The outer conditional contains two branches. The first branch contains a simple statement. The second branch contains another `if` statement, which has two branches of its own. Those two branches are both simple statements, although they could have been conditional statements as well.¹⁰

Figure 3.4: Nested If Statements 2

Although the indentation of the statements makes the structure apparent, *nested conditionals* become difficult to read very quickly. In general, it is a good idea to avoid them when you can. 3

Logical operators often provide a way to simplify nested conditional statements. 4 For example, we can rewrite the following code using a single conditional:

```
if 0 < x:  
    if x < 10:  
        print('x is a positive single-digit number.')5
```

The `print` statement is executed only when we pass both conditionals. We can 6 get the same effect with the `and` operator:

```
if 0 < x and x < 10:7  
    print('x is a positive single-digit number.)
```

3.7 Catching exceptions using `try` and `except` 8

Earlier we saw a code segment where we used the `input` and `int` functions to read 9 and parse an integer number entered by the user. We also saw how treacherous doing this could be:

```
>>> prompt = "What is the air velocity of an unladen swallow?\n"10  
>>> speed = input(prompt)  
What is the air velocity of an unladen swallow?  
What do you mean, an African or a European swallow?  
>>> int(speed)  
ValueError: invalid literal for int() with  
base 10: 'What do you mean, an African or a European swallow?'  
>>>
```

When we are executing these statements in the Python interpreter, we get a new 1 prompt from the interpreter, think “oops”, and move on to our next statement.

However if you place this code in a Python script and this error occurs, your script 2 immediately stops in its tracks with a traceback. It does not execute the following statement.

Here is a sample program to convert a Fahrenheit temperature to a Celsius tem- 3 perature:

```
inp = input('Enter Fahrenheit Temperature: ') 4
fahr = float(inp)
cel = (fahr - 32.0) * 5.0 / 9.0
print(cel)
```

Code: <https://www.py4e.com/code3/fahren.py> 5

If we execute this code and give it invalid input, it simply fails with an unfriendly 6 error message:

```
python fahren.py 7
Enter Fahrenheit Temperature:72
22.22222222222222
```

```
python fahren.py 8
Enter Fahrenheit Temperature:fred
Traceback (most recent call last):
  File "fahren.py", line 2, in <module>
    fahr = float(inp)
ValueError: could not convert string to float: 'fred'
```

There is a conditional execution structure built into Python to handle these types 9 of expected and unexpected errors called “try / except”. The purpose of `try` and `except` is that you know that some sequence of instruction(s) may have a problem and you want to add some statements to be executed if an error occurs. These extra statements (the `except` block) are ignored if there is no error.

You can think of the `try` and `except` feature in Python as an “insurance policy” 10 on a sequence of statements.

We can rewrite our temperature converter as follows: 11

```
inp = input('Enter Fahrenheit Temperature:') 12
try:
    fahr = float(inp)
    cel = (fahr - 32.0) * 5.0 / 9.0
    print(cel)
except:
    print('Please enter a number')
```

Code: <https://www.py4e.com/code3/fahren2.py> 13

Python starts by executing the sequence of statements in the `try` block. If all goes well, it skips the `except` block and proceeds. If an exception occurs in the `try` block, Python jumps out of the `try` block and executes the sequence of statements in the `except` block.

```
python fahren2.py          2
Enter Fahrenheit Temperature:72
22.22222222222222
```

```
python fahren2.py          3
Enter Fahrenheit Temperature:fred
Please enter a number
```

Handling an exception with a `try` statement is called *catching* an exception. In this example, the `except` clause prints an error message. In general, catching an exception gives you a chance to fix the problem, or try again, or at least end the program gracefully.

3.8 Short-circuit evaluation of logical expressions⁵

When Python is processing a logical expression such as `x >= 2 and (x/y) > 2`, it evaluates the expression from left to right. Because of the definition of `and`, if `x` is less than 2, the expression `x >= 2` is `False` and so the whole expression is `False` regardless of whether `(x/y) > 2` evaluates to `True` or `False`.

When Python detects that there is nothing to be gained by evaluating the rest of a logical expression, it stops its evaluation and does not do the computations in the rest of the logical expression. When the evaluation of a logical expression stops because the overall value is already known, it is called *short-circuiting* the evaluation.

While this may seem like a fine point, the short-circuit behavior leads to a clever technique called the *guardian pattern*. Consider the following code sequence in the Python interpreter:

```
>>> x = 6                      9
>>> y = 2
>>> x >= 2 and (x/y) > 2
True
>>> x = 1
>>> y = 0
>>> x >= 2 and (x/y) > 2
False
>>> x = 6
>>> y = 0
>>> x >= 2 and (x/y) > 2
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ZeroDivisionError: division by zero
>>>
```

The third calculation failed because Python was evaluating (x/y) and y was zero, which causes a runtime error. But the first and the second examples did *not* fail because in the first calculation y was non zero and in the second one the first part of these expressions $x \geq 2$ evaluated to `False` so the (x/y) was not ever executed due to the *short-circuit* rule and there was no error. 1

We can construct the logical expression to strategically place a *guard* evaluation just before the evaluation that might cause an error as follows: 2

```
>>> x = 1
>>> y = 0
>>> x >= 2 and y != 0 and (x/y) > 2
False
>>> x = 6
>>> y = 0
>>> x >= 2 and y != 0 and (x/y) > 2
False
>>> x >= 2 and (x/y) > 2 and y != 0
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ZeroDivisionError: division by zero
>>>
```

3

In the first logical expression, $x \geq 2$ is `False` so the evaluation stops at the `and`. 4
 In the second logical expression, $x \geq 2$ is `True` but $y \neq 0$ is `False` so we never reach (x/y) .

In the third logical expression, the $y \neq 0$ is *after* the (x/y) calculation so the expression fails with an error. 5

In the second expression, we say that $y \neq 0$ acts as a *guard* to ensure that we only execute (x/y) if y is non-zero. 6

3.9 Debugging 7

The traceback Python displays when an error occurs contains a lot of information, 8
 but it can be overwhelming. The most useful parts are usually:

- What kind of error it was, and 9
- Where it occurred.

Syntax errors are usually easy to find, but there are a few gotchas. Whitespace errors can be tricky because spaces and tabs are invisible and we are used to ignoring them. 10

```
>>> x = 5
>>> y = 6
  File "<stdin>", line 1
    y = 6
    ^
```

11

`IndentationError: unexpected indent` 12

In this example, the problem is that the second line is indented by one space. But the error message points to `y`, which is misleading. In general, error messages indicate where the problem was discovered, but the actual error might be earlier in the code, sometimes on a previous line. 1

In general, error messages tell you where the problem was discovered, but that is often not where it was caused. 2

3.10 Glossary³

body The sequence of statements within a compound statement. 4

boolean expression An expression whose value is either `True` or `False`. 5

branch One of the alternative sequences of statements in a conditional statement. 6

chained conditional A conditional statement with a series of alternative branches. 7

comparison operator One of the operators that compares its operands: `==`, `!=`, `>`, `<`, `>=`, and `<=`. 8

conditional statement A statement that controls the flow of execution depending on some condition. 9

condition The boolean expression in a conditional statement that determines which branch is executed. 10

compound statement A statement that consists of a header and a body. The header ends with a colon (`:`). The body is indented relative to the header. 11

guardian pattern Where we construct a logical expression with additional comparisons to take advantage of the short-circuit behavior. 12

logical operator One of the operators that combines boolean expressions: `and`, `or`, and `not`. 13

nested conditional A conditional statement that appears in one of the branches of another conditional statement. 14

traceback A list of the functions that are executing, printed when an exception occurs. 15

short circuit When Python is part-way through evaluating a logical expression and stops the evaluation because Python knows the final value for the expression without needing to evaluate the rest of the expression. 16

3.11 Exercises¹⁷

Exercise 1: Rewrite your pay computation to give the employee 1.5 times the hourly rate for hours worked above 40 hours. 18

`Enter Hours: 45` 19

`Enter Rate: 10` 20

`Pay: 475.0` 21

Exercise 2: Rewrite your pay program using `try` and `except` so that your program handles non-numeric input gracefully by printing a message and exiting the program. The following shows two executions of the program: 22

```
Enter Hours: 20 1  
Enter Rate: nine 2  
Error, please enter numeric input 3
```

```
Enter Hours: forty 4  
Error, please enter numeric input 5
```

Exercise 3: Write a program to prompt for a score between 0.0 and 1.0. If the score is out of range, print an error message. If the score is between 0.0 and 1.0, print a grade using the following table: 6

Score	Grade
≥ 0.9	A
≥ 0.8	B
≥ 0.7	C
≥ 0.6	D
< 0.6	F

```
Enter score: 0.95 8  
A 9
```

```
Enter score: perfect 10  
Bad score 11
```

```
Enter score: 10.0 12  
Bad score 13
```

```
Enter score: 0.75 14  
C 15
```

```
Enter score: 0.5 16  
F 17
```

Run the program repeatedly as shown above to test the various different values for input. 18

Chapter 4¹

Functions²

4.1 Function calls³

In the context of programming, a *function* is a named sequence of statements that ⁴ performs a computation. When you define a function, you specify the name and the sequence of statements. Later, you can “call” the function by name. We have already seen one example of a *function call*:

```
>>> type(32) 5
<class 'int'>
```

The name of the function is `type`. The expression in parentheses is called the ⁶ *argument* of the function. The argument is a value or variable that we are passing into the function as input to the function. The result, for the `type` function, is the type of the argument.

It is common to say that a function “takes” an argument and “returns” a result.⁷ The result is called the *return value*.

4.2 Built-in functions⁸

Python provides a number of important built-in functions that we can use without ⁹ needing to provide the function definition. The creators of Python wrote a set of functions to solve common problems and included them in Python for us to use.

The `max` and `min` functions give us the largest and smallest values in a list, respectively:¹⁰

```
>>> max('Hello world') 11
'w'
>>> min('Hello world')
' '
>>>
```

The `max` function tells us the “largest character” in the string (which turns out to be the letter “w”) and the `min` function shows us the smallest character (which turns out to be a space).¹

Another very common built-in function is the `len` function which tells us how many items are in its argument. If the argument to `len` is a string, it returns the number of characters in the string.²

```
>>> len('Hello world') 3
11
>>>
```

These functions are not limited to looking at strings. They can operate on any set⁴ of values, as we will see in later chapters.

You should treat the names of built-in functions as reserved words (i.e., avoid using⁵ “max” as a variable name).

4.3 Type conversion functions⁶

Python also provides built-in functions that convert values from one type to another. The `int` function takes any value and converts it to an integer, if it can, or complains otherwise:⁷

```
>>> int('32') 8
32
>>> int('Hello')
ValueError: invalid literal for int() with base 10: 'Hello'
```

`int` can convert floating-point values to integers, but it doesn’t round off; it chops⁹ off the fraction part:

```
>>> int(3.99999) 10
3
>>> int(-2.3)
-2
```

`float` converts integers and strings to floating-point numbers:¹¹

```
>>> float(32) 12
32.0
>>> float('3.14159')
3.14159
```

Finally, `str` converts its argument to a string:¹³

```
>>> str(32) 14
'32'
>>> str(3.14159)
'3.14159'
```

4.4 Math functions¹

Python has a `math` module that provides most of the familiar mathematical functions. Before we can use the module, we have to import it:²

```
>>> import math3
```

This statement creates a *module object* named `math`. If you print the module⁴ object, you get some information about it:

```
>>> print(math)5  
<module 'math' (built-in)>
```

The module object contains the functions and variables defined in the module. To⁶ access one of the functions, you have to specify the name of the module and the name of the function, separated by a dot (also known as a period). This format is called *dot notation*.

```
>>> ratio = signal_power / noise_power7  
>>> decibels = 10 * math.log10(ratio)  
  
>>> radians = 0.78  
>>> height = math.sin(radians)
```

The first example computes the logarithm base 10 of the signal-to-noise ratio. The⁹ `math` module also provides a function called `log` that computes logarithms base e.

The second example finds the sine of `radians`. The name of the variable is a hint¹⁰ that `sin` and the other trigonometric functions (`cos`, `tan`, etc.) take arguments in radians. To convert from degrees to radians, divide by 360 and multiply by 2π :

```
>>> degrees = 4511  
>>> radians = degrees / 360.0 * 2 * math.pi  
>>> math.sin(radians)  
0.7071067811865476
```

The expression `math.pi` gets the variable `pi` from the `math` module. The value of¹² this variable is an approximation of π , accurate to about 15 digits.

If you know your trigonometry, you can check the previous result by comparing it¹³ to the square root of two divided by two:

```
>>> math.sqrt(2) / 2.014  
0.7071067811865476
```

4.5 Random numbers¹

Given the same inputs, most computer programs generate the same outputs every ² time, so they are said to be *deterministic*. Determinism is usually a good thing, since we expect the same calculation to yield the same result. For some applications, though, we want the computer to be unpredictable. Games are an obvious example, but there are more.

Making a program truly nondeterministic turns out to be not so easy, but there ³ are ways to make it at least seem nondeterministic. One of them is to use *algorithms* that generate *pseudorandom* numbers. Pseudorandom numbers are not truly random because they are generated by a deterministic computation, but just by looking at the numbers it is all but impossible to distinguish them from random.

The `random` module provides functions that generate pseudorandom numbers ⁴ (which I will simply call “random” from here on).

The function `random` returns a random float between 0.0 and 1.0 (including 0.0 ⁵ but not 1.0). Each time you call `random`, you get the next number in a long series. To see a sample, run this loop:

```
import random 6

for i in range(10): 7
    x = random.random()
    print(x)
```

This program produces the following list of 10 random numbers between 0.0 and ⁸ up to but not including 1.0.

```
0.11132867921152356 9
0.5950949227890241
0.04820265884996877
0.841003109276478
0.997914947094958
0.04842330803368111
0.7416295948208405
0.510535245390327
0.27447040171978143
0.028511805472785867
```

Exercise 1: Run the program on your system and see what numbers you get. Run ¹⁰ the program more than once and see what numbers you get.

The `random` function is only one of many functions that handle random numbers. ¹¹ The function `randint` takes the parameters `low` and `high`, and returns an integer between `low` and `high` (including both).

```
>>> random.randint(5, 10) 12
5
>>> random.randint(5, 10)
9
```

To choose an element from a sequence at random, you can use `choice`:¹

```
>>> t = [1, 2, 3] 2
>>> random.choice(t)
2
>>> random.choice(t)
3
```

The `random` module also provides functions to generate random values from continuous distributions including Gaussian, exponential, gamma, and a few more.³

4.6 Adding new functions⁴

So far, we have only been using the functions that come with Python, but it is also possible to add new functions. A *function definition* specifies the name of a new function and the sequence of statements that execute when the function is called. Once we define a function, we can reuse the function over and over throughout our program.⁵

Here is an example:⁶

```
def print_lyrics():
    print("I'm a lumberjack, and I'm okay.")
    print('I sleep all night and I work all day.') 7
```

`def` is a keyword that indicates that this is a function definition. The name of the function is `print_lyrics`. The rules for function names are the same as for variable names: letters, numbers and some punctuation marks are legal, but the first character can't be a number. You can't use a keyword as the name of a function, and you should avoid having a variable and a function with the same name.⁸

The empty parentheses after the name indicate that this function doesn't take any arguments. Later we will build functions that take arguments as their inputs.⁹

The first line of the function definition is called the *header*; the rest is called the *body*. The header has to end with a colon and the body has to be indented. By convention, the indentation is always four spaces. The body can contain any number of statements.¹⁰

If you type a function definition in interactive mode, the interpreter prints ellipses (...) to let you know that the definition isn't complete:¹¹

```
>>> def print_lyrics():
...     print("I'm a lumberjack, and I'm okay.")
...     print('I sleep all night and I work all day.')
... 12
```

To end the function, you have to enter an empty line (this is not necessary in a script).¹³

Defining a function creates a variable with the same name.¹⁴

```
>>> print(print_lyrics)          1
<function print_lyrics at 0xb7e99e9c>
>>> print(type(print_lyrics))
<class 'function'>
```

The value of `print_lyrics` is a *function object*, which has type “function”. 2

The syntax for calling the new function is the same as for built-in functions: 3

```
>>> print_lyrics()              4
I'm a lumberjack, and I'm okay.
I sleep all night and I work all day.
```

Once you have defined a function, you can use it inside another function. For example, to repeat the previous refrain, we could write a function called `repeat_lyrics`: 5

```
def repeat_lyrics(): 6
    print_lyrics()
    print_lyrics()
```

And then call `repeat_lyrics`: 7

```
>>> repeat_lyrics()             8
I'm a lumberjack, and I'm okay.
I sleep all night and I work all day.
I'm a lumberjack, and I'm okay.
I sleep all night and I work all day.
```

But that's not really how the song goes. 9

4.7 Definitions and uses 10

Pulling together the code fragments from the previous section, the whole program 11 looks like this:

```
def print_lyrics(): 12
    print("I'm a lumberjack, and I'm okay.")
    print('I sleep all night and I work all day.')
```

```
def repeat_lyrics(): 13
    print_lyrics()
    print_lyrics()
```

`repeat_lyrics()` 14

Code: <https://www.py4e.com/code3/lyrics.py> 15

This program has two function definitions: `print_lyrics` and `repeat_lyrics`.¹ Function definitions get executed just like other statements, but the effect is to create function objects. The statements inside the function do not get executed until the function is called, and the function definition generates no output.

As you might expect, you have to create a function before you can execute it. In² other words, the function definition has to be executed before the first time it is called.

Exercise 2: Move the last line of this program to the top, so the function call³ appears before the definitions. Run the program and see what error message you get.

Exercise 3: Move the function call back to the bottom and move the definition⁴ of `print_lyrics` after the definition of `repeat_lyrics`. What happens when you run this program?

4.8 Flow of execution⁵

In order to ensure that a function is defined before its first use, you have to know⁶ the order in which statements are executed, which is called the *flow of execution*.

Execution always begins at the first statement of the program. Statements are⁷ executed one at a time, in order from top to bottom.

Function *definitions* do not alter the flow of execution of the program, but re⁸member that statements inside the function are not executed until the function is called.

A function call is like a detour in the flow of execution. Instead of going to the next⁹ statement, the flow jumps to the body of the function, executes all the statements there, and then comes back to pick up where it left off.

That sounds simple enough, until you remember that one function can call another.¹⁰ While in the middle of one function, the program might have to execute the statements in another function. But while executing that new function, the program might have to execute yet another function!

Fortunately, Python is good at keeping track of where it is, so each time a function¹¹ completes, the program picks up where it left off in the function that called it. When it gets to the end of the program, it terminates.

What's the moral of this sordid tale? When you read a program, you don't always¹² want to read from top to bottom. Sometimes it makes more sense if you follow the flow of execution.

4.9 Parameters and arguments¹³

Some of the built-in functions we have seen require arguments. For example, when¹⁴ you call `math.sin` you pass a number as an argument. Some functions take more than one argument: `math.pow` takes two, the base and the exponent.

Inside the function, the arguments are assigned to variables called *parameters*.¹ Here is an example of a user-defined function that takes an argument:

```
def print_twice(bruce):2
    print(bruce)
    print(bruce)
```

This function assigns the argument to a parameter named `bruce`. When the function is called, it prints the value of the parameter (whatever it is) twice.³

This function works with any value that can be printed.⁴

```
>>> print_twice('Spam')5
Spam
Spam
>>> print_twice(17)
17
17
>>> import math
>>> print_twice(math.pi)
3.141592653589793
3.141592653589793
```

The same rules of composition that apply to built-in functions also apply to user-defined functions, so we can use any kind of expression as an argument for `print_twice`:⁶

```
>>> print_twice('Spam '*4)7
Spam Spam Spam Spam
Spam Spam Spam Spam
>>> print_twice(math.cos(math.pi))
-1.0
-1.0
```

The argument is evaluated before the function is called, so in the examples the expressions '`'Spam '`*4' and `math.cos(math.pi)` are only evaluated once.⁸

You can also use a variable as an argument:⁹

```
>>> michael = 'Eric, the half a bee.'10
>>> print_twice(michael)
Eric, the half a bee.
Eric, the half a bee.
```

The name of the variable we pass as an argument (`michael`) has nothing to do with the name of the parameter (`bruce`). It doesn't matter what the value was called back home (in the caller); here in `print_twice`, we call everybody `bruce`.¹¹

4.10 Fruitful functions and void functions¹

Some of the functions we are using, such as the math functions, yield results;² for lack of a better name, I call them *fruitful functions*. Other functions, like `print_twice`, perform an action but don't return a value. They are called *void functions*.

When you call a fruitful function, you almost always want to do something with³ the result; for example, you might assign it to a variable or use it as part of an expression:

```
x = math.cos(radians)          4  
golden = (math.sqrt(5) + 1) / 2
```

When you call a function in interactive mode, Python displays the result:⁵

```
>>> math.sqrt(5) 6  
2.23606797749979
```

But in a script, if you call a fruitful function and do not store the result of the⁷ function in a variable, the return value vanishes into the mist!

```
math.sqrt(5) 8
```

This script computes the square root of 5, but since it doesn't store the result in⁹ a variable or display the result, it is not very useful.

Void functions might display something on the screen or have some other effect,¹⁰ but they don't have a return value. If you try to assign the result to a variable, you get a special value called `None`.

```
>>> result = print_twice('Bing') 11  
Bing  
Bing  
>>> print(result)  
None
```

The value `None` is not the same as the string "None". It is a special value that has¹² its own type:

```
>>> print(type(None)) 13  
<class 'NoneType'>
```

To return a result from a function, we use the `return` statement in our function.¹⁴ For example, we could make a very simple function called `addtwo` that adds two numbers together and returns a result.

```
def addtwo(a, b): 1
    added = a + b
    return added

x = addtwo(3, 5)
print(x)
```

Code: <https://www.py4e.com/code3/addtwo.py> **2**

When this script executes, the `print` statement will print out “8” because the **3** `addtwo` function was called with 3 and 5 as arguments. Within the function, the parameters `a` and `b` were 3 and 5 respectively. The function computed the sum of the two numbers and placed it in the local function variable named `added`. Then it used the `return` statement to send the computed value back to the calling code as the function result, which was assigned to the variable `x` and printed out.

4.11 Why functions?⁴

It may not be clear why it is worth the trouble to divide a program into functions.⁵ There are several reasons:

- Creating a new function gives you an opportunity to name a group of statements, which makes your program easier to read, understand, and debug.⁶
- Functions can make a program smaller by eliminating repetitive code. Later, if you make a change, you only have to make it in one place.
- Dividing a long program into functions allows you to debug the parts one at a time and then assemble them into a working whole.
- Well-designed functions are often useful for many programs. Once you write and debug one, you can reuse it.

Throughout the rest of the book, often we will use a function definition to explain a concept. Part of the skill of creating and using functions is to have a function properly capture an idea such as “find the smallest value in a list of values”. Later we will show you code that finds the smallest in a list of values and we will present it to you as a function named `min` which takes a list of values as its argument and returns the smallest value in the list.⁷

4.12 Debugging⁸

If you are using a text editor to write your scripts, you might run into problems with spaces and tabs. The best way to avoid these problems is to use spaces exclusively (no tabs). Most text editors that know about Python do this by default, but some don’t.⁹

Tabs and spaces are usually invisible, which makes them hard to debug, so try to find an editor that manages indentation for you.¹⁰

Also, don't forget to save your program before you run it. Some development environments do this automatically, but some don't. In that case, the program you are looking at in the text editor is not the same as the program you are running. 1

Debugging can take a long time if you keep running the same incorrect program over and over! 2

Make sure that the code you are looking at is the code you are running. If you're not sure, put something like `print("hello")` at the beginning of the program and run it again. If you don't see `hello`, you're not running the right program! 3

4.13 Glossary 4

algorithm A general process for solving a category of problems. 5

argument A value provided to a function when the function is called. This value is assigned to the corresponding parameter in the function. 6

body The sequence of statements inside a function definition. 7

composition Using an expression as part of a larger expression, or a statement as part of a larger statement. 8

deterministic Pertaining to a program that does the same thing each time it runs, given the same inputs. 9

dot notation The syntax for calling a function in another module by specifying the module name followed by a dot (period) and the function name. 10

flow of execution The order in which statements are executed during a program run. 11

fruitful function A function that returns a value. 12

function A named sequence of statements that performs some useful operation. 13
Functions may or may not take arguments and may or may not produce a result.

function call A statement that executes a function. It consists of the function name followed by an argument list. 14

function definition A statement that creates a new function, specifying its name, parameters, and the statements it executes. 15

function object A value created by a function definition. The name of the function is a variable that refers to a function object. 16

header The first line of a function definition. 17

import statement A statement that reads a module file and creates a module object. 18

module object A value created by an `import` statement that provides access to the data and code defined in a module. 19

parameter A name used inside a function to refer to the value passed as an argument. 20

pseudorandom Pertaining to a sequence of numbers that appear to be random, but are generated by a deterministic program. 21

return value The result of a function. If a function call is used as an expression, the return value is the value of the expression. 22

void function A function that does not return a value. 23

4.14 Exercises ¹

Exercise 4: What is the purpose of the “def” keyword in Python? ²

- a) It is slang that means “the following code is really cool” 3
- b) It indicates the start of a function
- c) It indicates that the following indented section of code is to be stored for later
- d) b and c are both true
- e) None of the above

Exercise 5: What will the following Python program print out? ⁴

```
def fred(): 5
    print("Zap")
```

```
def jane(): 6
    print("ABC")
```

```
jane() 7
fred() 8
jane() 9
```

- a) Zap ABC jane fred jane 10
- b) Zap ABC Zap
- c) ABC Zap jane
- d) ABC Zap ABC
- e) Zap Zap Zap

Exercise 6: Rewrite your pay computation with time-and-a-half for overtime and 11 create a function called `computepay` which takes two parameters (`hours` and `rate`).

```
Enter Hours: 45 12
Enter Rate: 10 13
Pay: 475.0 14
```

Exercise 7: Rewrite the grade program from the previous chapter using a function 15 called `computegrade` that takes a score as its parameter and returns a grade as a string.

Score	Grade
≥ 0.9	A
≥ 0.8	B
≥ 0.7	C
≥ 0.6	D
< 0.6	F

16

```
Enter score: 0.95 17
A
```

```
Enter score: perfect 18
Bad score
```

Enter score: 10.0 **2**
Bad score

Enter score: 0.75 **3**
C

Enter score: 0.5 **4**
F

Run the program repeatedly to test the various different values for input. **5**

Chapter 5¹

Iteration²

5.1 Updating variables³

A common pattern in assignment statements is an assignment statement that updates a variable, where the new value of the variable depends on the old.⁴

```
x = x + 15
```

This means “get the current value of `x`, add 1, and then update `x` with the new value.”⁶

If you try to update a variable that doesn’t exist, you get an error, because Python evaluates the right side before it assigns a value to `x`:⁷

```
>>> x = x + 18
NameError: name 'x' is not defined9
```

Before you can update a variable, you have to *initialize* it, usually with a simple assignment:¹⁰

```
>>> x = 011
>>> x = x + 1
```

Updating a variable by adding 1 is called an *increment*; subtracting 1 is called a *decrement*.¹²

5.2 The `while` statement¹³

Computers are often used to automate repetitive tasks. Repeating identical or similar tasks without making errors is something that computers do well and people do poorly. Because iteration is so common, Python provides several language features to make it easier.¹⁴

One form of iteration in Python is the `while` statement. Here is a simple program that counts down from five and then says “Blastoff!”.¹⁵

```
n = 5          1
while n > 0:
    print(n)
    n = n - 1
print('Blastoff!')
```

You can almost read the `while` statement as if it were English. It means, “While 2 `n` is greater than 0, display the value of `n` and then reduce the value of `n` by 1. When you get to 0, exit the `while` statement and display the word `Blastoff!`”

More formally, here is the flow of execution for a `while` statement: 3

1. Evaluate the condition, yielding `True` or `False`. 4
2. If the condition is false, exit the `while` statement and continue execution at the next statement.
3. If the condition is true, execute the body and then go back to step 1.

This type of flow is called a *loop* because the third step loops back around to the top. We call each time we execute the body of the loop an *iteration*. For the above loop, we would say, “It had five iterations”, which means that the body of the loop was executed five times. 5

The body of the loop should change the value of one or more variables so that eventually the condition becomes false and the loop terminates. We call the variable that changes each time the loop executes and controls when the loop finishes the *iteration variable*. If there is no iteration variable, the loop will repeat forever, resulting in an *infinite loop*. 6

5.3 Infinite loops 7

An endless source of amusement for programmers is the observation that the directions on shampoo, “Lather, rinse, repeat,” are an infinite loop because there is no *iteration variable* telling you how many times to execute the loop. 8

In the case of `countdown`, we can prove that the loop terminates because we know that the value of `n` is finite, and we can see that the value of `n` gets smaller each time through the loop, so eventually we have to get to 0. Other times a loop is obviously infinite because it has no iteration variable at all. 9

Sometimes you don’t know it’s time to end a loop until you get half way through the body. In that case you can write an infinite loop on purpose and then use the `break` statement to jump out of the loop. 10

This loop is obviously an *infinite loop* because the logical expression on the `while` statement is simply the logical constant `True`: 11

```
n = 10          12
while True:
    print(n, end=' ')
    n = n - 1
print('Done!')
```

If you make the mistake and run this code, you will learn quickly how to stop a runaway Python process on your system or find where the power-off button is on your computer. This program will run forever or until your battery runs out because the logical expression at the top of the loop is always true by virtue of the fact that the expression is the constant value `True`. 1

While this is a dysfunctional infinite loop, we can still use this pattern to build useful loops as long as we carefully add code to the body of the loop to explicitly exit the loop using `break` when we have reached the exit condition. 2

For example, suppose you want to take input from the user until they type `done`. You could write: 3

```
while True:          4
    line = input('> ')
    if line == 'done':
        break
    print(line)
print('Done!')
```

Code: <https://www.py4e.com/code3/copytildone1.py> 5

The loop condition is `True`, which is always true, so the loop runs repeatedly until it hits the `break` statement. 6

Each time through, it prompts the user with an angle bracket. If the user types `done`, the `break` statement exits the loop. Otherwise the program echoes whatever the user types and goes back to the top of the loop. Here's a sample run: 7

```
> hello there 8
hello there
> finished
finished
> done
Done!
```

This way of writing `while` loops is common because you can check the condition anywhere in the loop (not just at the top) and you can express the stop condition affirmatively (“stop when this happens”) rather than negatively (“keep going until that happens.”). 9

5.4 Finishing iterations with `continue` 10

Sometimes you are in an iteration of a loop and want to finish the current iteration and immediately jump to the next iteration. In that case you can use the `continue` statement to skip to the next iteration without finishing the body of the loop for the current iteration. 11

Here is an example of a loop that copies its input until the user types “done”, but treats lines that start with the hash character as lines not to be printed (kind of like Python comments). 12

```
while True:
    line = input('> ')
    if line[0] == '#':
        continue
    if line == 'done':
        break
    print(line)
print('Done!')
```

1

Code: <https://www.py4e.com/code3/copytildone2.py> 2

Here is a sample run of this new program with `continue` added. 3

```
> hello there      4
hello there
> # don't print this
> print this!
print this!
> done
Done!
```

4

All the lines are printed except the one that starts with the hash sign because 5 when the `continue` is executed, it ends the current iteration and jumps back to the `while` statement to start the next iteration, thus skipping the `print` statement.

5.5 Definite loops using `for` 6

Sometimes we want to loop through a *set* of things such as a list of words, the lines 7 in a file, or a list of numbers. When we have a list of things to loop through, we can construct a *definite* loop using a `for` statement. We call the `while` statement an *indefinite* loop because it simply loops until some condition becomes `False`, whereas the `for` loop is looping through a known set of items so it runs through as many iterations as there are items in the set.

The syntax of a `for` loop is similar to the `while` loop in that there is a `for` 8 statement and a loop body:

```
friends = ['Joseph', 'Glenn', 'Sally'] 9
for friend in friends:
    print('Happy New Year:', friend)
print('Done!')
```

In Python terms, the variable `friends` is a list¹ 10 of three strings and the `for` loop goes through the list and executes the body once for each of the three strings in the list resulting in this output:

¹We will examine lists in more detail in a later chapter.

```
Happy New Year: Joseph 1
Happy New Year: Glenn
Happy New Year: Sally
Done!
```

Translating this `for` loop to English is not as direct as the `while`, but if you think 2 of friends as a *set*, it goes like this: “Run the statements in the body of the `for` loop once for each friend *in* the set named `friends`.”

Looking at the `for` loop, `for` and `in` are reserved Python keywords, and `friend`3 and `friends` are variables.

```
for friend in friends: 4
    print('Happy New Year:', friend)
```

In particular, `friend` is the *iteration variable* for the `for` loop. The variable `friend`5 changes for each iteration of the loop and controls when the `for` loop completes. The *iteration variable* steps successively through the three strings stored in the `friends` variable.

5.6 Loop patterns⁶

Often we use a `for` or `while` loop to go through a list of items or the contents of 7 a file and we are looking for something such as the largest or smallest value of the data we scan through.

These loops are generally constructed by:⁸

- Initializing one or more variables before the loop starts9
- Performing some computation on each item in the loop body, possibly changing the variables in the body of the loop
- Looking at the resulting variables when the loop completes

We will use a list of numbers to demonstrate the concepts and construction of these 10 loop patterns.

5.6.1 Counting and summing loops¹¹

For example, to count the number of items in a list, we would write the following 12 `for` loop:

```
count = 0 13
for itervar in [3, 41, 12, 9, 74, 15]:
    count = count + 1
print('Count: ', count)
```

We set the variable `count` to zero before the loop starts, then we write a `for` loop **1** to run through the list of numbers. Our *iteration* variable is named `itervar` and while we do not use `itervar` in the loop, it does control the loop and cause the loop body to be executed once for each of the values in the list.

In the body of the loop, we add 1 to the current value of `count` for each of the **2** values in the list. While the loop is executing, the value of `count` is the number of values we have seen “so far”.

Once the loop completes, the value of `count` is the total number of items. The **3** total number “falls in our lap” at the end of the loop. We construct the loop so that we have what we want when the loop finishes.

Another similar loop that computes the total of a set of numbers is as follows: **4**

```
total = 0 5
for itervar in [3, 41, 12, 9, 74, 15]:
    total = total + itervar
print('Total: ', total)
```

In this loop we *do* use the *iteration variable*. Instead of simply adding one to the **6** `count` as in the previous loop, we add the actual number (3, 41, 12, etc.) to the running total during each loop iteration. If you think about the variable `total`, it contains the “running total of the values so far”. So before the loop starts `total` is zero because we have not yet seen any values, during the loop `total` is the running total, and at the end of the loop `total` is the overall total of all the values in the list.

As the loop executes, `total` accumulates the sum of the elements; a variable used **7** this way is sometimes called an *accumulator*.

Neither the counting loop nor the summing loop are particularly useful in practice **8** because there are built-in functions `len()` and `sum()` that compute the number of items in a list and the total of the items in the list respectively.

5.6.2 Maximum and minimum loops **9**

To find the largest value in a list or sequence, we construct the following loop: **10**

```
largest = None 11
print('Before:', largest)
for itervar in [3, 41, 12, 9, 74, 15]:
    if largest is None or itervar > largest :
        largest = itervar
    print('Loop:', itervar, largest)
print('Largest:', largest)
```

When the program executes, the output is as follows: **12**

```
Before: None 13
Loop: 3 3
```

```
Loop: 41 41 1
Loop: 12 41 2
Loop: 9 41 3
Loop: 74 74 4
Loop: 15 74 5
Largest: 74 6
```

The variable `largest` is best thought of as the “largest value we have seen so far”.⁷ Before the loop, we set `largest` to the constant `None`. `None` is a special constant value which we can store in a variable to mark the variable as “empty”.

Before the loop starts, the largest value we have seen so far is `None` since we have⁸ not yet seen any values. While the loop is executing, if `largest` is `None` then we take the first value we see as the largest so far. You can see in the first iteration when the value of `itervar` is 3, since `largest` is `None`, we immediately set `largest` to be 3.

After the first iteration, `largest` is no longer `None`, so the second part of the⁹ compound logical expression that checks `itervar > largest` triggers only when we see a value that is larger than the “largest so far”. When we see a new “even larger” value we take that new value for `largest`. You can see in the program output that `largest` progresses from 3 to 41 to 74.

At the end of the loop, we have scanned all of the values and the variable `largest`¹⁰ now does contain the largest value in the list.

To compute the smallest number, the code is very similar with one small change:¹¹

```
smallest = None                                12
print('Before:', smallest)
for itervar in [3, 41, 12, 9, 74, 15]:
    if smallest is None or itervar < smallest:
        smallest = itervar
    print('Loop:', itervar, smallest)
print('Smallest:', smallest)
```

Again, `smallest` is the “smallest so far” before, during, and after the loop executes.¹³ When the loop has completed, `smallest` contains the minimum value in the list.

Again as in counting and summing, the built-in functions `max()` and `min()` make¹⁴ writing these exact loops unnecessary.

The following is a simple version of the Python built-in `min()` function:¹⁵

```
def min(values):
    smallest = None                                16
    for value in values:
        if smallest is None or value < smallest:
            smallest = value
    return smallest
```

In the function version of the smallest code, we removed all of the `print` statements¹⁷ so as to be equivalent to the `min` function which is already built in to Python.

5.7 Debugging¹

As you start writing bigger programs, you might find yourself spending more time ² debugging. More code means more chances to make an error and more places for bugs to hide.

One way to cut your debugging time is “debugging by bisection.” For example, if ³ there are 100 lines in your program and you check them one at a time, it would take 100 steps.

Instead, try to break the problem in half. Look at the middle of the program, ⁴ or near it, for an intermediate value you can check. Add a `print` statement (or something else that has a verifiable effect) and run the program.

If the mid-point check is incorrect, the problem must be in the first half of the ⁵ program. If it is correct, the problem is in the second half.

Every time you perform a check like this, you halve the number of lines you have ⁶ to search. After six steps (which is much less than 100), you would be down to one or two lines of code, at least in theory.

In practice it is not always clear what the “middle of the program” is and not ⁷ always possible to check it. It doesn’t make sense to count lines and find the exact midpoint. Instead, think about places in the program where there might be errors and places where it is easy to put a check. Then choose a spot where you think the chances are about the same that the bug is before or after the check.

5.8 Glossary⁸

accumulator A variable used in a loop to add up or accumulate a result.⁹

counter A variable used in a loop to count the number of times something hap-¹⁰ pened. We initialize a counter to zero and then increment the counter each time we want to “count” something.

decrement An update that decreases the value of a variable.¹¹

initialize An assignment that gives an initial value to a variable that will be ¹² updated.

increment An update that increases the value of a variable (often by one).¹³

infinite loop A loop in which the terminating condition is never satisfied or for ¹⁴ which there is no terminating condition.

iteration Repeated execution of a set of statements using either a function that ¹⁵ calls itself or a loop.

5.9 Exercises¹⁶

Exercise 1: Write a program which repeatedly reads integers until the user enters ¹⁷ “done”. Once “done” is entered, print out the total, count, and average of the integers. If the user enters anything other than an integer, detect their mistake using `try` and `except` and print an error message and skip to the next integers.

```
Enter a number: 4  
Enter a number: 5  
Enter a number: bad data  
Invalid input  
Enter a number: 7  
Enter a number: done  
16 3 5.333333333333333
```

1

Exercise 2: Write another program that prompts for a list of numbers as above **2** and at the end prints out both the maximum and minimum of the numbers instead of the average.

Chapter 6¹

Strings²

6.1 A string is a sequence³

A string is a *sequence* of characters. You can access the characters one at a time⁴ with the bracket operator:

```
>>> fruit = 'banana' 5  
>>> letter = fruit[1]
```

The second statement extracts the character at index position 1 from the `fruit`⁶ variable and assigns it to the `letter` variable.

The expression in brackets is called an *index*. The index indicates which character⁷ in the sequence you want (hence the name).

But you might not get what you expect:⁸

```
>>> print(letter) 9  
a
```

For most people, the first letter of “banana” is “b”, not “a”. But in Python, the index is an offset from the beginning of the string, and the offset of the first letter is zero.

```
>>> letter = fruit[0] 11  
>>> print(letter)  
b
```

So “b” is the 0th letter (“zero-th”) of “banana”, “a” is the 1th letter (“one-th”),¹² and “n” is the 2th (“two-th”) letter.

You can use any expression, including variables and operators, as an index, but¹³ the value of the index has to be an integer. Otherwise you get:

```
>>> letter = fruit[1.5] 14  
TypeError: string indices must be integers
```

b	a	n	a	n	a	1
[0]	[1]	[2]	[3]	[4]	[5]	

Figure 6.1: String Indexes 2

6.2 Getting the length of a string using `len` 3

`len` is a built-in function that returns the number of characters in a string: 4

```
>>> fruit = 'banana' 5
>>> len(fruit)
6
```

To get the last letter of a string, you might be tempted to try something like this: 6

```
>>> length = len(fruit) 7
>>> last = fruit[length]
IndexError: string index out of range
```

The reason for the `IndexError` is that there is no letter in “banana” with the index 8 6. Since we started counting at zero, the six letters are numbered 0 to 5. To get the last character, you have to subtract 1 from `length`:

```
>>> last = fruit[length-1] 9
>>> print(last)
a
```

Alternatively, you can use negative indices, which count backward from the end of 10 the string. The expression `fruit[-1]` yields the last letter, `fruit[-2]` yields the second to last, and so on.

6.3 Traversal through a string with a loop 11

A lot of computations involve processing a string one character at a time. Often 12 they start at the beginning, select each character in turn, do something to it, and continue until the end. This pattern of processing is called a *traversal*. One way to write a traversal is with a `while` loop:

```
index = 0 13
while index < len(fruit):
    letter = fruit[index]
    print(letter)
    index = index + 1
```

This loop traverses the string and displays each letter on a line by itself. The 1 loop condition is `index < len(fruit)`, so when `index` is equal to the length of the string, the condition is false, and the body of the loop is not executed. The last character accessed is the one with the index `len(fruit)-1`, which is the last character in the string.

Exercise 1: Write a `while` loop that starts at the last character in the string and 2 works its way backwards to the first character in the string, printing each letter on a separate line, except backwards.

Another way to write a traversal is with a `for` loop: 3

```
for char in fruit: 4
    print(char)
```

Each time through the loop, the next character in the string is assigned to the 5 variable `char`. The loop continues until no characters are left.

6.4 String slices 6

A segment of a string is called a *slice*. Selecting a slice is similar to selecting a 7 character:

```
>>> s = 'Monty Python' 8
>>> print(s[0:5])
Monty
>>> print(s[6:12])
Python
```

The operator `[n:m]` returns the part of the string from the “n-th” character to the 9 “m-th” character, including the first but excluding the last.

If you omit the first index (before the colon), the slice starts at the beginning of 10 the string. If you omit the second index, the slice goes to the end of the string:

```
>>> fruit = 'banana' 11
>>> fruit[:3]
'ban'
>>> fruit[3:]
'ana'
```

If the first index is greater than or equal to the second the result is an *empty string*, 12 represented by two quotation marks:

```
>>> fruit = 'banana' 13
>>> fruit[3:3]
''
```

An empty string contains no characters and has length 0, but other than that, it 14 is the same as any other string.

Exercise 2: Given that `fruit` is a string, what does `fruit[:]` mean? 15

6.5 Strings are immutable¹

It is tempting to use the operator on the left side of an assignment, with the intention of changing a character in a string. For example:

```
>>> greeting = 'Hello, world!'
>>> greeting[0] = 'J'  
TypeError: 'str' object does not support item assignment
```

The “object” in this case is the string and the “item” is the character you tried to assign. For now, an *object* is the same thing as a value, but we will refine that definition later. An *item* is one of the values in a sequence.

The reason for the error is that strings are *immutable*, which means you can’t change an existing string. The best you can do is create a new string that is a variation on the original:

```
>>> greeting = 'Hello, world!'  
>>> new_greeting = 'J' + greeting[1:]  
>>> print(new_greeting)  
Jello, world!
```

This example concatenates a new first letter onto a slice of `greeting`. It has no effect on the original string.

6.6 Looping and counting⁸

The following program counts the number of times the letter “a” appears in a string:

```
word = 'banana'  
count = 0  
for letter in word:  
    if letter == 'a':  
        count = count + 1  
print(count)
```

This program demonstrates another pattern of computation called a *counter*. The variable `count` is initialized to 0 and then incremented each time an “a” is found. When the loop exits, `count` contains the result: the total number of a’s.

Exercise 3: Encapsulate this code in a function named `count`, and generalize it so that it accepts the string and the letter as arguments.

6.7 The in operator¹

The word `in` is a boolean operator that takes two strings and returns `True` if the first appears as a substring in the second:²

```
>>> 'a' in 'banana'      3
True
>>> 'seed' in 'banana'
False
```

6.8 String comparison⁴

The comparison operators work on strings. To see if two strings are equal:⁵

```
if word == 'banana':          6
    print('All right, bananas.')
```

Other comparison operations are useful for putting words in alphabetical order:⁷

```
if word < 'banana':          8
    print('Your word, ' + word + ', comes before banana.')
elif word > 'banana':
    print('Your word, ' + word + ', comes after banana.')
else:
    print('All right, bananas.')
```

Python does not handle uppercase and lowercase letters the same way that people do. All the uppercase letters come before all the lowercase letters, so:⁹

```
Your word, Pineapple, comes before banana.10
```

A common way to address this problem is to convert strings to a standard format,¹¹ such as all lowercase, before performing the comparison. Keep that in mind in case you have to defend yourself against a man armed with a Pineapple.

6.9 String methods¹²

Strings are an example of Python *objects*. An object contains both data (the actual string itself) and *methods*, which are effectively functions that are built into the object and are available to any *instance* of the object.¹³

Python has a function called `dir` which lists the methods available for an object.¹⁴ The `type` function shows the type of an object and the `dir` function shows the available methods.

```
>>> stuff = 'Hello world'  
>>> type(stuff)  
<class 'str'>  
>>> dir(stuff)  
[... 'capitalize', 'casefold', 'center', 'count', 'encode',  
'endswith', 'expandtabs', 'find', 'format', 'format_map',  
'index', 'isalnum', 'isalpha', 'isdecimal', 'isdigit',  
'isidentifier', 'islower', 'isnumeric', 'isprintable',  
'isspace', 'istitle', 'isupper', 'join', 'ljust', 'lower',  
'lstrip', 'maketrans', 'partition', 'replace', 'rfind',  
'rindex', 'rjust', 'rpartition', 'rsplit', 'rstrip',  
'split', 'splitlines', 'startswith', 'strip', 'swapcase',  
'title', 'translate', 'upper', 'zfill']  
>>> help(str.capitalize)  
Help on method_descriptor:  
  
capitalize(self, /)  
    Return a capitalized version of the string.  
  
    More specifically, make the first character have upper  
    case and the rest lower case.  
>>>
```

While the `dir` function lists the methods, and you can use `help` to get some simple **2** documentation on a method, a better source of documentation for string methods would be

[**3**](https://docs.python.org/library/stdtypes.html#string-methods)

Calling a *method* is similar to calling a function (it takes arguments and returns **4** a value) but the syntax is different. We call a method by appending the method name to the variable name using the period as a delimiter.

For example, the method `upper` takes a string and returns a new string with all **5** uppercase letters:

Instead of the function syntax `upper(word)`, it uses the method syntax **6** `word.upper()`.

```
>>> word = 'banana' 7  
>>> new_word = word.upper()  
>>> print(new_word)  
BANANA
```

This form of dot notation specifies the name of the method, `upper`, and the name **8** of the string to apply the method to, `word`. The empty parentheses indicate that this method takes no argument.

A method call is called an *invocation*; in this case, we would say that we are **9** invoking `upper` on the `word`.

For example, there is a string method named `find` that searches for the position **10** of one string within another:

```
>>> word = 'banana' 1
>>> index = word.find('a')
>>> print(index)
1
```

In this example, we invoke `find` on `word` and pass the letter we are looking for as 2 a parameter.

The `find` method can find substrings as well as characters: 3

```
>>> word.find('na') 4
2
```

It can take as a second argument the index where it should start: 5

```
>>> word.find('na', 3) 6
4
```

One common task is to remove white space (spaces, tabs, or newlines) from the 7 beginning and end of a string using the `strip` method:

```
>>> line = ' Here we go ' 8
>>> line.strip()
'Here we go'
```

Some methods such as `startswith` return boolean values. 9

```
>>> line = 'Have a nice day' 10
>>> line.startswith('Have')
True
>>> line.startswith('h')
False
```

You will note that `startswith` requires case to match, so sometimes we take a line 11 and map it all to lowercase before we do any checking using the `lower` method.

```
>>> line = 'Have a nice day' 12
>>> line.startswith('h')
False
>>> line.lower()
'have a nice day'
>>> line.lower().startswith('h')
True
```

In the last example, the method `lower` is called and then we use `startswith` to 13 see if the resulting lowercase string starts with the letter "h". As long as we are careful with the order, we can make multiple method calls in a single expression.

Exercise 4: There is a string method called `count` that is similar to the function 14 in the previous exercise. Read the documentation of this method at:

<https://docs.python.org/library/stdtypes.html#string-methods> 15

Write an invocation that counts the number of times the letter a occurs in "banana". 16

6.10 Parsing strings¹

Often, we want to look into a string and find a substring. For example if we were ² presented a series of lines formatted as follows:

```
From stephen.marquard@uct.ac.za Sat Jan 5 09:14:16 20083
```

and we wanted to pull out only the second half of the address (i.e., uct.ac.za)⁴ from each line, we can do this by using the `find` method and string slicing.

First, we will find the position of the at-sign in the string. Then we will find the ⁵ position of the first space *after* the at-sign. And then we will use string slicing to extract the portion of the string which we are looking for.

```
>>> data = 'From stephen.marquard@uct.ac.za Sat Jan 5 09:14:16 20086
>>> atpos = data.find('@')
>>> print(atpos)
21
>>> spos = data.find(' ',atpos)
>>> print(spos)
31
>>> host = data[atpos+1:spos]
>>> print(host)
uct.ac.za
>>>
```

We use a version of the `find` method which allows us to specify a position in ⁷ the string where we want `find` to start looking. When we slice, we extract the characters from “one beyond the at-sign through up to *but not including* the space character”.

The documentation for the `find` method is available at ⁸

[⁹](https://docs.python.org/library/stdtypes.html#string-methods)

6.11 Formatted String Literals¹⁰

A formatted string literal (often referred to simply as an f-string) allows Python ¹¹ expressions to be used within string literals. This is accomplished by prepending an `f` to the string literal and enclosing expressions in curly braces `{}`.

For example, wrapping a variable name in curly braces inside an f-string will cause ¹² it to be replaced by its value:

```
>>> camels = 4213
>>> f'{camels}'
'42'
```

The result is the string ‘42’, which is not to be confused with the integer value 42.¹⁴

An expression can appear anywhere in the string, so you can embed a value in a ¹⁵ sentence:

```
>>> camels = 42
>>> f'I have spotted {camels} camels.'
'I have spotted 42 camels.'
```

Several expressions can be included within a single string literal in order to create more complex strings.²

```
>>> years = 3
>>> count = .1
>>> species = 'camels'
>>> f'In {years} years I have spotted {count} {species}.'
'In 3 years I have spotted 0.1 camels.'
```

Formatted string literals are powerful, and they can do even more than is covered here. You can read more about them at

<https://docs.python.org/3/tutorial/inputoutput.html#formatted-string-literals>⁵

6.12 Debugging⁶

A skill that you should cultivate as you program is always asking yourself, “What could go wrong here?” or alternatively, “What crazy thing might our user do to crash our (seemingly) perfect program?”⁷

For example, look at the program which we used to demonstrate the `while` loop in the chapter on iteration:⁸

```
while True:          9
    line = input('> ')
    if line[0] == '#':
        continue
    if line == 'done':
        break
    print(line)
print('Done!')
```

Code: <https://www.py4e.com/code3/copytildone2.py>¹⁰

Look what happens when the user enters an empty line of input:¹¹

```
> hello there          12
hello there
> # don't print this
> print this!
print this!
>
Traceback (most recent call last):
  File "copytildone.py", line 3, in <module>
    if line[0] == '#':
IndexError: string index out of range
```

The code works fine until it is presented an empty line. Then there is no zero-th character, so we get a traceback. There are two solutions to this to make line three “safe” even if the line is empty.¹

One possibility is to simply use the `startswith` method which returns `False` if the string is empty.²

```
if line.startswith('#'):3
```

Another way is to safely write the `if` statement using the *guardian* pattern and make sure the second logical expression is evaluated only where there is at least one character in the string:⁴

```
if len(line) > 0 and line[0] == '#':5
```

6.13 Glossary⁶

counter A variable used to count something, usually initialized to zero and then incremented.⁷

empty string A string with no characters and length 0, represented by two quotation marks.⁸

flag A boolean variable used to indicate whether a condition is true or false.⁹

invocation A statement that calls a method.¹⁰

immutable The property of a sequence whose items cannot be assigned.¹¹

index An integer value used to select an item in a sequence, such as a character in a string.¹²

item One of the values in a sequence.¹³

method A function that is associated with an object and called using dot notation.¹⁴

object Something a variable can refer to. For now, you can use “object” and “value” interchangeably.¹⁵

search A pattern of traversal that stops when it finds what it is looking for.¹⁶

sequence An ordered set; that is, a set of values where each value is identified by an integer index.¹⁷

slice A part of a string specified by a range of indices.¹⁸

traverse To iterate through the items in a sequence, performing a similar operation on each.¹⁹

6.14 Exercises²⁰

Exercise 5: Slicing strings²¹

Take the following Python code that stores a string:²²

```
str = 'X-DSPAM-Confidence: 0.8475'23
```

Use `find` and string slicing to extract the portion of the string after the colon character and then use the `float` function to convert the extracted string into a floating point number.²⁴

Exercise 6: String methods 1

Read the documentation of the string methods at 2

<https://docs.python.org/library/stdtypes.html#string-methods>. 3

You might want to experiment with some of them to make sure you understand 4 how they work. `strip` and `replace` are particularly useful.

The documentation uses a syntax that might be confusing. For example, in 5 `find(sub[, start[, end]])`, the brackets indicate optional arguments. So `sub` is required, but `start` is optional, and if you include `start`, then `end` is optional.

Chapter 7¹

Files²

7.1 Persistence³

So far, we have learned how to write programs and communicate our intentions to ⁴ the *Central Processing Unit* using conditional execution, functions, and iterations. We have learned how to create and use data structures in the *Main Memory*. The CPU and memory are where our software works and runs. It is where all of the “thinking” happens.

But if you recall from our hardware architecture discussions, once the power is ⁵ turned off, anything stored in either the CPU or main memory is erased. So up to now, our programs have just been transient fun exercises to learn Python.

In this chapter, we start to work with *Secondary Memory* (or files). Secondary ⁶ memory is not erased when the power is turned off. Or in the case of a USB flash drive, the data we write from our programs can be removed from the system and transported to another system.

We will primarily focus on reading and writing text files such as those we create in ⁷ a text editor. Later we will see how to work with database files which are binary files, specifically designed to be read and written through database software.

7.2 Opening files⁸

When we want to read or write a file (say on your hard drive), we first must ⁹ open the file. Opening the file communicates with your operating system, which knows where the data for each file is stored. When you open a file, you are asking the operating system to find the file by name and make sure the file exists. In this example, we open the file *mbox.txt*, which should be stored in the same folder that you are in when you start Python. You can download this file from www.py4e.com/code3/mbox.txt

```
>>> fhand = open('mbox.txt')  
>>> print(fhand)  
<_io.TextIOWrapper name='mbox.txt' mode='r' encoding='cp1252'>
```




Figure 7.1: Secondary Memory 2




Figure 7.2: A File Handle 4

If the `open` is successful, the operating system returns us a *file handle*. The file handle is not the actual data contained in the file, but instead it is a “handle” that we can use to read the data. You are given a handle if the requested file exists and you have the proper permissions to read the file.

If the file does not exist, `open` will fail with a traceback and you will not get a handle to access the contents of the file:

```
>>> fhand = open('stuff.txt') 7
Traceback (most recent call last):
File "<stdin>", line 1, in <module>
FileNotFoundError: [Errno 2] No such file or directory: 'stuff.txt'
```

Later we will use `try` and `except` to deal more gracefully with the situation where we attempt to open a file that does not exist.

7.3 Text files and lines¹

A text file can be thought of as a sequence of lines, much like a Python string can² be thought of as a sequence of characters. For example, this is a sample of a text file which records mail activity from various individuals in an open source project development team:

```
From stephen.marquard@uct.ac.za Sat Jan  5 09:14:16 2008      3
Return-Path: <postmaster@collab.sakaiproject.org>
Date: Sat, 5 Jan 2008 09:12:18 -0500
To: source@collab.sakaiproject.org
From: stephen.marquard@uct.ac.za
Subject: [sakai] svn commit: r39772 - content/branches/
Details: http://source.sakaiproject.org/viewsvn/?view=rev&rev=39772
...
...
```

The entire file of mail interactions is available from⁴

www.py4e.com/code3/mbox.txt⁵

and a shortened version of the file is available from⁶

www.py4e.com/code3/mbox-short.txt⁷

These files are in a standard format for a file containing multiple mail messages.⁸ The lines which start with “From” separate the messages and the lines which start with “From:” are part of the messages. For more information about the mbox format, see <https://en.wikipedia.org/wiki/Mbox>.

To break the file into lines, there is a special character that represents the “end of⁹ the line” called the *newline* character.

In Python, we represent the *newline* character as a backslash-n in string constants.¹⁰ Even though this looks like two characters, it is actually a single character. When we look at the variable by entering “stuff” in the interpreter, it shows us the \n in the string, but when we use `print` to show the string, we see the string broken into two lines by the newline character.

```
>>> stuff = 'Hello\nWorld!' 11
>>> stuff
'Hello\nWorld!'
>>> print(stuff)
Hello
World!
>>> stuff = 'X\nY'
>>> print(stuff)
X
Y
>>> len(stuff)
3
```

You can also see that the length of the string X\nY is *three* characters because the¹² newline character is a single character.

So when we look at the lines in a file, we need to *imagine* that there is a special invisible character called the newline at the end of each line that marks the end of the line. 1

So the newline character separates the characters in the file into lines. 2

7.4 Reading files 3

While the *file handle* does not contain the data for the file, it is quite easy to 4 construct a **for** loop to read through and count each of the lines in a file:

```
fhand = open('mbox-short.txt') 5
count = 0
for line in fhand:
    count = count + 1
print('Line Count:', count)
```

Code: <https://www.py4e.com/code3/open.py> 6

We can use the file handle as the sequence in our **for** loop. Our **for** loop simply 7 counts the number of lines in the file and prints them out. The rough translation of the **for** loop into English is, “for each line in the file represented by the file handle, add one to the **count** variable.”

The reason that the **open** function does not read the entire file is that the file might 8 be quite large with many gigabytes of data. The **open** statement takes the same amount of time regardless of the size of the file. The **for** loop actually causes the data to be read from the file.

When the file is read using a **for** loop in this manner, Python takes care of splitting 9 the data in the file into separate lines using the newline character. Python reads each line through the newline and includes the newline as the last character in the **line** variable for each iteration of the **for** loop.

Because the **for** loop reads the data one line at a time, it can efficiently read and 10 count the lines in very large files without running out of main memory to store the data. The above program can count the lines in any size file using very little memory since each line is read, counted, and then discarded.

If you know the file is relatively small compared to the size of your main memory, 11 you can read the whole file into one string using the **read** method on the file handle.

```
>>> fhand = open('mbox-short.txt') 12
>>> inp = fhand.read()
>>> print(len(inp))
94626
>>> print(inp[:20])
From stephen.marquar
```

In this example, the entire contents (all 94,626 characters) of the file *mbox-short.txt* 13 are read directly into the variable **inp**. We use string slicing to print out the first 20 characters of the string data stored in **inp**.

When the file is read in this manner, all the characters including all of the lines and newline characters are one big string in the variable `inp`. It is a good idea to store the output of `read` as a variable because each call to `read` exhausts the resource: 1

```
>>> fhand = open('mbox-short.txt') 2
>>> print(len(fhand.read()))
94626
>>> print(len(fhand.read()))
0
```

Remember that this form of the `open` function should only be used if the file data will fit comfortably in the main memory of your computer. If the file is too large to fit in main memory, you should write your program to read the file in chunks using a `for` or `while` loop. 3

7.5 Searching through a file⁴

When you are searching through data in a file, it is a very common pattern to read through a file, ignoring most of the lines and only processing lines which meet a particular condition. We can combine the pattern for reading a file with string methods to build simple search mechanisms. 5

For example, if we wanted to read a file and only print out lines which started with the prefix “From:”, we could use the string method `startswith` to select only those lines with the desired prefix: 6

```
fhand = open('mbox-short.txt') 7
for line in fhand:
    if line.startswith('From:'):
        print(line)

# Code: https://www.py4e.com/code3/search1.py 8
```

When this program runs, we get the following output: 9

```
From: stephen.marquard@uct.ac.za 10
From: louis@media.berkeley.edu 11
From: zqian@umich.edu 12
From: rjlowe@iupui.edu 13
...
```

The output looks great since the only lines we are seeing are those which start with “From:”, but why are we seeing the extra blank lines? This is due to that invisible `newline` character. Each of the lines ends with a newline, so the `print` statement 14

prints the string in the variable `line` which includes a newline and then `print` adds 1
another newline, resulting in the double spacing effect we see.

We could use line slicing to print all but the last character, but a simpler approach 2
is to use the `rstrip` method which strips whitespaces from the right side of a string
as follows:

```
fhand = open('mbox-short.txt') 3
for line in fhand:
    line = line.rstrip()
    if line.startswith('From:'):
        print(line)
```

Code: <https://www.py4e.com/code3/search2.py> 4

When this program runs, we get the following output: 5

```
From: stephen.marquard@uct.ac.za 6
From: louis@media.berkeley.edu
From: zqian@umich.edu
From: rjlowe@iupui.edu
From: zqian@umich.edu
From: rjlowe@iupui.edu
From: cwen@iupui.edu
...
...
```

As your file processing programs get more complicated, you may want to structure 7
your search loops using `continue`. The basic idea of the search loop is that you are
looking for “interesting” lines and effectively skipping “uninteresting” lines. And
then when we find an interesting line, we do something with that line.

We can structure the loop to follow the pattern of skipping uninteresting lines as 8
follows:

```
fhand = open('mbox-short.txt') 9
for line in fhand:
    line = line.rstrip()
    # Skip 'uninteresting' lines
    if not line.startswith('From:'):
        continue
    # Process our 'interesting' line
    print(line)
```

Code: <https://www.py4e.com/code3/search3.py> 10

The output of the program is the same. In English, the uninteresting lines are 11
those which do not start with “From:”, which we skip using `continue`. For the
“interesting” lines (i.e., those that start with “From:”) we perform the processing.

We can use the `find` string method to simulate a text editor search that finds lines 12
where the search string is anywhere in the line. Since `find` looks for an occurrence

of a string within another string and either returns the position of the string or -1 if the string was not found, we can write the following loop to show lines which contain the string “@uct.ac.za” (i.e., they come from the University of Cape Town in South Africa):

```
fhand = open('mbox-short.txt')  
for line in fhand:  
    line = line.rstrip()  
    if line.find('@uct.ac.za') == -1: continue  
    print(line)
```

Code: <https://www.py4e.com/code3/search4.py> 3

Which produces the following output:

```
From stephen.marquard@uct.ac.za Sat Jan  5 09:14:16 2008  
X-Authentication-Warning: set sender to stephen.marquard@uct.ac.za using -f  
From: stephen.marquard@uct.ac.za  
Author: stephen.marquard@uct.ac.za  
From david.horwitz@uct.ac.za Fri Jan  4 07:02:32 2008  
X-Authentication-Warning: set sender to david.horwitz@uct.ac.za using -f  
From: david.horwitz@uct.ac.za  
Author: david.horwitz@uct.ac.za  
...
```

Here we also use the contracted form of the `if` statement where we put the `continue` on the same line as the `if`. This contracted form of the `if` functions the same as if the `continue` were on the next line and indented.

7.6 Letting the user choose the file name

We really do not want to have to edit our Python code every time we want to process a different file. It would be more usable to ask the user to enter the file name string each time the program runs so they can use our program on different files without changing the Python code.

This is quite simple to do by reading the file name from the user using `input` as follows:

```
fname = input('Enter the file name: ')  
fhand = open(fname)  
count = 0  
for line in fhand:  
    if line.startswith('Subject:'): 10  
        count = count + 1  
print('There were', count, 'subject lines in', fname)
```

Code: <https://www.py4e.com/code3/search6.py> 11

We read the file name from the user and place it in a variable named `fname` and **1** open that file. Now we can run the program repeatedly on different files.

```
python search6.py 2
Enter the file name: mbox.txt 3
There were 1797 subject lines in mbox.txt 4

python search6.py 5
Enter the file name: mbox-short.txt 6
There were 27 subject lines in mbox-short.txt 7
```

Before peeking at the next section, take a look at the above program and ask **8** yourself, “What could go possibly wrong here?” or “What might our friendly user do that would cause our nice little program to ungracefully exit with a traceback, making us look not-so-cool in the eyes of our users?”

7.7 Using `try`, `except`, and `open`**9**

I told you not to peek. This is your last chance.**10**

What if our user types something that is not a file name?**11**

```
python search6.py 12
Enter the file name: missing.txt 13
Traceback (most recent call last):14
  File "search6.py", line 2, in <module> 15
    fhndl = open(fname)
FileNotFoundError: [Errno 2] No such file or directory: 'missing.txt' 16

python search6.py 17
Enter the file name: na na boo boo 18
Traceback (most recent call last):19
  File "search6.py", line 2, in <module> 20
    fhndl = open(fname)
FileNotFoundError: [Errno 2] No such file or directory: 'na na boo boo' 21
```

Do not laugh. Users will eventually do every possible thing they can do to break **22** your programs, either mistakenly or with malicious intent. As a matter of fact, an important part of any software development team is a person or group called *Quality Assurance* (or QA for short) whose very job it is to do the craziest things possible in an attempt to break the software that the programmer has created.

The QA team is responsible for finding the flaws in programs before we have **23** delivered the program to the end users who may be purchasing the software or paying our salary to write the software. So the QA team is the programmer’s best friend.

So now that we see the flaw in the program, we can elegantly fix it using the **24** `try/except` structure. We need to assume that the `open` call might fail and add recovery code when the `open` fails as follows:

```

fname = input('Enter the file name: ') 1
try:
    fhand = open(fname)
except:
    print('File cannot be opened:', fname)
    exit()
count = 0
for line in fhand:
    if line.startswith('Subject:'):
        count = count + 1
print('There were', count, 'subject lines in', fname)

# Code: https://www.py4e.com/code3/search7.py 2

```

The `exit` function terminates the program. It is a function that we call that never returns. Now when our user (or QA team) types in silliness or bad file names, we “catch” them and recover gracefully:

```

python search7.py 4
Enter the file name: mbox.txt
There were 1797 subject lines in mbox.txt

```

```

python search7.py 5
Enter the file name: na na boo boo
File cannot be opened: na na boo boo

```

Protecting the `open` call is a good example of the proper use of `try` and `except` 6 in a Python program. We use the term “Pythonic” when we are doing something the “Python way”. We might say that the above example is the Pythonic way to open a file.

Once you become more skilled in Python, you can engage in repartee with other Python programmers to decide which of two equivalent solutions to a problem is “more Pythonic”. The goal to be “more Pythonic” captures the notion that programming is part engineering and part art. We are not always interested in just making something work, we also want our solution to be elegant and to be appreciated as elegant by our peers.

7.8 Writing files⁸

To write a file, you have to open it with mode “w” as a second parameter:

```

>>> fout = open('output.txt', 'w') 10
>>> print(fout)
<_io.TextIOWrapper name='output.txt' mode='w' encoding='cp1252'>

```

If the file already exists, opening it in write mode clears out the old data and starts fresh, so be careful! If the file doesn’t exist, a new one is created.

The `write` method of the file handle object puts data into the file, returning the number of characters written. The default write mode is text for writing (and reading) strings. 1

```
>>> line1 = "This here's the wattle,\n" 2  
>>> fout.write(line1)  
24
```

Again, the file object keeps track of where it is, so if you call `write` again, it adds the new data to the end. 3

We must make sure to manage the ends of lines as we write to the file by explicitly inserting the newline character when we want to end a line. The `print` statement automatically appends a newline, but the `write` method does not add the newline automatically. 4

```
>>> line2 = 'the emblem of our land.\n' 5  
>>> fout.write(line2)  
24
```

When you are done writing, you have to close the file to make sure that the last bit of data is physically written to the disk so it will not be lost if the power goes off. 6

```
>>> fout.close() 7
```

We could close the files which we open for reading as well, but we can be a little sloppy if we are only opening a few files since Python makes sure that all open files are closed when the program ends. When we are writing files, we want to explicitly close the files so as to leave nothing to chance. 8

7.9 Debugging 9

When you are reading and writing files, you might run into problems with whitespace. These errors can be hard to debug because spaces, tabs, and newlines are normally invisible: 10

```
>>> s = '1 2\t 3\n 4' 11  
>>> print(s)  
1 2 3  
4
```

The built-in function `repr` can help. It takes any object as an argument and returns a string representation of the object. For strings, it represents whitespace characters with backslash sequences: 12

```
>>> print(repr(s)) 13  
'1 2\t 3\n 4'
```

This can be helpful for debugging.¹

One other problem you might run into is that different systems use different characters to indicate the end of a line. Some systems use a newline, represented \n. Others use a return character, represented \r. Some use both. If you move files between different systems, these inconsistencies might cause problems.²

For most systems, there are applications to convert from one format to another. You can find them (and read more about this issue) at <https://www.wikipedia.org/wiki/Newline>. Or, of course, you could write one yourself.³

7.10 Glossary⁴

catch To prevent an exception from terminating a program using the try and except statements.⁵

newline A special character used in files and strings to indicate the end of a line.⁶

Pythonic A technique that works elegantly in Python. “Using try and except is the *Pythonic* way to recover from missing files”.⁷

Quality Assurance A person or team focused on ensuring the overall quality of a software product. QA is often involved in testing a product and identifying problems before the product is released.⁸

text file A sequence of characters stored in permanent storage like a hard drive.⁹

7.11 Exercises¹⁰

Exercise 1: Write a program to read through a file and print the contents of the file (line by line) all in upper case. Executing the program will look as follows:¹¹

```
python shout.py 12
Enter a file name: mbox-short.txt 13
FROM STEPHEN.MARQUARD@UCT.AC.ZA SAT JAN 5 09:14:16 2008 14
RETURN-PATH: <POSTMASTER@COLLAB.SAKAIPROJECT.ORG> 15
RECEIVED: FROM MURDER (MAIL.UMICH.EDU [141.211.14.90]) 16
      BY FRANKENSTEIN.MAIL.UMICH.EDU (CYRUS V2.3.8) WITH LMTPA; 17
      SAT, 05 JAN 2008 09:14:16 -0500 18
```

You can download the file from www.py4e.com/code3/mbox-short.txt¹⁹

Exercise 2: Write a program to prompt for a file name, and then read through the file and look for lines of the form:²⁰

```
X-DSPAM-Confidence: 0.8475 21
```

When you encounter a line that starts with “X-DSPAM-Confidence:” pull apart the line to extract the floating-point number on the line. Count these lines and then compute the total of the spam confidence values from these lines. When you reach the end of the file, print out the average spam confidence.²²

```
Enter the file name: mbox.txt      1  
Average spam confidence: 0.894128046745
```

```
Enter the file name: mbox-short.txt 2  
Average spam confidence: 0.750718518519
```

Test your file on the *mbox.txt* and *mbox-short.txt* files. 3

Exercise 3: 4

Sometimes when programmers get bored or want to have a bit of fun, they add a **5** harmless *Easter Egg* to their program. Modify the program that prompts the user for the file name so that it prints a funny message when the user types in the exact file name “na na boo boo”. The program should behave normally for all other files which exist and don’t exist. Here is a sample execution of the program:

```
python egg.py      6  
Enter the file name: mbox.txt  
There were 1797 subject lines in mbox.txt
```

```
python egg.py      7  
Enter the file name: missing.tyxt  
File cannot be opened: missing.tyxt
```

```
python egg.py      8  
Enter the file name: na na boo boo  
NA NA BOO BOO TO YOU - You have been punk'd!
```

We are not encouraging you to put Easter Eggs in your programs; this is just an **9** exercise.

Chapter 8¹

Lists²

8.1 A list is a sequence³

Like a string, a *list* is a sequence of values. In a string, the values are characters;⁴ in a list, they can be any type. The values in lists are called *elements* or sometimes *items*.

There are several ways to create a new list; the simplest is to enclose the elements⁵ in square brackets ("[" and "]"):

```
[10, 20, 30, 40]6  
['crunchy frog', 'ram bladder', 'lark vomit']7
```

The first example is a list of four integers. The second is a list of three strings.⁸ The elements of a list don't have to be the same type. The following list contains a string, a float, an integer, and (lo!) another list:

```
['spam', 2.0, 5, [10, 20]]9
```

A list within another list is *nested*.¹⁰

A list that contains no elements is called an empty list; you can create one with¹¹ empty brackets, [].

As you might expect, you can assign list values to variables:¹²

```
>>> cheeses = ['Cheddar', 'Edam', 'Gouda']13  
>>> numbers = [17, 123]  
>>> empty = []  
>>> print(cheeses, numbers, empty)  
['Cheddar', 'Edam', 'Gouda'] [17, 123] []
```

8.2 Lists are mutable¹

The syntax for accessing the elements of a list is the same as for accessing the characters of a string: the bracket operator. The expression inside the brackets specifies the index. Remember that the indices start at 0:²

```
>>> print(cheeses[0])3  
Cheddar4
```

Unlike strings, lists are mutable because you can change the order of items in a list or reassign an item in a list. When the bracket operator appears on the left side of an assignment, it identifies the element of the list that will be assigned.⁵

```
>>> numbers = [17, 123]6  
>>> numbers[1] = 57  
>>> print(numbers)8  
[17, 5]9
```

The one-th element of `numbers`, which used to be 123, is now 5.¹⁰

You can think of a list as a relationship between indices and elements. This relationship is called a *mapping*; each index “maps to” one of the elements.¹¹

List indices work the same way as string indices:¹²

- Any integer expression can be used as an index.¹³
- If you try to read or write an element that does not exist, you get an `IndexError`.
- If an index has a negative value, it counts backward from the end of the list.

The `in` operator also works on lists.¹⁴

```
>>> cheeses = ['Cheddar', 'Edam', 'Gouda']15  
>>> 'Edam' in cheeses16  
True17  
>>> 'Brie' in cheeses18  
False19
```

8.3 Traversing a list²⁰

The most common way to traverse the elements of a list is with a `for` loop. The syntax is the same as for strings:²¹

```
for cheese in cheeses:22  
    print(cheese)23
```

This works well if you only need to read the elements of the list. But if you want to write or update the elements, you need the indices. A common way to do that is to combine the functions `range` and `len`: 1

```
for i in range(len(numbers)): 2
    numbers[i] = numbers[i] * 2
```

This loop traverses the list and updates each element. `len` returns the number of elements in the list. `range` returns a list of indices from 0 to $n - 1$, where n is the length of the list. Each time through the loop, `i` gets the index of the next element. The assignment statement in the body uses `i` to read the old value of the element and to assign the new value. 3

A `for` loop over an empty list never executes the body: 4

```
for x in empty: 5
    print('This never happens.')
```

Although a list can contain another list, the nested list still counts as a single element. The length of this list is four: 6

```
['spam', 1, ['Brie', 'Roquefort', 'Pol le Veq'], [1, 2, 3]] 7
```

8.4 List operations 8

The `+` operator concatenates lists: 9

```
>>> a = [1, 2, 3] 10
>>> b = [4, 5, 6]
>>> c = a + b
>>> print(c)
[1, 2, 3, 4, 5, 6]
```

Similarly, the `*` operator repeats a list a given number of times: 11

```
>>> [0] * 4 12
[0, 0, 0, 0]
>>> [1, 2, 3] * 3
[1, 2, 3, 1, 2, 3, 1, 2, 3]
```

The first example repeats four times. The second example repeats the list three times. 13

8.5 List slices¹

The slice operator also works on lists:²

```
>>> t = ['a', 'b', 'c', 'd', 'e', 'f']3
>>> t[1:3]
['b', 'c']
>>> t[:4]
['a', 'b', 'c', 'd']
>>> t[3:]
['d', 'e', 'f']
```

If you omit the first index, the slice starts at the beginning. If you omit the second,⁴ the slice goes to the end. So if you omit both, the slice is a copy of the whole list.

```
>>> t[:]5
['a', 'b', 'c', 'd', 'e', 'f']
```

Since lists are mutable, it is often useful to make a copy before performing operations that fold, spindle, or mutilate lists.⁶

A slice operator on the left side of an assignment can update multiple elements:⁷

```
>>> t = ['a', 'b', 'c', 'd', 'e', 'f']8
>>> t[1:3] = ['x', 'y']
>>> print(t)
['a', 'x', 'y', 'd', 'e', 'f']
```

8.6 List methods⁹

Python provides methods that operate on lists. For example, `append` adds a new¹⁰ element to the end of a list:

```
>>> t = ['a', 'b', 'c']11
>>> t.append('d')
>>> print(t)
['a', 'b', 'c', 'd']
```

`extend` takes a list as an argument and appends all of the elements:¹²

```
>>> t1 = ['a', 'b', 'c']13
>>> t2 = ['d', 'e']
>>> t1.extend(t2)
>>> print(t1)
['a', 'b', 'c', 'd', 'e']
```

This example leaves `t2` unmodified.¹⁴

`sort` arranges the elements of the list from low to high:¹⁵

```
>>> t = ['d', 'c', 'e', 'b', 'a']1
>>> t.sort()
>>> print(t)
['a', 'b', 'c', 'd', 'e']
```

Most list methods are void; they modify the list and return `None`. If you accidentally write `t = t.sort()`, you will be disappointed with the result.²

8.7 Deleting elements³

There are several ways to delete elements from a list. If you know the index of the element you want, you can use `pop`:⁴

```
>>> t = ['a', 'b', 'c']5
>>> x = t.pop(1)
>>> print(t)
['a', 'c']
>>> print(x)
b
```

`pop` modifies the list and returns the element that was removed. If you don't provide an index, it deletes and returns the last element.⁶

If you don't need the removed value, you can use the `del` statement:⁷

```
>>> t = ['a', 'b', 'c']8
>>> del t[1]
>>> print(t)
['a', 'c']
```

If you know the element you want to remove (but not the index), you can use `remove`:⁹

```
>>> t = ['a', 'b', 'c']10
>>> t.remove('b')
>>> print(t)
['a', 'c']
```

The return value from `remove` is `None`.¹¹

To remove more than one element, you can use `del` with a slice index:¹²

```
>>> t = ['a', 'b', 'c', 'd', 'e', 'f']13
>>> del t[1:5]
>>> print(t)
['a', 'f']
```

As usual, the slice selects all the elements up to, but not including, the second index.¹⁴

8.8 Lists and functions¹

There are a number of built-in functions that can be used on lists that allow you² to quickly look through a list without writing your own loops:

```
>>> nums = [3, 41, 12, 9, 74, 15]3
>>> print(len(nums))
6
>>> print(max(nums))
74
>>> print(min(nums))
3
>>> print(sum(nums))
154
>>> print(sum(nums)/len(nums))
25
```

The `sum()` function only works when the list elements are numbers. The other⁴ functions (`max()`, `len()`, etc.) work with lists of strings and other types that can be comparable.

We could rewrite an earlier program that computed the average of a list of numbers⁵ entered by the user using a list.

First, the program to compute an average without a list:⁶

```
total = 0
count = 0
while (True):
    inp = input('Enter a number: ')
    if inp == 'done': break
    value = float(inp)
    total = total + value
    count = count + 1

average = total / count
print('Average:', average)
```

Code: <https://www.py4e.com/code3/avenum.py>⁸

In this program, we have `count` and `total` variables to keep the number and⁹ running total of the user's numbers as we repeatedly prompt the user for a number.

We could simply remember each number as the user entered it and use built-in¹⁰ functions to compute the sum and count at the end.

```
numlist = list()11
while (True):
    inp = input('Enter a number: ')
    if inp == 'done': break
    value = float(inp)
```

```

numlist.append(value) 1

average = sum(numlist) / len(numlist) 2
print('Average:', average)

# Code: https://www.py4e.com/code3/avelist.py 3

```

We make an empty list before the loop starts, and then each time we have a number,⁴ we append it to the list. At the end of the program, we simply compute the sum of the numbers in the list and divide it by the count of the numbers in the list to come up with the average.

8.9 Lists and strings⁵

A string is a sequence of characters and a list is a sequence of values, but a list⁶ of characters is not the same as a string. To convert from a string to a list of characters, you can use `list`:

```

>>> s = 'spam' 7
>>> t = list(s)
>>> print(t)
['s', 'p', 'a', 'm']

```

Because `list` is the name of a built-in function, you should avoid using it as⁸ a variable name. I also avoid the letter “l” because it looks too much like the number “1”. So that’s why I use “t”.

The `list` function breaks a string into individual letters. If you want to break a⁹ string into words, you can use the `split` method:

```

>>> s = 'pining for the fjords' 10
>>> t = s.split()
>>> print(t)
['pining', 'for', 'the', 'fjords']
>>> print(t[2])
the

```

Once you have used `split` to break the string into a list of words, you can use the¹¹ index operator (square bracket) to look at a particular word in the list.

You can call `split` with an optional argument called a *delimiter* that specifies¹² which characters to use as word boundaries. The following example uses a hyphen as a delimiter:

```

>>> s = 'spam-spam-spam' 13
>>> delimiter = '-'
>>> s.split(delimiter)
['spam', 'spam', 'spam']

```

`join` is the inverse of `split`. It takes a list of strings and concatenates the elements. 1
`join` is a string method, so you have to invoke it on the delimiter and pass the list as a parameter:

```
>>> t = ['pining', 'for', 'the', 'fjords']2
>>> delimiter = ' '
>>> delimiter.join(t)
'pining for the fjords'
```

In this case the delimiter is a space character, so `join` puts a space between words. 3
To concatenate strings without spaces, you can use the empty string, "", as a delimiter.

8.10 Parsing lines 4

Usually when we are reading a file we want to do something to the lines other than just printing the whole line. Often we want to find the “interesting lines” and then *parse* the line to find some interesting *part* of the line. What if we wanted to print out the day of the week from those lines that start with “From”? 5

```
From stephen.marquard@uct.ac.za Sat Jan 5 09:14:16 20086
```

The `split` method is very effective when faced with this kind of problem. We can 7 write a small program that looks for lines where the line starts with “From”, `split` those lines, and then print out the third word in the line:

```
fhand = open('mbox-short.txt')8
for line in fhand:
    line = line.rstrip()
    if not line.startswith('From '): continue
    words = line.split()
    print(words[2])

# Code: https://www.py4e.com/code3/search5.py
```

The program produces the following output: 9

```
Sat10
Fri
Fri
Fri
...

```

Later, we will learn increasingly sophisticated techniques for picking the lines to 11 work on and how we pull those lines apart to find the exact bit of information we are looking for.




Figure 8.1: Variables and Objects 2

8.11 Objects and values³

If we execute these assignment statements:⁴

```
a = 'banana'5
b = 'banana'
```

we know that `a` and `b` both refer to a string, but we don't know whether they refer⁶ to the *same* string. There are two possible states:

In one case, `a` and `b` refer to two different objects that have the same value. In the⁷ second case, they refer to the same object.

To check whether two variables refer to the same object, you can use the `is` operator.⁸

```
>>> a = 'banana'9
>>> b = 'banana'
>>> a is b
True
```

In this example, Python only created one string object, and both `a` and `b` refer to¹⁰ it.

But when you create two lists, you get two objects:¹¹

```
>>> a = [1, 2, 3]12
>>> b = [1, 2, 3]
>>> a is b
False
```

In this case we would say that the two lists are *equivalent*, because they have the same elements, but not *identical*, because they are not the same object. If two objects are identical, they are also equivalent, but if they are equivalent, they are not necessarily identical.

Until now, we have been using “object” and “value” interchangeably, but it is more precise to say that an object has a value. If you execute `a = [1, 2, 3]`, `a` refers to a list object whose value is a particular sequence of elements. If another list has the same elements, we would say it has the same value.¹⁴

8.12 Aliasing¹

If `a` refers to an object and you assign `b = a`, then both variables refer to the same ² object:

```
>>> a = [1, 2, 3]3
```

```
>>> b = a4
```

```
>>> b is a
```

True ⁵

The association of a variable with an object is called a *reference*. In this example, ⁶ there are two references to the same object.

An object with more than one reference has more than one name, so we say that ⁷ the object is *aliased*.

If the aliased object is mutable, changes made with one alias affect the other: ⁸

```
>>> b[0] = 179
```

```
>>> print(a)10
```

[17, 2, 3] ¹¹

Although this behavior can be useful, it is error-prone. In general, it is safer to ¹² avoid aliasing when you are working with mutable objects.

For immutable objects like strings, aliasing is not as much of a problem. In this ¹³ example:

```
a = 'banana'14
```

```
b = 'banana'15
```

it almost never makes a difference whether `a` and `b` refer to the same string or not. ¹⁶

8.13 List arguments¹⁷

When you pass a list to a function, the function gets a reference to the list. If ¹⁸ the function modifies a list parameter, the caller sees the change. For example, `delete_head` removes the first element from a list:

```
def delete_head(t):19
```

```
    del t[0]20
```

Here's how it is used: ²¹

```
>>> letters = ['a', 'b', 'c']22
```

```
>>> delete_head(letters)23
```

```
>>> print(letters)
```

['b', 'c'] ²⁴

The parameter `t` and the variable `letters` are aliases for the same object.¹

It is important to distinguish between operations that modify lists and operations² that create new lists. For example, the `append` method modifies a list, but the `+` operator creates a new list:

```
>>> t1 = [1, 2]      3
>>> t2 = t1.append(3)
>>> print(t1)
[1, 2, 3]
>>> print(t2)
None

>>> t3 = t1 + [3]
>>> print(t3)
[1, 2, 3]
>>> t1 is t3
False
```

This difference is important when you write functions that are supposed to modify⁴ lists. For example, this function *does not* delete the head of a list:

```
def bad_delete_head(t):      5
    t = t[1:]                 # WRONG!
```

The slice operator creates a new list and the assignment makes `t` refer to it, but⁶ none of that has any effect on the list that was passed as an argument.

An alternative is to write a function that creates and returns a new list. For⁷ example, `tail` returns all but the first element of a list:

```
def tail(t):      8
    return t[1:]
```

This function leaves the original list unmodified. Here's how it is used:⁹

```
>>> letters = ['a', 'b', 'c'] 10
>>> rest = tail(letters)
>>> print(rest)
['b', 'c']
```

Exercise 1: Write a function called `chop` that takes a list and modifies it, removing¹¹ the first and last elements, and returns `None`. Then write a function called `middle` that takes a list and returns a new list that contains all but the first and last elements.

8.14 Debugging ¹

Careless use of lists (and other mutable objects) can lead to long hours of debugging.²
Here are some common pitfalls and ways to avoid them:

1. Don't forget that most list methods modify the argument and return `None`.³
This is the opposite of the string methods, which return a new string and leave the original alone.

If you are used to writing string code like this:⁴

```
word = word.strip()5
```

It is tempting to write list code like this:⁶

```
t = t.sort() # WRONG!7
```

Because `sort` returns `None`, the next operation you perform with `t` is likely⁸ to fail.

Before using list methods and operators, you should read the documentation⁹ carefully and then test them in interactive mode. The methods and operators that lists share with other sequences (like strings) are documented at:

docs.python.org/library/stdtypes.html#common-sequence-operations¹⁰

The methods and operators that only apply to mutable sequences are documented at:¹¹

docs.python.org/library/stdtypes.html#mutable-sequence-types¹²

2. Pick an idiom and stick with it.¹³

Part of the problem with lists is that there are too many ways to do things.¹⁴
For example, to remove an element from a list, you can use `pop`, `remove`, `del`, or even a slice assignment.

To add an element, you can use the `append` method or the `+` operator. But¹⁵ don't forget that these are right:

```
t.append(x)16
```

```
t = t + [x]17
```

And these are wrong:¹⁸

```
t.append([x]) # WRONG!19
```

```
t = t.append(x) # WRONG!20
```

```
t + [x] # WRONG!21
```

```
t = t + x # WRONG!22
```

Try out each of these examples in interactive mode to make sure you understand what they do. Notice that only the last one causes a runtime error; the other three are legal, but they do the wrong thing.²³

3. Make copies to avoid aliasing.²⁴

If you want to use a method like `sort` that modifies the argument, but you²⁵ need to keep the original list as well, you can make a copy.

```
orig = t[:] 1
t.sort()
```

In this example you could also use the built-in function `sorted`, which returns 2
a new, sorted list and leaves the original alone. But in that case you should
avoid using `sorted` as a variable name!

4. Lists, `split`, and files 3

When we read and parse files, there are many opportunities to encounter 4
input that can crash our program so it is a good idea to revisit the *guardian*
pattern when it comes to writing programs that read through a file and look
for a “needle in the haystack”.

Let’s revisit our program that is looking for the day of the week on the from 5
lines of our file:

```
From stephen.marquard@uct.ac.za Sat Jan 5 09:14:16 2008 6
```

Since we are breaking this line into words, we could dispense with the use 7
of `startswith` and simply look at the first word of the line to determine if
we are interested in the line at all. We can use `continue` to skip lines that
don’t have “From” as the first word as follows:

```
fhand = open('mbox-short.txt') 8
for line in fhand:
    words = line.split()
    if words[0] != 'From' : continue
    print(words[2])
```

This looks much simpler and we don’t even need to do the `rstrip` to remove 9
the newline at the end of the file. But is it better?

```
python search8.py 10
Sat
Traceback (most recent call last):
  File "search8.py", line 5, in <module>
    if words[0] != 'From' : continue
IndexError: list index out of range
```

It kind of works and we see the day from the first line (Sat), but then the 11
program fails with a traceback error. What went wrong? What messed-up
data caused our elegant, clever, and very Pythonic program to fail?

You could stare at it for a long time and puzzle through it or ask someone 12
for help, but the quicker and smarter approach is to add a `print` statement.
The best place to add the `print` statement is right before the line where the
program failed and print out the data that seems to be causing the failure.

Now this approach may generate a lot of lines of output, but at least you will 13
immediately have some clue as to the problem at hand. So we add a `print` of
the variable `words` right before line five. We even add a prefix “Debug:” to
the line so we can keep our regular output separate from our debug output.

```

for line in fhand:          1
    words = line.split()
    print('Debug:', words)
    if words[0] != 'From' : continue
    print(words[2])

```

When we run the program, a lot of output scrolls off the screen but at the end, we see our debug output and the traceback so we know what happened just before the traceback. 2

```

Debug: ['X-DSPAM-Confidence:', '0.8475'] 3
Debug: ['X-DSPAM-Probability:', '0.0000']
Debug: []
Traceback (most recent call last):
  File "search9.py", line 6, in <module>
    if words[0] != 'From' : continue
IndexError: list index out of range

```

Each debug line is printing the list of words which we get when we split 4 the line into words. When the program fails, the list of words is empty []. If we open the file in a text editor and look at the file, at that point it looks as follows:

```

X-DSPAM-Result: Innocent      5
X-DSPAM-Processed: Sat Jan  5 09:14:16 2008
X-DSPAM-Confidence: 0.8475
X-DSPAM-Probability: 0.0000

```

Details: <http://source.sakaiproject.org/viewsvn/?view=rev&rev=39772> 6

The error occurs when our program encounters a blank line! Of course there 7 are “zero words” on a blank line. Why didn’t we think of that when we were writing the code? When the code looks for the first word (`word[0]`) to check to see if it matches “From”, we get an “index out of range” error.

This of course is the perfect place to add some *guardian* code to avoid checking 8 the first word if the first word is not there. There are many ways to protect this code; we will choose to check the number of words we have before we look at the first word:

```

fhand = open('mbox-short.txt')      9
count = 0
for line in fhand:
    words = line.split()
    # print('Debug:', words)
    if len(words) == 0 : continue
    if words[0] != 'From' : continue
    print(words[2])

```

First we commented out the debug print statement instead of removing it, 10 in case our modification fails and we need to debug again. Then we added a guardian statement that checks to see if we have zero words, and if so, we use `continue` to skip to the next line in the file.

We can think of the two `continue` statements as helping us refine the set of lines which are “interesting” to us and which we want to process some more. A line which has no words is “uninteresting” to us so we skip to the next line. A line which does not have “From” as its first word is uninteresting to us so we skip it.

The program as modified runs successfully, so perhaps it is correct. Our guardian statement does make sure that the `words[0]` will never fail, but perhaps it is not enough. When we are programming, we must always be thinking, “What might go wrong?”

Exercise 2: Figure out which line of the above program is still not properly guarded. See if you can construct a text file which causes the program to fail and then modify the program so that the line is properly guarded and test it to make sure it handles your new text file.

Exercise 3: Rewrite the guardian code in the above example without two `if` statements. Instead, use a compound logical expression using the `or` logical operator with a single `if` statement.

8.15 Glossary⁵

aliasing A circumstance where two or more variables refer to the same object. ⁶

delimiter A character or string used to indicate where a string should be split.

element One of the values in a list (or other sequence); also called items. ⁷

equivalent Having the same value.

index An integer value that indicates an element in a list. ⁸

identical Being the same object (which implies equivalence). ⁹

list A sequence of values. ¹⁰

list traversal The sequential accessing of each element in a list. ¹¹

nested list A list that is an element of another list. ¹²

object Something a variable can refer to. An object has a type and a value. ¹³

reference The association between a variable and its value. ¹⁴

8.16 Exercises¹⁵

Exercise 4: Find all unique words in a file¹⁶

Shakespeare used over 20,000 words in his works. But how would you determine that? How would you produce the list of all the words that Shakespeare used? Would you download all his work, read it and track all unique words by hand?

Let’s use Python to achieve that instead. List all unique words, sorted in alphabetical order, that are stored in a file `romeo.txt` containing a subset of Shakespeare’s work.

To get started, download a copy of the file www.py4e.com/code3/romeo.txt. Create a list of unique words, which will contain the final result. Write a program to open the file `romeo.txt` and read it line by line. For each line, split the line into a list of words using the `split` function. For each word, check to see if the word is

already in the list of unique words. If the word is not in the list of unique words, add it to the list. When the program completes, sort and print the list of unique words in alphabetical order.

```
Enter file: romeo.txt
['Arise', 'But', 'It', 'Juliet', 'Who', 'already',
'and', 'breaks', 'east', 'envious', 'fair', 'grief',
'is', 'kill', 'light', 'moon', 'pale', 'sick', 'soft',
'sun', 'the', 'through', 'what', 'window',
'with', 'yonder']
```

2

Exercise 5: Minimalist Email Client.3

MBOX (mail box) is a popular file format to store and share a collection of emails.⁴ This was used by early email servers and desktop apps. Without getting into too many details, MBOX is a text file, which stores emails consecutively. Emails are separated by a special line which starts with `From` (notice the space). Importantly, lines starting with `From:` (notice the colon) describes the email itself and does not act as a separator. Imagine you wrote a minimalist email app, that lists the email of the senders in the user's Inbox and counts the number of emails.

Write a program to read through the mail box data and when you find a line that starts with “`From`”, you will split the line into words using the `split` function. We are interested in who sent the message, which is the second word on the `From` line.

```
From stephen.marquard@uct.ac.za Sat Jan 5 09:14:16 2008 6
```

You will parse the `From` line and print out the second word for each `From` line,⁷ then you will also count the number of `From` (not `From:`) lines and print out a count at the end. This is a good sample output with a few lines removed:

```
python fromcount.py
Enter a file name: mbox-short.txt
stephen.marquard@uct.ac.za
louis@media.berkeley.edu
zqian@umich.edu
```

8

```
[...some output removed...]
ray@media.berkeley.edu
cwen@iupui.edu
cwen@iupui.edu
cwen@iupui.edu
There were 27 lines in the file with From as the first word
```

9

Exercise 6: 10

Rewrite the program that prompts the user for a list of numbers and prints out the maximum and minimum of the numbers at the end when the user enters “done”. Write the program to store the numbers the user enters in a list and use the `max()` and `min()` functions to compute the maximum and minimum numbers after the loop completes.

11

```
Enter a number: 6 1
Enter a number: 2 2
Enter a number: 9 3
Enter a number: 3 4
Enter a number: 5 5
Enter a number: done 6
Maximum: 9.0 7
Minimum: 2.0 8
```


Chapter 9¹

Dictionaries²

A *dictionary* is like a list, but more general. In a list, the index positions have to ³ be integers; in a dictionary, the indices can be (almost) any type.

You can think of a dictionary as a mapping between a set of indices (which are ⁴ called *keys*) and a set of values. Each key maps to a value. The association of a key and a value is called a *key-value pair* or sometimes an *item*.

As an example, we'll build a dictionary that maps from English to Spanish words, ⁵ so the keys and the values are all strings.

The function `dict` creates a new dictionary with no items. Because `dict` is the ⁶ name of a built-in function, you should avoid using it as a variable name.

```
>>> eng2sp = dict() 7
>>> print(eng2sp)
{}
```

The curly brackets, {}, represent an empty dictionary. To add items to the dictionary, you can use square brackets: ⁸

```
>>> eng2sp['one'] = 'uno' 9
```

This line creates an item that maps from the key 'one' to the value "uno". If we ¹⁰ print the dictionary again, we see a key-value pair with a colon between the key and value:

```
>>> print(eng2sp) 11
{'one': 'uno'}
```

This output format is also an input format. For example, you can create a new ¹² dictionary with three items.

```
>>> eng2sp = {'one': 'uno', 'two': 'dos', 'three': 'tres'} 13
>>> print(eng2sp)
{'one': 'uno', 'two': 'dos', 'three': 'tres'}
```

Since Python 3.7x the order of key-value pairs is the same as their input order, i.e. dictionaries are now ordered structures. 1

But that doesn't really matter because the elements of a dictionary are never indexed with integer indices. Instead, you use the keys to look up the corresponding values: 2

```
>>> print(eng2sp['two']) 3
'dos'
```

The key 'two' always maps to the value "dos" so the order of the items doesn't matter. 4

If the key isn't in the dictionary, you get an exception: 5

```
>>> print(eng2sp['four']) 6
KeyError: 'four'
```

The `len` function works on dictionaries; it returns the number of key-value pairs: 7

```
>>> len(eng2sp) 8
3
```

The `in` operator works on dictionaries; it tells you whether something appears as a *key* in the dictionary (appearing as a value is not good enough). 9

```
>>> 'one' in eng2sp 10
True
>>> 'uno' in eng2sp
False
```

To see whether something appears as a value in a dictionary, you can use the `values` method, which returns the values as a type that can be converted to a list, and then use the `in` operator: 11

```
>>> vals = list(eng2sp.values()) 12
>>> 'uno' in vals
True
```

The `in` operator uses different algorithms for lists and dictionaries. For lists, it uses a linear search algorithm. As the list gets longer, the search time gets longer in direct proportion to the length of the list. For dictionaries, Python uses an algorithm called a *hash table* that has a remarkable property: the `in` operator takes about the same amount of time no matter how many items there are in a dictionary. I won't explain why hash functions are so magical, but you can read more about it at [wikipedia.org/wiki/Hash_table](https://en.wikipedia.org/wiki/Hash_table).¹ 13

¹If you do want to learn more about hash tables, there is a course at <https://www.cc4e.com> that explores how the programming language C implements a Python dictionary.

Exercise 1: Download a copy of the file 1

www.py4e.com/code3/words.txt 2

Write a program that reads the words in *words.txt* and stores them as keys in 3 a dictionary. It doesn't matter what the values are. Then you can use the `in` operator as a fast way to check whether a string is in the dictionary.

9.1 Dictionary as a set of counters⁴

Suppose you are given a string and you want to count how many times each letter 5 appears. There are several ways you could do it:

1. You could create 26 variables, one for each letter of the alphabet. Then you 6 could traverse the string and, for each character, increment the corresponding counter, probably using a chained conditional.
2. You could create a list with 26 elements. Then you could convert each character to a number (using the built-in function `ord`), use the number as an index into the list, and increment the appropriate counter.
3. You could create a dictionary with characters as keys and counters as the corresponding values. The first time you see a character, you would add an item to the dictionary. After that you would increment the value of an existing item.

Each of these options performs the same computation, but each of them implements 7 that computation in a different way.

An *implementation* is a way of performing a computation; some implementations 8 are better than others. For example, an advantage of the dictionary implementation is that we don't have to know ahead of time which letters appear in the string and we only have to make room for the letters that do appear.

Here is what the code might look like: 9

```
word = 'brontosaurus' 10
d = dict()
for c in word:
    if c not in d:
        d[c] = 1
    else:
        d[c] = d[c] + 1
print(d)
```

We are effectively computing a *histogram*, which is a statistical term for a set of 11 counters (or frequencies).

The `for` loop traverses the string. Each time through the loop, if `c` is not in the dictionary, we create a new item with key `c` and the initial value 1 (since we have seen this letter once). If `c` is already in the dictionary we increment `d[c]`.

Here's the output of the program: 13

```
{'b': 1, 'r': 2, 'o': 2, 'n': 1, 't': 1, 's': 2, 'a': 1, 'u': 2} 1
```

The histogram indicates that the letters “a” and “b” appear once; “o” appears 2 twice, and so on.

Dictionaries have a method called `get` that takes a key and a default value. If the 3 key appears in the dictionary, `get` returns the corresponding value; otherwise it returns the default value. For example:

```
>>> counts = { 'chuck' : 1 , 'annie' : 42, 'jan': 100} 4
>>> print(counts.get('jan', 0))
100
>>> print(counts.get('tim', 0))
0
```

We can use `get` to write our histogram loop more concisely. Because the get 5 method automatically handles the case where a key is not in a dictionary, we can reduce four lines down to one and eliminate the `if` statement.

```
word = 'brontosaurus' 6
d = dict()
for c in word:
    d[c] = d.get(c,0) + 1
print(d)
```

The use of the `get` method to simplify this counting loop ends up being a very 7 commonly used “idiom” in Python and we will use it many times in the rest of the book. So you should take a moment and compare the loop using the `if` statement and `in` operator with the loop using the `get` method. They do exactly the same thing, but one is more succinct.

9.2 Dictionaries and files⁸

One of the common uses of a dictionary is to count the occurrence of words in a 9 file with some written text. Let’s start with a very simple file of words taken from the text of *Romeo and Juliet*.

For the first set of examples, we will use a shortened and simplified version of 10 the text with no punctuation. Later we will work with the text of the scene with punctuation included.

```
But soft what light through yonder window breaks 11
It is the east and Juliet is the sun
Arise fair sun and kill the envious moon
Who is already sick and pale with grief
```

We will write a Python program to read through the lines of the file, break each 12 line into a list of words, and then loop through each of the words in the line and count each word using a dictionary.

You will see that we have two `for` loops. The outer loop is reading the lines of the file and the inner loop is iterating through each of the words on that particular line. This is an example of a pattern called *nested loops* because one of the loops is the *outer* loop and the other loop is the *inner* loop. 1

Because the inner loop executes all of its iterations each time the outer loop makes a single iteration, we think of the inner loop as iterating “more quickly” and the outer loop as iterating more slowly. 2

The combination of the two nested loops ensures that we will count every word on every line of the input file. 3

```
fname = input('Enter the file name: ') 4
try:
    fhand = open(fname)
except:
    print('File cannot be opened:', fname)
    exit()

counts = dict()
for line in fhand:
    words = line.split()
    for word in words:
        if word not in counts:
            counts[word] = 1
        else:
            counts[word] += 1

print(counts)
```

Code: <https://www.py4e.com/code3/count1.py> 5

In our `else` statement, we use the more compact alternative for incrementing a variable. `counts[word] += 1` is equivalent to `counts[word] = counts[word] + 1`. Either method can be used to change the value of a variable by any desired amount. Similar alternatives exist for `-=`, `*=`, and `/=`. 6

When we run the program, we see a raw dump of all of the counts in unsorted hash order. (the *romeo.txt* file is available at www.py4e.com/code3/romeo.txt) 7

```
python count1.py
Enter the file name: romeo.txt 8
{'But': 1, 'soft': 1, 'what': 1, 'light': 1, 'through': 1, 'yonder': 1,
>window': 1, 'breaks': 1, 'It': 1, 'is': 3, 'the': 3, 'east': 1, 'and': 3,
'Juliet': 1, 'sun': 2, 'Arise': 1, 'fair': 1, 'kill': 1, 'envious': 1,
'moon': 1, 'Who': 1, 'already': 1, 'sick': 1, 'pale': 1, 'with': 1,
'grief': 1}
```

It is a bit inconvenient to look through the dictionary to find the most common words and their counts, so we need to add some more Python code to get us the output that will be more helpful. 9

9.3 Looping and dictionaries¹

If you use a dictionary as the sequence in a `for` statement, it traverses the keys of ² the dictionary. This loop prints each key and the corresponding value:

```
counts = { 'chuck' : 1 , 'annie' : 42, 'jan': 100}3
for key in counts:
    print(key, counts[key])
```

Here's what the output looks like: ⁴

```
chuck 1 5
annie 42
jan 100
```

Again, the keys are ordered. ⁶

We can use this pattern to implement the various loop idioms that we have described earlier. For example if we wanted to find all the entries in a dictionary with a value above ten, we could write the following code: ⁷

```
counts = { 'chuck' : 1 , 'annie' : 42, 'jan': 100}8
for key in counts:
    if counts[key] > 10 :
        print(key, counts[key])
```

The `for` loop iterates through the *keys* of the dictionary, so we must use the index operator to retrieve the corresponding *value* for each key. Here's what the output looks like: ⁹

```
annie 42 10
jan 100
```

We see only the entries with a value above 10. ¹¹

If you want to print the keys in alphabetical order, you first make a list of the keys in the dictionary using the `keys` method available in dictionary objects, and then sort that list and loop through the sorted list, looking up each key and printing out key-value pairs in sorted order as follows: ¹²

```
counts = { 'chuck' : 1 , 'annie' : 42, 'jan': 100}13
lst = list(counts.keys())
print(lst)
lst.sort()
print(lst)
for key in lst:
    print(key, counts[key])
```

Here's what the output looks like: ¹⁴

```
['chuck', 'annie', 'jan']1
['annie', 'chuck', 'jan']
annie 42
chuck 1
jan 100
```

First you see the list of keys in non-alphabetical order that we get from the `keys`² method. Then we see the key-value pairs in alphabetical order from the `for` loop.

9.4 Advanced text parsing³

In the above example using the file *romeo.txt*, we made the file as simple as possible⁴ by removing all punctuation by hand. The actual text has lots of punctuation, as shown below.

```
But, soft! what light through yonder window breaks?5
It is the east, and Juliet is the sun.6
Arise, fair sun, and kill the envious moon,7
Who is already sick and pale with grief,8
```

Since the Python `split` function looks for spaces and treats words as tokens separated by spaces, we would treat the words “soft!” and “soft” as *different* words and create a separate dictionary entry for each word.⁹

Also since the file has capitalization, we would treat “who” and “Who” as different¹⁰ words with different counts.

We can solve both these problems by using the string methods `lower`, `punctuation`,¹¹ and `translate`. The `translate` is the most subtle of the methods. Here is the documentation for `translate`:

```
line.translate(str.maketrans(fromstr, tostr, deletestr))12
```

Replace the characters in `fromstr` with the character in the same position in `tostr` and delete all characters that are in `deletestr`. The `fromstr` and `tostr` can be empty strings and the `deletestr` parameter can be omitted.¹³

We will not specify the `tostr` but we will use the `deletestr` parameter to delete all of the punctuation. We will even let Python tell us the list of characters that it considers “punctuation”:¹⁴

```
>>> import string15
>>> string.punctuation
'!"#$%&\\"()*+, -./:;<=>?@[\\"]^_`{|}~'
```

The parameters used by `translate` were different in Python 2.0.¹⁶

We make the following modifications to our program:¹⁷

```
import string18
```

```

fname = input('Enter the file name: ')
try:
    fhand = open(fname)
except:
    print('File cannot be opened:', fname)
    exit()

counts = dict()
for line in fhand:
    line = line.rstrip()
    # First two parameters are empty strings
    line = line.translate(line.maketrans("", "", string.punctuation))
    line = line.lower()
    words = line.split()
    for word in words:
        if word not in counts:
            counts[word] = 1
        else:
            counts[word] += 1

print(counts)

```

Code: <https://www.py4e.com/code3/count2.py> 2

Part of learning the “Art of Python” or “Thinking Pythonically” is realizing that 3 Python often has built-in capabilities for many common data analysis problems. Over time, you will see enough example code and read enough of the documentation to know where to look to see if someone has already written something that makes your job much easier.

The following is an abbreviated version of the output: 4

```

Enter the file name: romeo-full.txt
{'romeo': 40, 'and': 42, 'juliet': 32, 'act': 1, '2': 2, 'scene': 2,
'ii': 1, 'capulets': 1, 'orchard': 2, 'enter': 1, 'he': 5, 'jest': 1,
'at': 9, 'scars': 1, 'that': 30, 'never': 2, 'felt': 1, 'a': 24,
>wound': 1, 'appears': 1, 'above': 6, 'window': 2, 'but': 18,
'soft': 1, 'what': 11, 'light': 5, 'through': 2, 'yonder': 2,
'breaks': 1, ...}

```

Looking through this output is still unwieldy and we can use Python to give us 6 exactly what we are looking for, but to do so, we need to learn about Python *tuples*. We will pick up this example once we learn about tuples.

9.5 Debugging⁷

As you work with bigger datasets it can become unwieldy to debug by printing and 8 checking data by hand. Here are some suggestions for debugging large datasets:

Scale down the input If possible, reduce the size of the dataset. For example 1
if the program reads a text file, start with just the first 10 lines, or with the
smallest example you can find. You can either edit the files themselves, or
(better) modify the program so it reads only the first n lines.

If there is an error, you can reduce n to the smallest value that manifests the 2
error, and then increase it gradually as you find and correct errors.

Check summaries and types Instead of printing and checking the entire 3
dataset, consider printing summaries of the data: for example, the number
of items in a dictionary or the total of a list of numbers.

A common cause of runtime errors is a value that is not the right type. For 4
debugging this kind of error, it is often enough to print the type of a value.

Write self-checks Sometimes you can write code to check for errors automatically. 5
For example, if you are computing the average of a list of numbers, you could check that the result is not greater than the largest element in the list or less than the smallest. This is called a “sanity check” because it detects results that are “completely illogical”.

Another kind of check compares the results of two different computations to 6
see if they are consistent. This is called a “consistency check”.

Pretty print the output Formatting debugging output can make it easier to 7
spot an error.

Again, time you spend building scaffolding can reduce the time you spend debug- 8
ging.

9.6 Glossary⁹

dictionary A mapping from a set of keys to their corresponding values. 10

hashtable The algorithm used to implement Python dictionaries. 11

hash function A function used by a hashtable to compute the location for a key. 12

histogram A set of counters. 13

implementation A way of performing a computation. 14

item Another name for a key-value pair. 15

key An object that appears in a dictionary as the first part of a key-value pair. 16

key-value pair The representation of the mapping from a key to a value. 17

lookup A dictionary operation that takes a key and finds the corresponding value. 18

nested loops When there are one or more loops “inside” of another loop. The 19
inner loop runs to completion each time the outer loop runs once.

value An object that appears in a dictionary as the second part of a key-value 20
pair. This is more specific than our previous use of the word “value”.

9.7 Exercises²¹

Exercise 2: Write a program that categorizes each mail message by which day of 22
the week the commit was done. To do this look for lines that start with “From”,

then look for the third word and keep a running count of each of the days of the week. At the end of the program print out the contents of your dictionary (order does not matter).¹

Sample Line:²

```
From stephen.marquard@uct.ac.za Sat Jan  5 09:14:16 20083
```

Sample Execution:⁴

```
python dow.py5
Enter a file name: mbox-short.txt6
{'Fri': 20, 'Thu': 6, 'Sat': 1}7
```

Exercise 3: Write a program to read through a mail log, build a histogram using⁸ a dictionary to count how many messages have come from each email address, and print the dictionary.

```
Enter file name: mbox-short.txt9
{'gopal.ramasammycook@gmail.com': 1, 'louis@media.berkeley.edu': 3,10
 'cwen@iupui.edu': 5, 'antranig@caret.cam.ac.uk': 1,11
 'rjlowe@iupui.edu': 2, 'gsilver@umich.edu': 3,12
 'david.horwitz@uct.ac.za': 4, 'wagnermr@iupui.edu': 1,13
 'zqian@umich.edu': 4, 'stephen.marquard@uct.ac.za': 2,14
 'ray@media.berkeley.edu': 1}15
```

Exercise 4: Add code to the above program to figure out who has sent the most¹⁶ messages in the file. After all the data has been read and the dictionary has been created, look through the dictionary using a maximum loop (see Chapter 5: Maximum and minimum loops) to find who has the most messages and print how many messages the person has.

```
Enter a file name: mbox-short.txt17
cwen@iupui.edu 518
```

```
Enter a file name: mbox.txt19
zqian@umich.edu 19520
```

Exercise 5: This program records the domain name (instead of the address) where²¹ the message was sent from instead of who the mail came from (i.e., the whole email address). At the end of the program, print out the contents of your dictionary.

```
python schoolcount.py22
Enter a file name: mbox-short.txt23
{'media.berkeley.edu': 4, 'uct.ac.za': 6, 'umich.edu': 7,24
 'gmail.com': 1, 'caret.cam.ac.uk': 1, 'iupui.edu': 8}25
```

Chapter 10¹

Tuples²

10.1 Tuples are immutable³

A tuple¹ is a sequence of values much like a list. The values stored in a tuple can be any type, and they are indexed by integers. The important difference is that tuples are *immutable*. Tuples are also *comparable* and *hashable* so we can sort lists of them and use tuples as key values in Python dictionaries.⁴

Syntactically, a tuple is a comma-separated list of values:⁵

```
>>> t = 'a', 'b', 'c', 'd', 'e'6
```

Although it is not necessary, it is common to enclose tuples in parentheses to help⁷ us quickly identify tuples when we look at Python code:

```
>>> t = ('a', 'b', 'c', 'd', 'e')8
```

To create a tuple with a single element, you have to include the final comma:⁹

```
>>> t1 = ('a',)10
>>> type(t1)
<type 'tuple'>
```

Without the comma Python treats ('a') as an expression with a string in parentheses that evaluates to a string:¹¹

```
>>> t2 = ('a')12
>>> type(t2)
<type 'str'>
```

Another way to construct a tuple is the built-in function `tuple`. With no argument,¹³ it creates an empty tuple:

¹Fun fact: The word “tuple” comes from the names given to sequences of numbers of varying lengths: single, double, triple, quadruple, quintuple, sextuple, septuple, etc.

```
>>> t = tuple() 1
>>> print(t)
()
```

If the argument is a sequence (string, list, or tuple), the result of the call to `tuple` 2 is a tuple with the elements of the sequence:

```
>>> t = tuple('lupins') 3
>>> print(t)
('l', 'u', 'p', 'i', 'n', 's')
```

Because `tuple` is the name of a constructor, you should avoid using it as a variable 4 name.

Most list operators also work on tuples. The bracket operator indexes an element: 5

```
>>> t = ('a', 'b', 'c', 'd', 'e') 6
>>> print(t[0])
'a'
```

And the slice operator selects a range of elements. 7

```
>>> print(t[1:3]) 8
('b', 'c')
```

But if you try to modify one of the elements of the tuple, you get an error: 9

```
>>> t[0] = 'A' 10
TypeError: object doesn't support item assignment
```

You can't modify the elements of a tuple, but you can replace one tuple with 11 another:

```
>>> t = ('A',) + t[1:] 12
>>> print(t)
('A', 'b', 'c', 'd', 'e')
```

10.2 Comparing tuples 13

The comparison operators work with tuples and other sequences. Python starts by 14 comparing the first element from each sequence. If they are equal, it goes on to the next element, and so on, until it finds elements that differ. Subsequent elements are not considered (even if they are really big).

```
>>> (0, 1, 2) < (0, 3, 4) 15
True
>>> (0, 1, 2000000) < (0, 3, 4)
True
```

The `sort` function works the same way. It sorts primarily by first element, but in the case of a tie, it sorts by second element, and so on.

This feature lends itself to a pattern called *DSU* for 2

Decorate a sequence by building a list of tuples with one or more sort keys 3 preceding the elements from the sequence,

Sort the list of tuples using the Python built-in `sort`, and 4

Undecorate by extracting the sorted elements of the sequence. 5

For example, suppose you have a list of words and you want to sort them from 6 longest to shortest:

```
txt = 'but soft what light in yonder window breaks' 7
words = txt.split() 8
t = list() 9
for word in words: 10
    t.append((len(word), word)) 11

t.sort(reverse=True) 12

res = list() 13
for length, word in t: 14
    res.append(word) 15

print(res) 16

# Code: https://www.py4e.com/code3/soft.py 17
```

The first loop builds a list of tuples, where each tuple is a word preceded by its 18 length.

`sort` compares the first element, length, first, and only considers the second element to break ties. The keyword argument `reverse=True` tells `sort` to go in decreasing order. 19

The second loop traverses the list of tuples and builds a list of words in descending order of length. The four-character words are sorted in *reverse* alphabetical order, so “what” appears before “soft” in the following list. 20

The output of the program is as follows: 21

```
['yonder', 'window', 'breaks', 'light', 'what', 22
 'soft', 'but', 'in']
```

Of course the line loses much of its poetic impact when turned into a Python list 23 and sorted in descending word length order.

10.3 Tuple assignment¹

One of the unique syntactic features of the Python language is the ability to have a tuple on the left side and a sequence on the right side of an assignment statement. This allows you to assign more than one variable at a time to the given sequence.²

In this example we have a two-element tuple and assign the first and second elements of the tuple to the variables `x` and `y` in a single statement.³

```
>>> m = ('have', 'fun')4
>>> x, y = m
>>> x
'have'
>>> y
'fun'
>>>
```

This is more general than tuple-to-tuple assignment. Both tuples and lists are sequences, so this syntax works with a two element list as well.⁵

```
>>> m = ['have', 'fun']6
>>> x, y = m
>>> x
'have'
>>> y
'fun'
>>>
```

It is not magic, Python *roughly* translates the tuple assignment syntax to be the following:²⁷

```
>>> m = ('have', 'fun')8
>>> x = m[0]
>>> y = m[1]
>>> x
'have'
>>> y
'fun'
>>>
```

Stylistically when we use a tuple on the left side of the assignment statement, we omit the parentheses, but the following is an equally valid syntax:⁹

```
>>> m = ('have', 'fun')10
>>> (x, y) = m
>>> x
'have'
```

²Python does not translate the syntax literally. For example, if you try this with a dictionary, it will not work as you might expect.

```
>>> y 1
'fun'
>>>
```

A particularly clever application of tuple assignment allows us to *swap* the values² of two variables in a single statement:

```
>>> a, b = b, a 3
```

Both sides of this statement are tuples, but the left side is a tuple of variables;⁴ the right side is a tuple of expressions. Each value on the right side is assigned to its respective variable on the left side. All the expressions on the right side are evaluated before any of the assignments.

The number of variables on the left and the number of values on the right must be⁵ the same:

```
>>> a, b = 1, 2, 3 6
ValueError: too many values to unpack
```

More generally, the right side can be any kind of sequence (string, list, or tuple).⁷ For example, to split an email address into a user name and a domain, you could write:

```
>>> addr = 'monty@python.org' 8
>>> uname, domain = addr.split('@')
```

The return value from `split` is a list with two elements; the first element is assigned⁹ to `uname`, the second to `domain`.

```
>>> print(uname) 10
monty
>>> print(domain)
python.org
```

10.4 Dictionaries and tuples¹¹

Dictionaries have a method called `items` that returns a list of tuples, where each¹² tuple is a key-value pair:

```
>>> d = {'b':1, 'a':10, 'c':22} 13
>>> t = list(d.items())
>>> print(t)
[('b', 1), ('a', 10), ('c', 22)]
```

As you should expect from a dictionary, the items are in non-alphabetical order.¹⁴

However, since the list of tuples is a list, and tuples are comparable, we can now¹⁵ sort the list of tuples. Converting a dictionary to a list of tuples is a way for us to output the contents of a dictionary sorted by key:

```
>>> d = {'b':1, 'a':10, 'c':22} 1
>>> t = list(d.items())
>>> t
[('b', 1), ('a', 10), ('c', 22)]
>>> t.sort()
>>> t
[('a', 10), ('b', 1), ('c', 22)]
```

The new list is sorted in ascending alphabetical order by the key value. 2

10.5 Multiple assignment with dictionaries 3

Combining `items`, tuple assignment, and `for`, you can see a nice code pattern for 4 traversing the keys and values of a dictionary in a single loop:

```
d = {'a':10, 'b':1, 'c':22} 5
for key, val in d.items():
    print(val, key)
```

This loop has two *iteration variables* because `items` returns a list of tuples and `key`, 6 `val` is a tuple assignment that successively iterates through each of the key-value pairs in the dictionary.

For each iteration through the loop, both `key` and `val` are advanced to the next 7 key-value pair in the dictionary (still in hash order).

The output of this loop is: 8

```
10 a 9
1 b
22 c
```

Again, it is in hash key order (i.e., no particular order). 10

If we combine these two techniques, we can print out the contents of a dictionary 11 sorted by the *value* stored in each key-value pair.

To do this, we first make a list of tuples where each tuple is `(value, key)`. The 12 `items` method would give us a list of `(key, value)` tuples, but this time we want to sort by value, not key. Once we have constructed the list with the value-key tuples, it is a simple matter to sort the list in reverse order and print out the new, sorted list.

```
>>> d = {'a':10, 'b':1, 'c':22} 13
>>> l = list()
>>> for key, val in d.items() :
...     l.append( (val, key) )
...
>>> l
[(10, 'a'), (1, 'b'), (22, 'c')]
```

```
>>> l.sort(reverse=True)      1
>>> l
[(22, 'c'), (10, 'a'), (1, 'b')]
>>>
```

By carefully constructing the list of tuples to have the value as the first element 2 of each tuple, we can sort the list of tuples and get our dictionary contents sorted by value.

10.6 The most common words³

Coming back to our running example of the text from *Romeo and Juliet* Act 2,⁴ Scene 2, we can augment our program to use this technique to print the ten most common words in the text as follows:

```
import string
fhand = open('romeo-full.txt')
counts = dict()
for line in fhand:
    line = line.translate(str.maketrans('', '', string.punctuation))
    line = line.lower()
    words = line.split()
    for word in words:
        if word not in counts:
            counts[word] = 1
        else:
            counts[word] += 1

# Sort the dictionary by value      6
lst = list()
for key, val in list(counts.items()):
    lst.append((val, key))

lst.sort(reverse=True)

for key, val in lst[:10]:
    print(key, val)

# Code: https://www.py4e.com/code3/count3.py 7
```

The first part of the program which reads the file and computes the dictionary 8 that maps each word to the count of words in the document is unchanged. But instead of simply printing out `counts` and ending the program, we construct a list of `(val, key)` tuples and then sort the list in reverse order.

Since the value is first, it will be used for the comparisons. If there is more than 9 one tuple with the same value, it will look at the second element (the key), so tuples where the value is the same will be further sorted by the alphabetical order of the key.

At the end we write a nice `for` loop which does a multiple assignment iteration **1** and prints out the ten most common words by iterating through a slice of the list `(lst[:10])`.

So now the output finally looks like what we want for our word frequency analysis. **2**

```
61 i      3
42 and
40 romeo
34 to
34 the
32 thou
32 juliet
30 that
29 my
24 thee
```

The fact that this complex data parsing and analysis can be done with an easy-to-**4** understand 19-line Python program is one reason why Python is a good choice as a language for exploring information.

10.7 Using tuples as keys in dictionaries **5**

Because tuples are *hashable* and lists are not, if we want to create a *composite key* **6** to use in a dictionary we must use a tuple as the key.

We would encounter a composite key if we wanted to create a telephone directory **7** that maps from last-name, first-name pairs to telephone numbers. Assuming that we have defined the variables `last`, `first`, and `number`, we could write a dictionary assignment statement as follows:

```
directory[last,first] = number8
```

The expression in brackets is a tuple. We could use tuple assignment in a `for` loop **9** to traverse this dictionary.

```
for last, first in directory:          10
    print(first, last, directory[last,first])
```

This loop traverses the keys in `directory`, which are tuples. It assigns the elements **11** of each tuple to `last` and `first`, then prints the name and corresponding telephone number.

10.8 Sequences: strings, lists, and tuples - Oh **12** My!

I have focused on lists of tuples, but almost all of the examples in this chapter **13** also work with lists of lists, tuples of tuples, and tuples of lists. To avoid enumerating the possible combinations, it is sometimes easier to talk about sequences of sequences.

In many contexts, the different kinds of sequences (strings, lists, and tuples) can 1 be used interchangeably. So how and why do you choose one over the others?

To start with the obvious, strings are more limited than other sequences because 2 the elements have to be characters. They are also immutable. If you need the ability to change the characters in a string (as opposed to creating a new string), you might want to use a list of characters instead.

Lists are more common than tuples, mostly because they are mutable. But there 3 are a few cases where you might prefer tuples:

1. In some contexts, like a `return` statement, it is syntactically simpler to create 4 a tuple than a list. In other contexts, you might prefer a list.
2. If you want to use a sequence as a dictionary key, you have to use an immutable type like a tuple or string.
3. If you are passing a sequence as an argument to a function, using tuples reduces the potential for unexpected behavior due to aliasing.

Because tuples are immutable, they don't provide methods like `sort` and `reverse`, 5 which modify existing lists. However Python provides the built-in functions `sorted` and `reversed`, which take any sequence as a parameter and return a new sequence with the same elements in a different order.

10.9 List comprehension⁶

Sometimes you want to create a sequence by using data from another sequence. 7 You can achieve this by writing a for loop and appending one item at a time. For example, if you wanted to convert a list of strings – each string storing digits – into numbers that you can sum up, you would write:

```
list_of_ints_in_strings = ['42', '65', '12'] 8
list_of_ints = []
for x in list_of_ints_in_strings:
    list_of_ints.append(int(x))

print(sum(list_of_ints))
```

With list comprehension, the above code can be written in a more compact manner: 9

```
list_of_ints_in_strings = ['42', '65', '12'] 10
list_of_ints = [int(x) for x in list_of_ints_in_strings]
print(sum(list_of_ints))
```

10.10 Debugging¹

Lists, dictionaries and tuples are known generically as *data structures*; in this ² chapter we are starting to see compound data structures, like lists of tuples, and dictionaries that contain tuples as keys and lists as values. Compound data structures are useful, but they are prone to what I call *shape errors*; that is, errors caused when a data structure has the wrong type, size, or composition, or perhaps you write some code and forget the shape of your data and introduce an error. For example, if you are expecting a list with one integer and I give you a plain old integer (not in a list), it won't work.

10.11 Glossary³

comparable A type where one value can be checked to see if it is greater than,⁴ less than, or equal to another value of the same type. Types which are comparable can be put in a list and sorted.

data structure A collection of related values, often organized in lists, dictionaries,⁵ tuples, etc.

DSU Abbreviation of “decorate-sort-undecorate”, a pattern that involves building⁶ a list of tuples, sorting, and extracting part of the result.

gather The operation of assembling a variable-length argument tuple.⁷

hashable A type that has a hash function. Immutable types like integers, floats,⁸ and strings are hashable; mutable types like lists and dictionaries are not.

scatter The operation of treating a sequence as a list of arguments.⁹

shape (of a data structure) A summary of the type, size, and composition of¹⁰ a data structure.

singleton A list (or other sequence) with a single element.¹¹

tuple An immutable sequence of elements.¹²

tuple assignment An assignment with a sequence on the right side and a tuple¹³ of variables on the left. The right side is evaluated and then its elements are assigned to the variables on the left.

10.12 Exercises¹⁴

Exercise 1: Revise a previous program as follows: Read and parse the “From”¹⁵ lines and pull out the addresses from the line. Count the number of messages from each person using a dictionary.

After all the data has been read, print the person with the most commits by¹⁶ creating a list of (count, email) tuples from the dictionary. Then sort the list in reverse order and print out the person who has the most commits.

Sample Line:¹⁷

From stephen.marquard@uct.ac.za Sat Jan 5 09:14:16 2008¹⁸

Enter a file name: mbox-short.txt¹⁹

cwen@iupui.edu 5²⁰

```
Enter a file name: mbox.txt 1  
zqian@umich.edu 195
```

Exercise 2: This program counts the distribution of the hour of the day for each of the messages. You can pull the hour from the “From” line by finding the time string and then splitting that string into parts using the colon character. Once you have accumulated the counts for each hour, print out the counts, one per line, sorted by hour as shown below. 2

```
python timeofday.py 3  
Enter a file name: mbox-short.txt  
04 3  
06 1  
07 1  
09 2  
10 3  
11 6  
14 1  
15 2  
16 4  
17 2  
18 1  
19 1
```

Exercise 3: Write a program that reads a file and prints the *letters* in decreasing 4 order of frequency.

Your program should convert all the input to lower case and only count the letters a-z. Your program should not count spaces, digits, punctuation, or anything other than the letters a-z. Find text samples from several different languages and see how letter frequency varies between languages. Compare your results with the tables at https://wikipedia.org/wiki/Letter_frequencies. 5

Chapter 11¹

Regular expressions²

So far we have been reading through files, looking for patterns and extracting³ various bits of lines that we find interesting. We have been using string methods like `split` and `find` and using lists and string slicing to extract portions of the lines.

This task of searching and extracting is so common that Python has a very powerful⁴ module called *regular expressions* that handles many of these tasks quite elegantly. The reason we have not introduced regular expressions earlier in the book is because while they are very powerful, they are a little complicated and their syntax takes some getting used to.

Regular expressions are almost their own little programming language for searching⁵ and parsing strings. As a matter of fact, entire books have been written on the topic of regular expressions. In this chapter, we will only cover the basics of regular expressions. For more detail on regular expressions, see:

https://en.wikipedia.org/wiki/Regular_expression⁶

<https://docs.python.org/library/re.html>⁷

You can also bookmark the following page to experiment with regular expressions⁸ as you learn:

<https://regex101.com/>⁹

The regular expression module `re` must be imported into your program before you¹⁰ can use it. The simplest use of the regular expression module is the `search()` function. The following program demonstrates a trivial use of the search function.

```
# Search for lines that contain 'From' 11
import re 12
hand = open('mbox-short.txt') 13
for line in hand: 14
    line = line.rstrip() 15
    if re.search('From:', line): 16
        print(line) 17

# Code: https://www.py4e.com/code3/re01.py 18
```

We open the file, loop through each line, and use the regular expression `search()`¹ to only print out lines that contain the string “From:”. This program does not use the real power of regular expressions, since we could have just as easily used `line.find()` to accomplish the same result.

The power of the regular expressions comes when we add special characters to² the search string that allow us to more precisely control which lines match the string. Adding these special characters to our regular expression allows us to do sophisticated matching and extraction while writing very little code.

For example, the caret character is used in regular expressions to match “the beginning” of a line. We could change our program to only match lines where “From:” was at the beginning of the line as follows:³

```
# Search for lines that start with 'From' 4
import re
hand = open('mbox-short.txt')
for line in hand:
    line = line.rstrip()
    if re.search('^From:', line):
        print(line)
```

Code: <https://www.py4e.com/code3/re02.py> ⁵

Now we will only match lines that *start with* the string “From:”. This is still a⁶ very simple example that we could have done equivalently with the `startswith()` method from the string module. But it serves to introduce the notion that regular expressions contain special action characters that give us more control as to what will match the regular expression.

11.1 Character matching in regular expressions⁷

There are a number of other special characters that let us build even more powerful⁸ regular expressions. The most commonly used special character is the period or full stop, which matches any character.

In the following example, the regular expression `F..m:` would match any of the⁹ strings “From:”, “Fxxm:”, “F12m:”, or “F!@m:” since the period characters in the regular expression match any character.

```
# Search for lines that start with 'F', followed by 10
# 2 characters, followed by 'm:'
import re
hand = open('mbox-short.txt')
for line in hand:
    line = line.rstrip()
    if re.search('^F..m:', line):
        print(line)
```

Code: <https://www.py4e.com/code3/re03.py> ¹¹

This is particularly powerful when combined with the ability to indicate that a character can be repeated any number of times using the * or + characters in your regular expression. These special characters mean that instead of matching a single character in the search string, they match zero-or-more characters (in the case of the asterisk) or one-or-more of the characters (in the case of the plus sign). 1

We can further narrow down the lines that we match using a repeated *wild card* character in the following example: 2

```
# Search for lines that start with From and have an at sign 3
import re
hand = open('mbox-short.txt')
for line in hand:
    line = line.rstrip()
    if re.search('^From:.+@', line):
        print(line)
```

Code: <https://www.py4e.com/code3/re04.py> 4

The search string ^From:.+@ will successfully match lines that start with “From:”, followed by one or more characters (.+), followed by an at-sign. So this will match the following line: 5

`From: stephen.marquard@uct.ac.za` 6

You can think of the .+ wildcard as expanding to match all the characters between the colon character and the at-sign. 7

`From:.+@` 8

It is good to think of the plus and asterisk characters as “pushy” or “greedy.” For example, the following string would match the last at-sign in the string as the .+ pushes outwards, the one in `cwen@iupui.edu`. 9

`From: stephen.marquard@uct.ac.za, csev@umich.edu, and cwen@iupui.edu` 10

It is possible to tell an asterisk or plus sign not to be so “greedy.” If you append ?, to make it .*? or .+?, your regex will match the very first instance rather than the last instance. In the example above, using .+?@ would instead match the first at-sign, the one in `stephen.marquard@uct.ac.za`. For more details, see the [Python documentation on greedy vs non-greedy quantifiers](#). 11

11.2 Extracting data using regular expressions 12

If we want to extract data from a string in Python we can use the `.findall()` method to extract all of the substrings which match a regular expression. Let’s use the example of wanting to extract anything that looks like an email address from any line regardless of format. For example, we want to pull the email addresses from each of the following lines: 13

```
From stephen.marquard@uct.ac.za Sat Jan  5 09:14:16 2008 1
Return-Path: <postmaster@collab.sakaiproject.org>
      for <source@collab.sakaiproject.org>;
Received: (from apache@localhost)
Author: stephen.marquard@uct.ac.za
```

We don't want to write code for each of the types of lines, splitting and slicing 2 differently for each line. This following program uses `.findall()` to find the lines with email addresses in them and extract one or more addresses from each of those lines.

```
import re
s = 'A message from csev@umich.edu to cwen@iupui.edu about meeting @2PM'
lst = re.findall(r'\S+@\S+', s)
print(lst)
```

Code: <https://www.py4e.com/code3/re05.py> 4

The `.findall()` method searches the string in the second argument and returns 5 a list of all of the strings that look like email addresses. We are using a two-character sequence that matches a non-whitespace character (`\S`). You will notice we've also appended the letter `r` just before our regular expression; this tells Python to interpret our regular expression as a raw string, and not to treat backslashes as escape characters (like we use with `\n`).

The output of the program would be: 6

```
['csev@umich.edu', 'cwen@iupui.edu'] 7
```

Translating the regular expression, we are looking for substrings that have at least 8 one non-whitespace character, followed by an at-sign, followed by at least one more non-whitespace character. The `\S+` matches as many non-whitespace characters as possible.

The regular expression would match twice (csev@umich.edu and cwen@iupui.edu), 9 but it would not match the string “@2PM” because there are no non-blank characters *before* the at-sign. We can use this regular expression in a program to read all the lines in a file and print out anything that looks like an email address as follows:

```
# Search for lines that have an at sign between characters 10
import re
hand = open('mbox-short.txt') 11
for line in hand:
    line = line.rstrip()
    x = re.findall(r'\S+@\S+', line)
    if len(x) > 0:
        print(x)
```

Code: <https://www.py4e.com/code3/re06.py> 12

We read each line and then extract all the substrings that match our regular expression. Since `.findall()` returns a list, we simply check if the number of elements in our returned list is more than zero to print only lines where we found at least one substring that looks like an email address. 1

If we run the program on `mbox-short.txt` we get the following output: 2

```
...
['<source@collab.sakaiproject.org>;']
['<source@collab.sakaiproject.org>;']
['apache@localhost')
['source@collab.sakaiproject.org;']
['cwen@iupui.edu']
['source@collab.sakaiproject.org']
['cwen@iupui.edu']
['cwen@iupui.edu']
['wagnermr@iupui.edu']
```

3

Some of our email addresses have incorrect characters like “<” or “;” at the beginning or end. Let’s declare that we are only interested in the portion of the string that starts and ends with a letter or a number. 4

To do this, we use another feature of regular expressions. Square brackets are used to indicate a set of multiple acceptable characters we are willing to consider matching. In a sense, the `\S` is asking to match the set of “non-whitespace characters”. Now we will be a little more explicit in terms of the characters we will match. 5

Here is our new regular expression: 6

```
[a-zA-Z0-9]\S*@\S*[a-zA-Z]
```

7

This is getting a little complicated and you can begin to see why regular expressions are their own little language unto themselves. Translating this regular expression, we are looking for substrings that start with a *single* lowercase letter, uppercase letter, or number “[a-zA-Z0-9]”, followed by zero or more non-blank characters (`\S*`), followed by an at-sign, followed by zero or more non-blank characters (`\S*`), followed by an uppercase or lowercase letter. Note that we switched from `+` to `*` to indicate zero or more non-blank characters since `[a-zA-Z0-9]` is already one non-blank character. Remember that the `*` or `+` applies to the single character immediately to the left of the plus or asterisk. 8

If we use this expression in our program, our data is much cleaner: 9

```
# Search for lines that have an at sign between characters
# The characters must be a letter or number
import re
hand = open('mbox-short.txt')
for line in hand:
    line = line.rstrip()
    x = re.findall(r'[a-zA-Z0-9]\S*@\S*[a-zA-Z]', line)
    if len(x) > 0:
        print(x)
```

10

`# Code: https://www.py4e.com/code3/re07.py` 11

```
...  
['wagnermr@iupui.edu']  
['cwen@iupui.edu']  
['postmaster@collab.sakaiproject.org']  
['200801032122.m03LMFo4005148@nakamura.uits.iupui.edu']  
['source@collab.sakaiproject.org']  
['source@collab.sakaiproject.org']  
['source@collab.sakaiproject.org']  
['apache@localhost']
```

1

Notice that on the `source@collab.sakaiproject.org` lines, our regular expression eliminated two letters at the end of the string (“>”). This is because when we append `[a-zA-Z]` to the end of our regular expression, we are demanding that whatever string the regular expression parser finds must end with a letter. So when it sees the “>” at the end of “`sakaiproject.org>`;” it simply stops at the last “matching” letter it found (i.e., the “g” was the last good match). 2

Also note that the output of the program is a Python list that has a string as the single element in the list. 3

11.3 Combining searching and extracting⁴

If we want to find numbers on lines that start with the string “X-” such as: 5

```
X-DSPAM-Confidence: 0.8475 6  
X-DSPAM-Probability: 0.0000
```

we don’t just want any floating-point numbers from any lines. We only want to 7
extract numbers from lines that have the above syntax.

We can construct the following regular expression to select the lines: 8

```
^X-.*: [0-9.]+ 9
```

Translating this, we are saying, we want lines that start with X-, followed by zero or more characters `(.*)`, followed by a colon `(:)` and then a space. After the space we are looking for one or more characters that are either a digit `(0-9)` or a period `[0-9.]`. Note that inside the square brackets, the period matches an actual period (i.e., it is not a wildcard between the square brackets). 10

This is a very tight expression that will pretty much match only the lines we are interested in as follows: 11

```
# Search for lines that start with 'X' followed by any non  
# whitespace characters and ':'  
# followed by a space and any number.  
# The number can include a decimal.  
import re  
hand = open('mbox-short.txt')  
for line in hand:
```

12

```
line = line.rstrip() 1
if re.search(r'^X-.*: [0-9.]+', line):
    print(line)
```

Code: <https://www.py4e.com/code3/re10.py> 2

When we run the program, we see the data nicely filtered to show only the lines 3 we are looking for.

```
X-DSPAM-Confidence: 0.8475 4
X-DSPAM-Probability: 0.0000
X-DSPAM-Confidence: 0.6178
X-DSPAM-Probability: 0.0000
...
```

But now we have to solve the problem of extracting the numbers. While it would 5 be simple enough to use `split`, we can use another feature of regular expressions to both search and parse the line at the same time.

Parentheses are another special character in regular expressions. When you add 6 parentheses to a regular expression, they are ignored when matching the string. But when you are using `findall()`, parentheses indicate that while you want the whole expression to match, you only are interested in extracting a portion of the substring that matches the regular expression.

So we make the following change to our program: 7

```
# Search for lines that start with 'X' followed by any 8
# non whitespace characters and ':' followed by a space
# and any number. The number can include a decimal.
# Then print the number if it is greater than zero.
import re
hand = open('mbox-short.txt')
for line in hand:
    line = line.rstrip()
    x = re.findall(r'^X-.*: ([0-9.]+)', line)
    if len(x) > 0:
        print(x)
```

Code: <https://www.py4e.com/code3/re11.py> 9

Instead of calling `search()`, we add parentheses around the part of the regular 10 expression that represents the floating-point number to indicate we only want `findall()` to give us back the floating-point number portion of the matching string.

The output from this program is as follows: 11

```
['0.8475'] 12
['0.0000']
['0.6178']
```

```
['0.0000'] 1
['0.6961']
['0.0000']
...
```

The numbers are still in a list and need to be converted from strings to floating point, but we have used the power of regular expressions to both search and extract the information we found interesting. 2

As another example of this technique, if you look at the file there are a number of 3 lines of the form:

```
Details: http://source.sakaiproject.org/viewsvn/?view=rev&rev=39772 4
```

If we wanted to extract all of the revision numbers (the integer number at the end 5 of these lines) using the same technique as above, we could write the following program:

```
# Search for lines that start with 'Details: rev=' 6
# followed by numbers
# Then print the number if one is found
import re
hand = open('mbox-short.txt')
for line in hand:
    line = line.rstrip()
    x = re.findall(r'^Details:.*rev=([0-9]+)$', line)
    if len(x) > 0:
        print(x)
```

```
# Code: https://www.py4e.com/code3/re12.py 7
```

Translating our regular expression, we are looking for lines that start with 8 Details:, followed by any number of characters (.*), followed by rev=, and then by one or more digits. We want to find lines that match the entire expression but we only want to extract the integer number at the end of the line, so we surround [0-9]+ with parentheses, and we add \$ to indicate matches specifically at the end of the line.

When we run the program, we get the following output: 9

```
['39772'] 10
['39771']
['39770']
['39769']
...
```

Remember that the [0-9]+ is “greedy” and it tries to make as large a string of 11 digits as possible before extracting those digits. This “greedy” behavior is why we get all five digits for each number. The regular expression module expands in both directions until it encounters a non-digit, or the beginning or the end of a line.

Now we can use regular expressions to redo an exercise from earlier in the book where we were interested in the time of day of each mail message. We looked for lines of the form:

```
From stephen.marquard@uct.ac.za Sat Jan 5 09:14:16 2008 2
```

and wanted to extract the hour of the day for each line. Previously we did this with two calls to `split`. First the line was split into words and then we pulled out the fifth word and split it again on the colon character to pull out the two characters we were interested in.

While this worked, it actually results in pretty brittle code that is assuming the lines are nicely formatted. If you were to add enough error checking (or a big `try/except` block) to ensure that your program never failed when presented with incorrectly formatted lines, the code would balloon to 10-15 lines of code that was pretty hard to read.

We can do this in a far simpler way with the following regular expression:

```
^From .* [0-9][0-9]: 6
```

The translation of this regular expression is that we are looking for lines that start with `From` (note the space), followed by any number of characters (`.*`), followed by a space, followed by two digits `[0-9][0-9]`, followed by a colon character. This is the definition of the kinds of lines we are looking for.

In order to pull out only the hour using `.findall()`, we add parentheses around the two digits as follows:

```
^From .* ([0-9][0-9]): 9
```

This results in the following program:

```
# Search for lines that start with From and a character
# followed by a two digit number between 00 and 99 followed by ':'
# Then print the number if one is found
import re
hand = open('mbox-short.txt')
for line in hand:
    line = line.rstrip()
    x = re.findall('^From .* ([0-9][0-9]):', line)
    if len(x) > 0: print(x)
```

`# Code: https://www.py4e.com/code3/re13.py`

When the program runs, it produces the following output:

```
['09'] 14
['18']
['16']
['15']
...
```

11.4 Escape character¹

Regular expressions utilize special characters like `^` to match the beginning of a² line, `$` for the end of a line, and `.` as a wildcard; however, sometimes we want to match those characters literally. We need a way to indicate that we want to match the actual character such as a caret symbol, dollar sign, or period.

We can indicate that we want to simply match a character by prefixing that character with a backslash. For example, we can find money amounts with the following regular expression.³

```
import re
x = 'We just received $10.00 for cookies.'
y = re.findall('\$[0-9.]+',x)4
```

Since we prefix the dollar sign with a backslash, it actually matches the dollar sign⁵ in the input string instead of matching the “end of line”, and the rest of the regular expression matches one or more digits or the period character. Remember, as we saw above, inside square brackets, characters are not “special”. So when we say `[0-9.]`, it really means digits or a period. Outside of square brackets, a period is the “wild-card” character and matches any character. Inside square brackets, the period is a period.

11.5 Summary⁶

While this only scratched the surface of regular expressions, we have learned a bit⁷ about the language of regular expressions. They are search strings with special characters in them that communicate your wishes to the regular expression system as to what defines “matching” and what is extracted from the matched strings. Here are some of those special characters and character sequences:

<code>^</code> Matches the beginning of the line ⁸
<code>\$</code> Matches the end of the line
<code>.</code> Matches any character (a wildcard)
<code>\s</code> Matches a whitespace character
<code>\S</code> Matches a non-whitespace character (opposite of <code>\s</code>)
<code>*</code> Applies to the immediately preceding character(s) and indicates to match zero or more times.
<code>*?</code> Applies to the immediately preceding character(s) and indicates to match zero or more times in “non-greedy mode”.
<code>+</code> Applies to the immediately preceding character(s) and indicates to match one or more times.
<code>+?</code> Applies to the immediately preceding character(s) and indicates to match one or more times in “non-greedy mode”.

? Applies to the immediately preceding character(s) and indicates to match zero or one time. 1

?? Applies to the immediately preceding character(s) and indicates to match zero or one time in “non-greedy mode”. 1

[aeiou] Matches a single character as long as that character is in the specified set. In this example, it would match “a”, “e”, “i”, “o”, or “u”, but no other characters. 2

[a-zA-Z0-9] You can specify ranges of characters using the minus sign. This example is a single character that must be a lowercase letter or a digit. 3

[^A-Za-z] When the first character in the set notation is a caret, it inverts the logic. This example matches a single character that is anything *other than* an uppercase or lowercase letter. 4

() When parentheses are added to a regular expression, they are ignored for the purpose of matching, but allow you to extract a particular subset of the matched string rather than the whole string when using `.findall()`. 5

\b Matches the empty string, but only at the start or end of a word.

\B Matches the empty string, but not at the start or end of a word.

\d Matches any decimal digit; equivalent to the set [0-9].

\D Matches any non-digit character; equivalent to the set [^0-9].

11.6 Bonus section for Unix / Linux users 6

Support for searching files using regular expressions was built into the Unix operating system since the 1960s and it is available in nearly all programming languages in one form or another. 7

As a matter of fact, there is a command-line program built into Unix called `grep` (Generalized Regular Expression Parser) that does pretty much the same as the `search()` examples in this chapter. So if you have a Macintosh or Linux system, you can try the following commands in your command-line window. 8

```
$ grep '^From:' mbox-short.txt 9
From: stephen.marquard@uct.ac.za
From: louis@media.berkeley.edu
From: zqian@umich.edu
From: rjlowe@iupui.edu
```

This tells `grep` to show you lines that start with the string “From:” in the file `mbox-short.txt`. If you experiment with the `grep` command a bit and read the documentation for `grep`, you will find some subtle differences between the regular expression support in Python and the regular expression support in `grep`. As an example, `grep` does not support the non-blank character `\S` so you will need to use the slightly more complex set notation `[^]`, which simply means match a character that is anything *other than* a space. 10

11.7 Debugging ¹

Python has some simple and rudimentary built-in documentation that can be quite ² helpful if you need a quick refresher to trigger your memory about the exact name of a particular method. This documentation can be viewed in the Python interpreter in interactive mode.

You can bring up an interactive help system using `help()`. ³

```
>>> help() 4
```

```
help> modules 5
```

If you know what module you want to use, you can use the `dir()` command to ⁶ find the methods in the module as follows:

```
>>> import re 7
>>> dir(re)
[.. 'compile', 'copy_reg', 'error', 'escape', 'findall',
'finditer', 'match', 'purge', 'search', 'split', 'sre_compile',
'sre_parse', 'sub', 'subn', 'sys', 'template']
```

You can also get a small amount of documentation on a particular method using ⁸ the help command combined with the desired method.

```
>>> help (re.search) 9
Help on function search in module re:
```

```
search(pattern, string, flags=0) 10
    Scan through string looking for a match to the pattern, returning
    a match object, or None if no match was found.
>>>
```

The built-in documentation is not very extensive, but it can be helpful when you ¹¹ are in a hurry or don't have access to a web browser or search engine.

11.8 Glossary ¹²

brittle code Code that works when the input data is in a particular format but ¹³ is prone to breakage if there is some deviation from the correct format. We call this “brittle code” because it is easily broken.

greedy matching The notion that the `+` and `*` characters in a regular expression ¹⁴ expand outward to match the largest possible string.

grep A command available in most Unix systems that searches through text files ¹⁵ looking for lines that match regular expressions. The command name stands for “Generalized Regular Expression Parser”.

regular expression A language for expressing more complex search strings. A regular expression may contain special characters that indicate that a search only matches at the beginning or end of a line or many other similar capabilities. 1

wild card A special character that matches any character. In regular expressions the wild-card character is the period. 2

11.9 Exercises 3

Exercise 1: Write a simple program to simulate the operation of the `grep` command on Unix. Ask the user to enter a regular expression and count the number of lines that matched the regular expression: 4

```
$ python grep.py 5
Enter a regular expression: ^Author6
mbox.txt had 1798 lines that matched ^Author7
```

```
$ python grep.py 8
Enter a regular expression: ^X-9
mbox.txt had 14368 lines that matched ^X-10
```

```
$ python grep.py 11
Enter a regular expression: java$ 12
mbox.txt had 4175 lines that matched java$ 13
```

Exercise 2: Write a program to look for lines of the form: 14

```
New Revision: 39772 15
```

Extract the number from each of the lines using a regular expression and the `findall()` method. Compute the average of the numbers and print out the average as an integer. 16

```
Enter file:mbox.txt 17
38549 18
```

```
Enter file:mbox-short.txt 19
39756 20
```


Chapter 12¹

Networked programs²

While many of the examples in this book have focused on reading files and looking for data in those files, there are many different sources of information when one considers the Internet.³

In this chapter we will pretend to be a web browser and retrieve web pages using⁴ the Hypertext Transfer Protocol (HTTP). Then we will read through the web page data and parse it.

12.1 Hypertext Transfer Protocol - HTTP⁵

The network protocol that powers the web is actually quite simple and there is⁶ built-in support in Python called `socket` which makes it very easy to make network connections and retrieve data over those sockets in a Python program.

A `socket` is much like a file, except that a single socket provides a two-way connection between two programs. You can both read from and write to the same socket. If you write something to a socket, it is sent to the application at the other end of the socket. If you read from the socket, you are given the data which the other application has sent.⁷

But if you try to read a socket¹ when the program on the other end of the socket⁸ has not sent any data, you just sit and wait. If the programs on both ends of the socket simply wait for some data without sending anything, they will wait for a very long time, so an important part of programs that communicate over the Internet is to have some sort of protocol.⁹

A protocol is a set of precise rules that determine who is to go first, what they are to do, and then what the responses are to that message, and who sends next, and so on. In a sense the two applications at either end of the socket are doing a dance and making sure not to step on each other's toes.¹⁰

There are many documents that describe these network protocols. The Hypertext Transfer Protocol is described in the following document:¹⁰

¹If you want to learn more about sockets, protocols or how web servers are developed, you can explore the course at <https://www.dj4e.com>.

<https://www.w3.org/Protocols/rfc2616/rfc2616.txt> 1

This is a long and complex 176-page document with a lot of detail. If you find it interesting, feel free to read it all. But if you take a look around page 36 of RFC2616 you will find the syntax for the GET request. To request a document from a web server, we make a connection, e.g. to the www.pr4e.org server on port 80, and then send a line of the form 2

GET <http://data.pr4e.org/romeo.txt> HTTP/1.0 3

where the second parameter is the web page we are requesting, and then we also 4 send a blank line. The web server will respond with some header information about the document and a blank line followed by the document content.

12.2 The world's simplest web browser⁵

Perhaps the easiest way to show how the HTTP protocol works is to write a very simple Python program that makes a connection to a web server and follows the rules of the HTTP protocol to request a document and display what the server sends back. 6

```
import socket 7

mysock = socket.socket(socket.AF_INET, socket.SOCK_STREAM) 8
mysock.connect(('data.pr4e.org', 80))
cmd = 'GET http://data.pr4e.org/romeo.txt HTTP/1.0\r\n\r\n'.encode()
mysock.send(cmd)

while True: 9
    data = mysock.recv(512)
    if len(data) < 1:
        break
    print(data.decode(), end=' ')
mysock.close()
```


Code: <https://www.py4e.com/code3/socket1.py> 10

First the program makes a connection to port 80 on the server www.pr4e.com. 11 Since our program is playing the role of the “web browser”, the HTTP protocol says we must send the GET command followed by a blank line. \r\n signifies an EOL (end of line), so \r\n\r\n signifies nothing between two EOL sequences. That is the equivalent of a blank line.

Once we send that blank line, we write a loop that receives data in 512-character 12 chunks from the socket and prints the data out until there is no more data to read (i.e., the recv() returns an empty string).

The program produces the following output: 13

HTTP/1.1 200 OK 14

Figure 12.1: A Socket Connection 2

```
Date: Wed, 11 Apr 2018 18:52:55 GMT
Server: Apache/2.4.7 (Ubuntu)
Last-Modified: Sat, 13 May 2017 11:22:22 GMT
ETag: "a7-54f6609245537"
Accept-Ranges: bytes
Content-Length: 167
Cache-Control: max-age=0, no-cache, no-store, must-revalidate
Pragma: no-cache
Expires: Wed, 11 Jan 1984 05:00:00 GMT
Connection: close
Content-Type: text/plain
```

But soft what light through yonder window breaks 4
 It is the east and Juliet is the sun
 Arise fair sun and kill the envious moon
 Who is already sick and pale with grief

The output starts with headers which the web server sends to describe the document. For example, the Content-Type header indicates that the document is a plain text document (`text/plain`). 5

After the server sends us the headers, it adds a blank line to indicate the end of the headers, and then sends the actual data of the file `romeo.txt`. 6

This example shows how to make a low-level network connection with sockets. 7
 Sockets can be used to communicate with a web server or with a mail server or many other kinds of servers. All that is needed is to find the document which describes the protocol and write the code to send and receive the data according to the protocol.

However, since the protocol that we use most commonly is the HTTP web protocol, Python has a special library specifically designed to support the HTTP protocol for the retrieval of documents and data over the web. 8

One of the requirements for using the HTTP protocol is the need to send and receive data as bytes objects, instead of strings. In the preceding example, the `encode()` and `decode()` methods convert strings into bytes objects and back again. 9

The next example uses `b''` notation to specify that a variable should be stored as a bytes object. `encode()` and `b''` are equivalent.

```
>>> b'Hello world'          2
b'Hello world'
>>> 'Hello world'.encode()
b'Hello world'
```

12.3 Retrieving an image over HTTP³

In the above example, we retrieved a plain text file which had newlines in the file ⁴ and we simply copied the data to the screen as the program ran. We can use a similar program to retrieve an image using HTTP. Instead of copying the data to the screen as the program runs, we accumulate the data in a string, trim off the headers, and then save the image data to a file as follows:

```
import socket
import time

HOST = 'data.pr4e.org'
PORT = 80
mysock = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
mysock.connect((HOST, PORT))
mysock.sendall(b'GET http://data.pr4e.org/cover3.jpg HTTP/1.0\r\n\r\n')
count = 0
picture = b""

while True:
    data = mysock.recv(5120)
    if len(data) < 1: break
    #time.sleep(0.25)
    count = count + len(data)
    print(len(data), count)
    picture = picture + data

mysock.close()

# Look for the end of the header (2 CRLF) 6
pos = picture.find(b"\r\n\r\n")
print('Header length', pos)
print(picture[:pos].decode())

# Skip past the header and save the picture data 7
picture = picture[pos+4:]
fhand = open("stuff.jpg", "wb")
fhand.write(picture)
fhand.close()

# Code: https://www.py4e.com/code3/urljpeg.py 8
```

When the program runs, it produces the following output:1

```
$ python urljpeg.py2
5120 5120
5120 10240
4240 14480
5120 19600
...
5120 214000
3200 217200
5120 222320
5120 227440
3167 230607
Header length 393
HTTP/1.1 200 OK
Date: Wed, 11 Apr 2018 18:54:09 GMT
Server: Apache/2.4.7 (Ubuntu)
Last-Modified: Mon, 15 May 2017 12:27:40 GMT
ETag: "38342-54f8f2e5b6277"
Accept-Ranges: bytes
Content-Length: 230210
Vary: Accept-Encoding
Cache-Control: max-age=0, no-cache, no-store, must-revalidate
Pragma: no-cache
Expires: Wed, 11 Jan 1984 05:00:00 GMT
Connection: close
Content-Type: image/jpeg
```

You can see that for this url, the Content-Type header indicates that body of the3 document is an image (image/jpeg). Once the program completes, you can view the image data by opening the file **stuff.jpg** in an image viewer.

As the program runs, you can see that we don't get 5120 characters each time4 we call the `recv()` method. We get as many characters as have been transferred across the network to us by the web server at the moment we call `recv()`. In this example, we either get as few as 3200 characters each time we request up to 5120 characters of data.

Your results may be different depending on your network speed. Also note that on5 the last call to `recv()` we get 3167 bytes, which is the end of the stream, and in the next call to `recv()` we get a zero-length string that tells us that the server has called `close()` on its end of the socket and there is no more data forthcoming.

We can slow down our successive `recv()` calls by uncommenting the call to6 `time.sleep()`. This way, we wait a quarter of a second after each call so that the server can “get ahead” of us and send more data to us before we call `recv()` again. With the delay, in place the program executes as follows:

```
$ python urljpeg.py7
5120 5120
5120 10240
5120 15360
...
5120 225280
```

```

5120 230400
207 230607
Header length 393
HTTP/1.1 200 OK
Date: Wed, 11 Apr 2018 21:42:08 GMT
Server: Apache/2.4.7 (Ubuntu)
Last-Modified: Mon, 15 May 2017 12:27:40 GMT
ETag: "38342-54f8f2e5b6277"
Accept-Ranges: bytes
Content-Length: 230210
Vary: Accept-Encoding
Cache-Control: max-age=0, no-cache, no-store, must-revalidate
Pragma: no-cache
Expires: Wed, 11 Jan 1984 05:00:00 GMT
Connection: close
Content-Type: image/jpeg

```

1

Now other than the first and last calls to `recv()`, we now get 5120 characters each 2 time we ask for new data.

There is a buffer between the server making `send()` requests and our application 3 making `recv()` requests. When we run the program with the delay in place, at some point the server might fill up the buffer in the socket and be forced to pause until our program starts to empty the buffer. The pausing of either the sending application or the receiving application is called “flow control.”

12.4 Retrieving web pages with `urllib`⁴

While we can manually send and receive data over HTTP using the socket library, 5 there is a much simpler way to perform this common task in Python by using the `urllib` library.

Using `urllib`, you can treat a web page much like a file. You simply indicate 6 which web page you would like to retrieve and `urllib` handles all of the HTTP protocol and header details.

The equivalent code to read the *romeo.txt* file from the web using `urllib` is as 7 follows:

```

import urllib.request 8

fhand = urllib.request.urlopen('http://data.pr4e.org/romeo.txt') 9
for line in fhand:
    print(line.decode().strip())

```

Code: https://www.py4e.com/code3/urllib1.py 10

Once the web page has been opened with `urllib.request.urlopen`, we can treat 11 it like a file and read through it using a `for` loop.

When the program runs, we only see the output of the contents of the file. The 12 headers are still sent, but the `urllib` code consumes the headers and only returns the data to us.

```
But soft what light through yonder window breaks 1
It is the east and Juliet is the sun
Arise fair sun and kill the envious moon
Who is already sick and pale with grief
```

As an example, we can write a program to retrieve the data for romeo.txt and 2 compute the frequency of each word in the file as follows:

```
import urllib.request, urllib.parse, urllib.error 3

fhand = urllib.request.urlopen('http://data.pr4e.org/romeo.txt') 4

counts = dict() 5
for line in fhand:
    words = line.decode().split()
    for word in words:
        counts[word] = counts.get(word, 0) + 1
print(counts)

# Code: https://www.py4e.com/code3/urlwords.py 6
```

Again, once we have opened the web page, we can read it like a local file. 7

12.5 Reading binary files using urllib⁸

Sometimes you want to retrieve a non-text (or binary) file such as an image or 9 video file. The data in these files is generally not useful to print out, but you can easily make a copy of a URL to a local file on your hard disk using `urllib`.

The pattern is to open the URL and use `read` to download the entire contents of 10 the document into a string variable (`img`) then write that information to a local file as follows:

```
import urllib.request, urllib.parse, urllib.error 11

img = urllib.request.urlopen('http://data.pr4e.org/cover3.jpg').read() 12
fhand = open('cover3.jpg', 'wb')
fhand.write(img)
fhand.close()

# Code: https://www.py4e.com/code3/curl1.py 13
```

This program reads all of the data in at once across the network and stores it in the 14 variable `img` in the main memory of your computer, then opens the file `cover.jpg` and writes the data out to your disk. The `wb` argument for `open()` opens a binary file for writing only. This program will work if the size of the file is less than the size of the memory of your computer.

However if this is a large audio or video file, this program may crash or at least 15 run extremely slowly when your computer runs out of memory. In order to avoid

running out of memory, we retrieve the data in blocks (or buffers) and then write each block to your disk before retrieving the next block. This way the program can read any size file without using up all of the memory you have in your computer. 1

```
import urllib.request, urllib.parse, urllib.error 2

img = urllib.request.urlopen('http://data.pr4e.org/cover3.jpg') 3
fhand = open('cover3.jpg', 'wb')
size = 0
while True:
    info = img.read(100000)
    if len(info) < 1: break
    size = size + len(info)
    fhand.write(info)

print(size, 'characters copied.')
fhand.close()
```

Code: [https://www.py4e.com/code3\(curl2.py](https://www.py4e.com/code3(curl2.py) 4

In this example, we read only 100,000 characters at a time and then write those characters to the cover3.jpg file before retrieving the next 100,000 characters of data from the web. 5

This program runs as follows: 6

```
python curl2.py 7
230210 characters copied. 8
```

12.6 Parsing HTML and scraping the web 9

One of the common uses of the `urllib` capability in Python is to *scrape* the web. 10
Web scraping is when we write a program that pretends to be a web browser and retrieves pages, then examines the data in those pages looking for patterns.

As an example, a search engine such as Google will look at the source of one web page and extract the links to other pages and retrieve those pages, extracting links, and so on. Using this technique, Google *spiders* its way through nearly all of the pages on the web. 11

Google also uses the frequency of links from pages it finds to a particular page as one measure of how “important” a page is and how high the page should appear in its search results. 12

12.7 Parsing HTML using regular expressions 13

One simple way to parse HTML is to use regular expressions to repeatedly search for and extract substrings that match a particular pattern. 14

Here is a simple web page: 15

```
<h1>The First Page</h1>
<p>
If you like, you can switch to the
<a href="http://www.dr-chuck.com/page2.htm">
Second Page</a>.
</p>
```

1

We can construct a well-formed regular expression to match and extract the link values from the above text as follows:

`href="http[s]?://.+?"` 3

Our regular expression looks for strings that start with “`href="http://"` or “`href="https://"`, followed by one or more characters `(.+?)`, followed by another double quote. The question mark behind the `[s]?` indicates to search for the string “`http`” followed by zero or one “`s`”.

The question mark added to the `.+?` indicates that the match is to be done in a “non-greedy” fashion instead of a “greedy” fashion. A non-greedy match tries to find the *smallest* possible matching string and a greedy match tries to find the *largest* possible matching string.

We add parentheses to our regular expression to indicate which part of our matched string we would like to extract, and produce the following program:

```
# Search for link values within URL input
import urllib.request, urllib.parse, urllib.error
import re
import ssl

# Ignore SSL certificate errors
ctx = ssl.create_default_context()
ctx.check_hostname = False
ctx.verify_mode = ssl.CERT_NONE

url = input('Enter - ')
html = urllib.request.urlopen(url, context=ctx).read()
links = re.findall(b'href="(http[s]?://.*?)"', html)
for link in links:
    print(link.decode())
```

7

Code: <https://www.py4e.com/code3/urlregex.py> 8

The `ssl` library allows this program to access web sites that strictly enforce HTTPS.⁹ The `read` method returns HTML source code as a bytes object instead of returning an `HTTPResponse` object. The `findall` regular expression method will give us a list of all of the strings that match our regular expression, returning only the link text between the double quotes.

When we run the program and input a URL, we get the following output:¹⁰

```
Enter - https://docs.python.org
https://docs.python.org/3/index.html
https://www.python.org/
https://docs.python.org/3.8/
https://docs.python.org/3.7/
https://docs.python.org/3.5/
https://docs.python.org/2.7/
https://www.python.org/doc/versions/
https://www.python.org/dev/peps/
https://wiki.python.org/moin/BeginnersGuide
https://wiki.python.org/moin/PythonBooks
https://www.python.org/doc/av/
https://www.python.org/
https://www.python.org/psf/donations/
http://sphinx.pocoo.org/
```

1

Regular expressions work very nicely when your HTML is well formatted and predictable. But since there are a lot of “broken” HTML pages out there, a solution only using regular expressions might either miss some valid links or end up with bad data.

This can be solved by using a robust HTML parsing library.³

12.8 Parsing HTML using BeautifulSoup⁴

Even though HTML looks like XML² and some pages are carefully constructed to be XML, most HTML is generally broken in ways that cause an XML parser to reject the entire page of HTML as improperly formed.

There are a number of Python libraries which can help you parse HTML and extract data from the pages. Each of the libraries has its strengths and weaknesses and you can pick one based on your needs.

As an example, we will simply parse some HTML input and extract links using the *BeautifulSoup* library. BeautifulSoup tolerates highly flawed HTML and still lets you easily extract the data you need. You can download and install the BeautifulSoup code from:

<https://pypi.python.org/pypi/beautifulsoup4>⁸

Information on installing BeautifulSoup with the Python Package Index tool pip⁹ is available at:

<https://packaging.python.org/tutorials/installing-packages/>¹⁰

We will use `urllib` to read the page and then use `BeautifulSoup` to extract the `href` attributes from the anchor (`a`) tags.

```
# To run this, download the BeautifulSoup zip file
# http://www.py4e.com/code3/bs4.zip
# or pip install beautifulsoup4 to ensure you have the latest version
# and unzip it in the same directory as this file
```

12

²The XML format is described in the next chapter.

```
import urllib.request, urllib.parse, urllib.error
from bs4 import BeautifulSoup
import ssl # defaults to certificate verification and most secure protocol (r

# Ignore SSL/TLS certificate errors 2
ctx = ssl.create_default_context()
ctx.check_hostname = False
ctx.verify_mode = ssl.CERT_NONE

url = input('Enter - ')
html = urllib.request.urlopen(url, context=ctx).read()
soup = BeautifulSoup(html, 'html.parser')

# Retrieve all of the anchor tags 3
tags = soup('a')
for tag in tags:
    print(tag.get('href', None))

# Code: https://www.py4e.com/code3/urllinks.py 4
```

The program prompts for a web address, then opens the web page, reads the data 5 and passes the data to the BeautifulSoup parser, and then retrieves all of the anchor tags and prints out the `href` attribute for each tag.

When the program runs, it produces the following output: 6

```
Enter - https://docs.python.org 7
genindex.html
py-modindex.html
https://www.python.org/
#
whatsnew/3.6.html
whatsnew/index.html
tutorial/index.html
library/index.html
reference/index.html
using/index.html
howto/index.html
installing/index.html
distributing/index.html
extending/index.html
c-api/index.html
faq/index.html
py-modindex.html
genindex.html
glossary.html
search.html
contents.html
bugs.html
about.html
license.html
copyright.html
```

```

download.html
https://docs.python.org/3.8/
https://docs.python.org/3.7/
https://docs.python.org/3.5/
https://docs.python.org/2.7/
https://www.python.org/doc/versions/
https://www.python.org/dev/peps/
https://wiki.python.org/moin/BeginnersGuide
https://wiki.python.org/moin/PythonBooks
https://www.python.org/doc/av/
genindex.html
py-modindex.html
https://www.python.org/
#
copyright.html
https://www.python.org/psf/donations/
bugs.html
http://sphinx.pocoo.org/

```

1

This list is much longer because some HTML anchor tags are relative paths (e.g., `tutorial/index.html`) or in-page references (e.g., `'#'`) that do not include “`http://`” or “`https://`”, which was a requirement in our regular expression.

You can use also BeautifulSoup to pull out various parts of each tag:³

```

# To run this, download the BeautifulSoup zip file 4
# http://www.py4e.com/code3/bs4.zip
# and unzip it in the same directory as this file

```

```

from urllib.request import urlopen 5
from bs4 import BeautifulSoup
import ssl

```

```

# Ignore SSL certificate errors 6
ctx = ssl.create_default_context()
ctx.check_hostname = False
ctx.verify_mode = ssl.CERT_NONE

url = input('Enter - ')
html = urlopen(url, context=ctx).read()
soup = BeautifulSoup(html, "html.parser")

```

```

# Retrieve all of the anchor tags 7
tags = soup('a')
for tag in tags:
    # Look at the parts of a tag
    print('TAG:', tag)
    print('URL:', tag.get('href', None))
    print('Contents:', tag.contents[0])
    print('Attrs:', tag.attrs)

```

```
# Code: https://www.py4e.com/code3/urllink2.py 8
```

```
python urllink2.py
Enter - http://www.dr-chuck.com/page1.htm
TAG: <a href="http://www.dr-chuck.com/page2.htm">
Second Page</a>
URL: http://www.dr-chuck.com/page2.htm
Content: ['\nSecond Page']
Attrs: [('href', 'http://www.dr-chuck.com/page2.htm')]
```

1

`html.parser` is the HTML parser included in the standard Python 3 library. Information on other HTML parsers is available at:

<http://www.crummy.com/software/BeautifulSoup/bs4/doc/#installing-a-parser>

These examples only begin to show the power of BeautifulSoup when it comes to parsing HTML.

12.9 Bonus section for Unix / Linux users⁵

If you have a Linux, Unix, or Macintosh computer, you probably have commands built in to your operating system that retrieves both plain text and binary files using the HTTP or File Transfer (FTP) protocols. One of these commands is `curl`:

```
$ curl -O http://www.py4e.com/cover.jpg
```

7

The command `curl` is short for “copy URL” and so the two examples listed earlier to retrieve binary files with `urllib` are cleverly named `curl1.py` and `curl2.py` on www.py4e.com/code3 as they implement similar functionality to the `curl` command. There is also a `curl3.py` sample program that does this task a little more effectively, in case you actually want to use this pattern in a program you are writing.

A second command that functions very similarly is `wget`:

```
$ wget http://www.py4e.com/cover.jpg
```

10

Both of these commands make retrieving webpages and remote files a simple task.

12.10 Glossary¹²

BeautifulSoup A Python library for parsing HTML documents and extracting data from HTML documents that compensates for most of the imperfections in the HTML that browsers generally ignore. You can download the BeautifulSoup code from www.crummy.com.

port A number that generally indicates which application you are contacting when you make a socket connection to a server. As an example, web traffic usually uses port 80 while email traffic uses port 25.

13

14

scrape When a program pretends to be a web browser and retrieves a web page,
then looks at the web page content. Often programs are following the links
in one page to find the next page so they can traverse a network of pages or
a social network.

socket A network connection between two applications where the applications can
send and receive data in either direction.

spider The act of a web search engine retrieving a page and then all the pages
linked from a page and so on until they have nearly all of the pages on the
Internet which they use to build their search index.

12.11 Exercises⁴

Exercise 1: Change the socket program `socket1.py` to prompt the user for the
URL so it can read any web page.⁵

You can use `split('/')` to break the URL into its component parts so you can
extract the host name for the socket `connect` call. Add error checking using `try`
and `except` to handle the condition where the user enters an improperly formatted
or non-existent URL.

Exercise 2: Change your socket program so that it counts the number of characters it
has received and stops displaying any text after it has shown 3000 characters.
The program should retrieve the entire document and count the total number of
characters and display the count of the number of characters at the end of the
document.⁷

Exercise 3: Use `urllib` to replicate the previous exercise of (1) retrieving the
document from a URL, (2) displaying up to 3000 characters, and (3) counting the
overall number of characters in the document. Don't worry about the headers for
this exercise, simply show the first 3000 characters of the document contents.⁸

Exercise 4: Change the `urllinks.py` program to extract and count paragraph (`p`)
tags from the retrieved HTML document and display the count of the paragraphs
as the output of your program. Do not display the paragraph text, only count
them. Test your program on several small web pages as well as some larger web
pages.⁹

Exercise 5: (Advanced) Change the socket program so that it only shows data
after the headers and a blank line have been received. Remember that `recv` receives
characters (newlines and all), not lines.¹⁰

Chapter 13¹

Using Web Services²

Once it became easy to retrieve documents and parse documents over HTTP using ³ programs, it did not take long to develop an approach where we started producing documents that were specifically designed to be consumed by other programs (i.e., not HTML to be displayed in a browser).

There are two common formats that we use when exchanging data across the web.⁴ eXtensible Markup Language (XML) has been in use for a very long time and is best suited for exchanging document-style data. When programs just want to exchange dictionaries, lists, or other internal information with each other, they use JavaScript Object Notation (JSON) (see www.json.org). We will look at both formats.

13.1 eXtensible Markup Language - XML⁵


XML looks very similar to HTML, but XML is more structured than HTML. Here ⁶ is a sample of an XML document:

```
<person>
  <name>Chuck</name>
  <phone type="intl">
    +1 734 303 4456
  </phone>
  <email hide="yes" />
</person>
```

⁷

Each pair of opening (e.g., `<person>`) and closing tags (e.g., `</person>`) represents ⁸ an *element* or *node* with the same name as the tag (e.g., `person`). Each element can have some text, some attributes (e.g., `hide`), and other nested elements. If an XML element is empty (i.e., has no content), then it may be depicted by a self-closing tag (e.g., `<email />`).

Often it is helpful to think of an XML document as a tree structure where there is ⁹ a top element (here: `person`), and other tags (e.g., `phone`) are drawn as *children* of their *parent* elements.

Figure 13.1: A Tree Representation of XML 2

13.2 Parsing XML 3

Here is a simple application that parses some XML and extracts some data elements 4 from the XML:

```

import xml.etree.ElementTree as ET 5

data = '''
<person>
    <name>Chuck</name>
    <phone type="intl">
        +1 734 303 4456
    </phone>
    <email hide="yes" />
</person>'''

tree = ET.fromstring(data)
print('Name:', tree.find('name').text)
print('Attr:', tree.find('email').get('hide'))
  
```

Code: <https://www.py4e.com/code3/xml1.py> 6

The triple single quote (''''), as well as the triple double quote (""""), allow for 7 the creation of strings that span multiple lines.

Calling `fromstring` converts the string representation of the XML into a “tree” of 8 XML elements. When the XML is in a tree, we have a series of methods we can call to extract portions of data from the XML string. The `find` function searches through the XML tree and retrieves the element that matches the specified tag.

Name: Chuck 9

Attr: yes 10

Using an XML parser such as `ElementTree` has the advantage that while the 11 XML in this example is quite simple, it turns out there are many rules regarding

valid XML, and using `ElementTree` allows us to extract data from XML without worrying about the rules of XML syntax. 1

13.3 Looping through nodes²

Often the XML has multiple nodes and we need to write a loop to process all of the nodes. In the following program, we loop through all of the `user` nodes: 3

```
import xml.etree.ElementTree as ET 4

input = '''
<stuff>
    <users>
        <user x="2">
            <id>001</id>
            <name>Chuck</name>
        </user>
        <user x="7">
            <id>009</id>
            <name>Brent</name>
        </user>
    </users>
</stuff>''' 5

stuff = ET.fromstring(input) 6
lst = stuff.findall('users/user')
print('User count:', len(lst))

for item in lst: 7
    print('Name', item.find('name').text)
    print('Id', item.find('id').text)
    print('Attribute', item.get('x'))

# Code: https://www.py4e.com/code3/xml2.py 8
```

The `findall` method retrieves a Python list of subtrees that represent the `user` structures in the XML tree. Then we can write a `for` loop that looks at each of the user nodes, and prints the `name` and `id` text elements as well as the `x` attribute from the `user` node. 9

```
User count: 2 10
Name Chuck
Id 001
Attribute 2
Name Brent
Id 009
Attribute 7
```

It is important to include all parent level elements in the `.findall` statement except **1** for the top level element (e.g., `users/user`). Otherwise, Python will not find any desired nodes.

```
import xml.etree.ElementTree as ET 2

3
input = '''
<stuff>
  <users>
    <user x="2">
      <id>001</id>
      <name>Chuck</name>
    </user>
    <user x="7">
      <id>009</id>
      <name>Brent</name>
    </user>
  </users>
</stuff>'''

4
stuff = ET.fromstring(input)

lst = stuff.findall('users/user')
print('User count:', len(lst))

lst2 = stuff.findall('user')
print('User count:', len(lst2))
```

1 `lst` stores all `user` elements that are nested within their `users` parent. **2** `lst2` looks for `user` elements that are not nested within the top level `stuff` element where there are none.

```
User count: 2 5
User count: 0
```

13.4 JavaScript Object Notation - JSON **6**

The JSON format was inspired by the object and array format used in the **7** JavaScript language. But since Python was invented before JavaScript, Python's syntax for dictionaries and lists influenced the syntax of JSON. So the format of JSON is nearly identical to a combination of Python lists and dictionaries.

Here is a JSON encoding that is roughly equivalent to the simple XML from above: **8**

```
{ 9
  "name" : "Chuck",
  "phone" : {
    "type" : "intl",
    "number" : "+1 734 303 4456"
```

```

    },
    "email" : {
        "hide" : "yes"
    }
}

```

You will notice some differences. First, in XML, we can add attributes like “`intl`” 2 to the “`phone`” tag. In JSON, we simply have key-value pairs. Also the XML “`person`” tag is gone, replaced by a set of outer curly braces.

In general, JSON structures are simpler than XML because JSON has fewer capabilities than XML. But JSON has the advantage that it maps *directly* to some combination of dictionaries and lists. And since nearly all programming languages have something equivalent to Python’s dictionaries and lists, JSON is a very natural format to have two cooperating programs exchange data.

JSON is quickly becoming the format of choice for nearly all data exchange between 3 applications because of its relative simplicity compared to XML.

13.5 Parsing JSON 5

We construct our JSON by nesting dictionaries and lists as needed. In this example, 6 we represent a list of users where each user is a set of key-value pairs (i.e., a dictionary). So we have a list of dictionaries.

In the following program, we use the built-in `json` library to parse the JSON and 7 read through the data. Compare this closely to the equivalent XML data and code above. The JSON has less detail, so we must know in advance that we are getting a list and that the list is of users and each user is a set of key-value pairs. The JSON is more succinct (an advantage) but also is less self-describing (a disadvantage).

```

import json
data = '''
[
    { "id" : "001",
      "x" : "2",
      "name" : "Chuck"
    },
    { "id" : "009",
      "x" : "7",
      "name" : "Brent"
    }
]'''

info = json.loads(data)
print('User count:', len(info))

for item in info:
    print('Name', item['name'])

```

```
print('Id', item['id']) 1
print('Attribute', item['x'])
```

Code: <https://www.py4e.com/code3/json2.py> 2

If you compare the code to extract data from the parsed JSON and XML you will 3 see that what we get from `json.loads()` is a Python list which we traverse with a `for` loop, and each item within that list is a Python dictionary. Once the JSON has been parsed, we can use the Python index operator to extract the various bits of data for each user. We don't have to use the JSON library to dig through the parsed JSON, since the returned data is simply native Python structures.

The output of this program is exactly the same as the XML version above. 4

```
User count: 2 5
Name Chuck
Id 001
Attribute 2
Name Brent
Id 009
Attribute 7
```

In general, there is an industry trend away from XML and towards JSON for web 6 services. Because the JSON is simpler and more directly maps to native data structures we already have in programming languages, the parsing and data extraction code is usually simpler and more direct when using JSON. But XML is more self-descriptive than JSON and so there are some applications where XML retains an advantage. For example, most word processors store documents internally using XML rather than JSON.

13.6 Application Programming Interfaces 7

We now have the ability to exchange data between applications using Hypertext Transport Protocol (HTTP) and a way to represent complex data that we are sending back and forth between these applications using eXtensible Markup Language (XML) or JavaScript Object Notation (JSON). 8

The next step is to begin to define and document “contracts” between applications 9 using these techniques. The general name for these application-to-application contracts is *Application Program Interfaces* (APIs). When we use an API, generally one program makes a set of *services* available for use by other applications and publishes the APIs (i.e., the “rules”) that must be followed to access the services provided by the program.

When we begin to build our programs where the functionality of our program includes access to services provided by other programs, we call the approach a *Service-oriented architecture* (SOA). An SOA approach is one where our overall application makes use of the services of other applications. A non-SOA approach is where the application is a single standalone application which contains all of the code necessary to implement the application. 10




Figure 13.2: Service-oriented architecture 2

We see many examples of SOA when we use the web. We can go to a single web site and book air travel, hotels, and automobiles all from a single site. The data for hotels is not stored on the airline computers. Instead, the airline computers contact the services on the hotel computers and retrieve the hotel data and present it to the user. When the user agrees to make a hotel reservation using the airline site, the airline site uses another web service on the hotel systems to actually make the reservation. And when it comes time to charge your credit card for the whole transaction, still other computers become involved in the process.

A Service-oriented architecture has many advantages, including: (1) we always maintain only one copy of data (this is particularly important for things like hotel reservations where we do not want to over-commit) and (2) the owners of the data can set the rules about the use of their data. With these advantages, an SOA system must be carefully designed to have good performance and meet the user's needs.

When an application makes a set of services in its API available over the web, we call these *web services*.

13.7 Security and API usage⁶

It is quite common that you need an API key to make use of a vendor's API. The general idea is that they want to know who is using their services and how much each user is using. Perhaps they have free and pay tiers of their services or have a policy that limits the number of requests that a single individual can make during a particular time period.

Sometimes once you get your API key, you simply include the key as part of POST data or perhaps as a parameter on the URL when calling the API.

Other times, the vendor wants increased assurance of the source of the requests **1** and so they expect you to send cryptographically signed messages using shared keys and secrets. A very common technology that is used to sign requests over the Internet is called *OAuth*. You can read more about the OAuth protocol at www.oauth.net.

Thankfully there are a number of convenient and free OAuth libraries so you can **2** avoid writing an OAuth implementation from scratch by reading the specification. These libraries are of varying complexity and have varying degrees of richness. The OAuth web site has information about various OAuth libraries.

13.8 Glossary³

API Application Program Interface - A contract between applications that defines **4** the patterns of interaction between two application components.

ElementTree A built-in Python library used to parse XML data. **5**

JSON JavaScript Object Notation - A format that allows for the markup of structured data based on the syntax of JavaScript Objects. **6**

SOA Service-Oriented Architecture - When an application is made of components **7** connected across a network.

XML eXtensible Markup Language - A format that allows for the markup of structured data. **8**

Chapter 14¹

Object-oriented² programming

14.1 Managing larger programs³

At the beginning of this book, we came up with four basic programming patterns⁴ which we use to construct programs:

- Sequential code
- Conditional code (if statements)
- Repetitive code (loops)
- Store and reuse (functions)

⁵

In later chapters, we explored simple variables as well as collection data structures⁶ like lists, tuples, and dictionaries.

As we build programs, we design data structures and write code to manipulate⁷ those data structures. There are many ways to write programs and by now, you probably have written some programs that are “not so elegant” and other programs that are “more elegant”. Even though your programs may be small, you are starting to see how there is a bit of art and aesthetic to writing code.

As programs get to be millions of lines long, it becomes increasingly important to⁸ write code that is easy to understand. If you are working on a million-line program, you can never keep the entire program in your mind at the same time. We need ways to break large programs into multiple smaller pieces so that we have less to look at when solving a problem, fix a bug, or add a new feature.

In a way, object oriented programming is a way to arrange your code so that you can zoom into 50 lines of the code and understand it while ignoring the other 999,950 lines of code for the moment.⁹

14.2 Getting started¹

Like many aspects of programming, it is necessary to learn the concepts of object oriented programming before you can use them effectively. You should approach this chapter as a way to learn some terms and concepts and work through a few simple examples to lay a foundation for future learning.²

The key outcome of this chapter is to have a basic understanding of how objects are constructed and how they function and most importantly how we make use of the capabilities of objects that are provided to us by Python and Python libraries.³

14.3 Using objects⁴

As it turns out, we have been using objects all along in this book. Python provides us with many built-in objects. Here is some simple code where the first few lines should feel very simple and natural to you.⁵

```
stuff = list()          6
stuff.append('python')
stuff.append('chuck')
stuff.sort()
print (stuff[0])
print (stuff.__getitem__(0))
print (list.__getitem__(stuff,0))
```

*# Code: <https://www.py4e.com/code3/objects.py>*⁷

Instead of focusing on what these lines accomplish, let's look at what is really happening from the point of view of object-oriented programming. Don't worry if the following paragraphs don't make any sense the first time you read them because we have not yet defined all of these terms.⁸

The first line *constructs* an object of type `list`, the second and third lines *call*⁹ the `append()` *method*, the fourth line calls the `sort()` method, and the fifth line *retrieves* the item at position 0.

The sixth line calls the `__getitem__()` method in the `stuff` list with a parameter¹⁰ of zero.

```
print (stuff.__getitem__(0))11
```

The seventh line is an even more verbose way of retrieving the 0th item in the list.¹²

```
| print (list.__getitem__(stuff,0))13
```

In this code, we call the `__getitem__` method in the `list` class and *pass* the list¹⁴ and the item we want retrieved from the list as parameters.

The last three lines of the program are equivalent, but it is more convenient to 1 simply use the square bracket syntax to look up an item at a particular position in a list.

We can take a look at the capabilities of an object by looking at the output of the 2 `dir()` function:

```
>>> stuff = list()
>>> dir(stuff)
['__add__', '__class__', '__contains__', '__delattr__',
 '__delitem__', '__dir__', '__doc__', '__eq__',
 '__format__', '__ge__', '__getattribute__', '__getitem__',
 '__gt__', '__hash__', '__iadd__', '__imul__', '__init__',
 '__iter__', '__le__', '__len__', '__lt__', '__mul__',
 '__ne__', '__new__', '__reduce__', '__reduce_ex__',
 '__repr__', '__reversed__', '__rmul__', '__setattr__',
 '__setitem__', '__sizeof__', '__str__', '__subclasshook__',
 'append', 'clear', 'copy', 'count', 'extend', 'index',
 'insert', 'pop', 'remove', 'reverse', 'sort']
>>>
```

3

The rest of this chapter will define all of the above terms so make sure to come 4 back after you finish the chapter and re-read the above paragraphs to check your understanding.

14.4 Starting with programs⁵

A program in its most basic form takes some input, does some processing, and 6 produces some output. Our elevator conversion program demonstrates a very short but complete program showing all three of these steps.

```
usf = input('Enter the US Floor Number: ')7
wf = int(usf) - 1
print('Non-US Floor Number is',wf)
```

Code: <https://www.py4e.com/code3/elev.py>⁸

If we think a bit more about this program, there is the “outside world” and the 9 program. The input and output aspects are where the program interacts with the outside world. Within the program we have code and data to accomplish the task the program is designed to solve.

One way to think about object-oriented programming is that it separates our program into multiple “zones.” Each zone contains some code and data (like a program) and has well defined interactions with the outside world and the other zones within the program.¹⁰

If we look back at the link extraction application where we used the BeautifulSoup library, we can see a program that is constructed by connecting different objects together to accomplish a task:¹¹

Figure 14.1: A Program 2

```

# To run this, download the BeautifulSoup zip file 3
# http://www.py4e.com/code3/bs4.zip
# or pip install beautifulsoup4 to ensure you have the latest version
# and unzip it in the same directory as this file

```

```

import urllib.request, urllib.parse, urllib.error
from bs4 import BeautifulSoup
import ssl # defaults to certificate verification and most secure protocol (now TLS)

```

```

# Ignore SSL/TLS certificate errors 5
ctx = ssl.create_default_context()
ctx.check_hostname = False
ctx.verify_mode = ssl.CERT_NONE

url = input('Enter - ')
html = urllib.request.urlopen(url, context=ctx).read()
soup = BeautifulSoup(html, 'html.parser')

```

```

# Retrieve all of the anchor tags 6
tags = soup('a')
for tag in tags:
    print(tag.get('href', None))

```

```
# Code: https://www.py4e.com/code3/urllinks.py 7
```

We read the URL into a string and then pass that into `urllib` to retrieve the data 8 from the web. The `urllib` library uses the `socket` library to make the actual network connection to retrieve the data. We take the string that `urllib` returns and hand it to `BeautifulSoup` for parsing. `BeautifulSoup` makes use of the object `html.parser`¹ and returns an object. We call the `tags()` method on the returned object that returns a dictionary of tag objects. We loop through the tags and call the `get()` method for each tag to print out the `href` attribute.

We can draw a picture of this program and how the objects work together. 9

The key here is not to understand perfectly how this program works but to see 10 how we build a network of interacting objects and orchestrate the movement of information between the objects to create a program. It is also important to note that when you looked at that program several chapters back, you could fully

¹<https://docs.python.org/3/library/html.parser.html>




Figure 14.2: A Program as Network of Objects 2




Figure 14.3: Ignoring Detail When Using an Object 4

understand what was going on in the program without even realizing that the program was “orchestrating the movement of data between objects.” It was just lines of code that got the job done.

14.5 Subdividing a problem⁶

One of the advantages of the object-oriented approach is that it can hide complexity.⁷ For example, while we need to know how to use the `urllib` and `BeautifulSoup` code, we do not need to know how those libraries work internally. This allows us to focus on the part of the problem we need to solve and ignore the other parts of the program.

This ability to focus exclusively on the part of a program that we care about and ignore the rest is also helpful to the developers of the objects that we use. For example, the programmers developing `BeautifulSoup` do not need to know or care about how we retrieve our HTML page, what parts we want to read, or what we plan to do with the data we extract from the web page.

14.6 Our first Python object⁹

At a basic level, an object is simply some code plus data structures that are smaller than a whole program. Defining a function allows us to store a bit of code and give it a name and then later invoke that code using the name of the function.




Figure 14.4: Ignoring Detail When Building an Object 2

An object can contain a number of functions (which we call *methods*) as well as 3 data that is used by those functions. We call data items that are part of the object *attributes*.

We use the `class` keyword to define the data and code that will make up each 4 of the objects. The class keyword includes the name of the class and begins an indented block of code where we include the attributes (data) and methods (code).

```

class PartyAnimal:5

    def __init__(self):6
        self.x = 0

    def party(self) :
        self.x = self.x + 1
        print("So far",self.x)

an = PartyAnimal()
an.party()
an.party()
an.party()
  
```

Code: <https://www.py4e.com/code3/party2.py> 7

Each method looks like a function, starting with the `def` keyword and consisting 8 of an indented block of code.

The first method is a specially-named method called `__init__()`. This method is 9 called to do any initial setup of the data we want to store in the object. In this class we allocate the `x` attribute using dot notation and initialize it to zero.

`self.x = 0` 10

The other method named `party`. The methods all have a special first parameter 11 that we name by convention `self`. The first parameter gives us access to the object instance so we can set attributes and call methods using dot notation.

Just as the `def` keyword does not cause function code to be executed, the `class` 12 keyword does not create an object. Instead, the `class` keyword defines a template

Figure 14.5: A Class and Two Objects ²

indicating what data and code will be contained in each object of type `PartyAnimal`. ³
The class is like a cookie cutter and the objects created using the class are the cookies². You don't put frosting on the cookie cutter; you put frosting on the cookies, and you can put different frosting on each cookie.

If we continue through this sample program, we see the first executable line of ⁴ code:

```
an = PartyAnimal()5
```

This is where we instruct Python to construct (i.e., create) an *object* or *instance* ⁶ of the class `PartyAnimal`. It looks like a function call to the class itself. Python constructs the object with the right data and methods and returns the object which is then assigned to the variable `an`. In a way this is quite similar to the following line which we have been using all along:

```
counts = dict()7
```

Here we instruct Python to construct an object using the `dict` template (already ⁸ present in Python), return the instance of dictionary, and assign it to the variable `counts`.

When the `PartyAnimal` class is used to construct an object, the variable `an` is used ⁹ to point to that object. We use `an` to access the code and data for that particular instance of the `PartyAnimal` class.

Each `Partyanimal` object/instance contains within it a variable `x` and a ¹⁰ method/function named `party`. We call the `party` method in this line:

```
an.party()11
```

When the `party` method is called, the first parameter (which we call by convention `self`) points to the particular instance of the `PartyAnimal` object that `party` is called from. Within the `party` method, we see the line: ¹²

²Cookie image copyright CC-BY <https://www.flickr.com/photos/dinnerseries/23570475099>

```
self.x = self.x + 1
```

This syntax using the *dot* operator is saying ‘the x within self.’ Each time `party()`² is called, the internal x value is incremented by 1 and the value is printed out.

The following line is another way to call the `party` method within the `an` object:³

```
PartyAnimal.party(an)
```

In this variation, we access the code from within the class and explicitly pass the `an` object pointer as the first parameter (i.e., `self` within the method). You can think of `an.party()` as shorthand for the above line.⁵

When the program executes, it produces the following output:⁶

```
So far 1
So far 2
So far 3
So far 4
```

The object is constructed, and the `party` method is called four times, both incrementing and printing the value for x within the `an` object.⁸

14.7 Classes as types⁹

As we have seen, in Python all variables have a type. We can use the built-in `dir`¹⁰ function to examine the capabilities of a variable. We can also use `type` and `dir` with the classes that we create.

```
class PartyAnimal:11

    def __init__(self):12
        self.x = 0

    def party(self) :
        self.x = self.x + 1
        print("So far",self.x)

an = PartyAnimal()
print ("Type", type(an))
print ("Dir ", dir(an))
print ("Type", type(an.x))
print ("Type", type(an.party))
```

Code: <https://www.py4e.com/code3/party3.py>¹³

When this program executes, it produces the following output:¹⁴

```
Type <class '__main__.PartyAnimal'>          1
Dir  ['__class__', '__delattr__', ...
      '__sizeof__', '__str__', '__subclasshook__',
      '__weakref__', 'party', 'x']
Type <class 'int'>
Type <class 'method'>
```

You can see that using the `class` keyword, we have created a new type. From the 2 `dir` output, you can see both the `x` integer attribute and the `party` method are available in the object.

14.8 Object lifecycle³

In the previous examples, we define a class (template), use that class to create 4 an instance of that class (object), and then use the instance. When the program finishes, all of the variables are discarded. Usually, we don't think much about the creation and destruction of variables, but often as our objects become more complex, we need to take some action within the object to set things up as the object is constructed and possibly clean things up as the object is discarded.

If we want our object to be aware of these moments of construction and destruction, 5 we add specially named methods to our object:

```
class PartyAnimal: 6

    def __init__(self): 7
        self.x = 0
        print('I am constructed')

    def party(self) :
        self.x = self.x + 1
        print('So far',self.x)

    def __del__(self):
        print('I am destructed', self.x)

an = PartyAnimal()
an.party()
an.party()
an = 42
print('an contains',an)
```

8 # Code: <https://www.py4e.com/code3/party4.py>

When this program executes, it produces the following output: 9

```
I am constructed 10
So far 1
So far 2
I am destructed 2
an contains 42
```

As Python constructs our object, it calls our `__init__` method to give us a chance **1** to set up some default or initial values for the object. When Python encounters the line:

```
an = 42 2
```

It actually “throws our object away” so it can reuse the `an` variable to store the **3** value `42`. Just at the moment when our `an` object is being “destroyed” our destructor code (`__del__`) is called. We cannot stop our variable from being destroyed, but we can do any necessary cleanup right before our object no longer exists.

When developing objects, it is quite common to add a constructor to an object to **4** set up initial values for the object. It is relatively rare to need a destructor for an object.

14.9 Multiple instances **5**

So far, we have defined a class, constructed a single object, used that object, **6** and then thrown the object away. However, the real power in object-oriented programming happens when we construct multiple instances of our class.

When we construct multiple objects from our class, we might want to set up different initial values for each of the objects. We can pass data to the constructors **7** to give each object a different initial value:

```
class PartyAnimal:8

    def __init__(self, nam):9
        self.x = 0
        self.name = nam
        print(self.name, 'constructed')

    def party(self) :
        self.x = self.x + 1
        print(self.name, 'party count', self.x)

s = PartyAnimal('Sally')
s.party()
j = PartyAnimal('Jim')

j.party()
s.party()
```

*# Code: <https://www.py4e.com/code3/party5.py> **10***

The constructor has both a `self` parameter that points to the object instance **11** and additional parameters that are passed into the constructor as the object is constructed:

```
s = PartyAnimal('Sally') 12
```

Within the constructor, the second line copies the parameter (`nam`) that is passed 1 into the `name` attribute within the object instance.

```
self.name = nam 2
```

The output of the program shows that each of the objects (`s` and `j`) contain their 3 own independent copies of `x` and `nam`:

```
Sally constructed 4
Sally party count 1
Jim constructed
Jim party count 1
Sally party count 2
```

14.10 Inheritance 5

Another powerful feature of object-oriented programming is the ability to create 6 a new class by extending an existing class. When extending a class, we call the original class the *parent class* and the new class the *child class*.

For this example, we move our `PartyAnimal` class into its own file. Then, we can 7 ‘import’ the `PartyAnimal` class in a new file and extend it, as follows:

```
from party import PartyAnimal 8

class CricketFan(PartyAnimal): 9

    def __init__(self, nam) :
        super().__init__(nam)
        self.points = 0 10

    def six(self):
        self.points = self.points + 6
        self.party()
        print(self.name,"points",self.points)

s = PartyAnimal("Sally")
s.party()
j = CricketFan("Jim")
j.party()
j.six()
print(dir(j))

# Code: https://www.py4e.com/code3/party6.py 11
```

When we define the `CricketFan` class, we indicate that we are extending the `PartyAnimal` class. This means that all of the variables (`x`) and methods (`party`) from the `PartyAnimal` class are *inherited* by the `CricketFan` class. For example, 12

within the `six` method in the `CricketFan` class, we call the `party` method from 1
the `PartyAnimal` class.

We use a special syntax in the `__init__()` method in the `CricketFan` class to 2
ensure that we call the `__init__()` method in the `PartyAnimal` so that whatever
setup that `PartyAnimal` needs is done in addition to the setup needed for the
`CricketFan` extensions.

```
def __init__(self, nam) :3
    super().__init__(nam)
    self.points = 0
```

The `super()` syntax is telling Python to call the `__init__` method in the class 4
that we are extending. `PartyAnimal` is the super (or parent) class and `CricketFan`
is the sub (or child) class.

As the program executes, we create `s` and `j` as independent instances of 5
`PartyAnimal` and `CricketFan`. The `j` object has additional capabilities beyond
the `s` object.

```
Sally constructed6
Sally party count 1
Jim constructed
Jim party count 1
Jim party count 2
Jim points 6
['__class__', '__delattr__', ... '__weakref__',
'name', 'party', 'points', 'six', 'x']
```

In the `dir` output for the `j` object (instance of the `CricketFan` class), we see that 7
it has the attributes and methods of the parent class, as well as the attributes and
methods that were added when the class was extended to create the `CricketFan`
class.

14.11 Summary 8

This is a very quick introduction to object-oriented programming that focuses 9
mainly on terminology and the syntax of defining and using objects. Let's quickly
review the code that we looked at in the beginning of the chapter. At this point
you should fully understand what is going on.

```
stuff = list()10
stuff.append('python')
stuff.append('chuck')
stuff.sort()
print (stuff[0])
print (stuff.__getitem__(0))
print (list.__getitem__(stuff,0))
```

Code: <https://www.py4e.com/code3/objects.py> 11

The first line constructs a `list object`. When Python creates the `list` object, it calls the `constructor` method (named `__init__`) to set up the internal data attributes that will be used to store the list data. We have not passed any parameters to the `constructor`. When the constructor returns, we use the variable `stuff` to point to the returned instance of the `list` class. 1

The second and third lines call the `append` method with one parameter to add a new item at the end of the list by updating the attributes within `stuff`. Then in the fourth line, we call the `sort` method with no parameters to sort the data within the `stuff` object. 2

We then print out the first item in the list using the square brackets which are a shortcut to calling the `__getitem__` method within the `stuff`. This is equivalent to calling the `__getitem__` method in the `list class` and passing the `stuff` object as the first parameter and the position we are looking for as the second parameter. 3

At the end of the program, the `stuff` object is discarded but not before calling the `destructor` (named `__del__`) so that the object can clean up any loose ends as necessary. 4

Those are the basics of object-oriented programming. There are many additional details as to how to best use object-oriented approaches when developing large applications and libraries that are beyond the scope of this chapter.³ 5

14.12 Glossary⁶

attribute A variable that is part of a class. 7

class A template that can be used to construct an object. Defines the attributes and methods that will make up the object. 8

child class A new class created when a parent class is extended. The child class inherits all of the attributes and methods of the parent class. 9

constructor An optional specially named method (`__init__`) that is called at the moment when a class is being used to construct an object. Usually this is used to set up initial values for the object. 10

destructor An optional specially named method (`__del__`) that is called at the moment just before an object is destroyed. Destructors are rarely used. 11

inheritance When we create a new class (child) by extending an existing class (parent). The child class has all the attributes and methods of the parent class plus additional attributes and methods defined by the child class. 12

method A function that is contained within a class and the objects that are constructed from the class. Some object-oriented patterns use ‘message’ instead of ‘method’ to describe this concept. 13

object A constructed instance of a class. An object contains all of the attributes and methods that were defined by the class. Some object-oriented documentation uses the term ‘instance’ interchangeably with ‘object’. 14

³If you are curious about where the `list` class is defined, take a look at (hopefully the URL won’t change) <https://github.com/python/cpython/blob/master/Objects/listobject.c> - the `list` class is written in a language called “C”. If you take a look at that source code and find it curious you might want to explore C programming with an eye towards building objects at <https://www.cc4e.com/>.

parent class The class which is being extended to create a new child class. The parent class contributes all of its methods and attributes to the new child class. 1

Chapter 15¹

Using Databases and SQL²

15.1 What is a database?³

A *database* is a file that is organized for storing data. Most databases are organized⁴ like a dictionary in the sense that they map from keys to values. The biggest difference is that the database is on disk (or other permanent storage), so it persists after the program ends. Because a database is stored on permanent storage, it can store far more data than a dictionary, which is limited to the size of the memory in the computer.

Like a dictionary, database software is designed to keep the inserting and accessing⁵ of data very fast, even for large amounts of data. Database software maintains its performance by building *indexes* as data is added to the database to allow the computer to jump quickly to a particular entry.

There are many different database systems which are used for a wide variety of purposes⁶ including: Oracle, MySQL, Microsoft SQL Server, PostgreSQL, and SQLite. We focus on SQLite in this book because it is a very common database and is already built into Python. SQLite is designed to be *embedded* into other applications to provide database support within the application. For example, the Firefox browser also uses the SQLite database internally as do many other products.

<http://sqlite.org/>⁷

SQLite is well suited to some of the data manipulation problems that we see in⁸ Informatics.

15.2 Database concepts⁹

When you first look at a database it looks like a spreadsheet with multiple sheets.¹⁰ The primary data structures in a database are: *tables*, *rows*, and *columns*.

In technical descriptions of relational databases the concepts of table, row, and¹¹ column are more formally referred to as *relation*, *tuple*, and *attribute*, respectively. We will use the less formal terms in this chapter.

The diagram illustrates the relationship between a table and a relation. On the left, a table is shown with three columns labeled "column" and one row labeled "row". The table contains three data cells: the top-left cell is labeled "Table", the middle-right cell is "2.3", and the bottom-right cell is empty. On the right, a relation is shown with three columns labeled "attribute" and one row labeled "tuple". The relation also has three rows of data, corresponding to the table's rows.

Figure 15.1: Relational Databases 2

15.3 Database Browser for SQLite 3

While this chapter will focus on using Python to work with data in SQLite database files, many operations can be done more conveniently using software called the *Database Browser for SQLite* which is freely available from:

<http://sqlitebrowser.org/> 5

Using the browser you can easily create tables, insert data, edit data, or run simple SQL queries on the data in the database. 6

In a sense, the database browser is similar to a text editor when working with text files. When you want to do one or very few operations on a text file, you can just open it in a text editor and make the changes you want. When you have many changes that you need to do to a text file, often you will write a simple Python program. You will find the same pattern when working with databases. You will do simple operations in the database manager and more complex operations will be most conveniently done in Python. 7

15.4 Creating a database table 8

Databases require more defined structure than Python lists or dictionaries¹. 9

When we create a database *table* we must tell the database in advance the names of each of the *columns* in the table and the type of data which we are planning to store in each *column*. When the database software knows the type of data in each column, it can choose the most efficient way to store and look up the data based on the type of data. 10

You can look at the various data types supported by SQLite at the following url: 11

<http://www.sqlite.org/datatypes.html> 12

Defining structure for your data up front may seem inconvenient at the beginning, but the payoff is fast access to your data even when the database contains a large amount of data. 13

¹SQLite actually does allow some flexibility in the type of data stored in a column, but we will keep our data types strict in this chapter so the concepts apply equally to other database systems such as MySQL.




Figure 15.2: A Database Cursor 2

The code to create a database file and a table named `Track` with two columns in 3 the database is as follows:

```
import sqlite3 4

conn = sqlite3.connect('music.sqlite')5
cur = conn.cursor()6

cur.execute('DROP TABLE IF EXISTS Track')7
cur.execute('CREATE TABLE Track (title TEXT, plays INTEGER)')8

conn.close()9

# Code: https://www.py4e.com/code3/db1.py 10
```

The `connect` operation makes a “connection” to the database stored in the file 11 `music.sqlite` in the current directory. If the file does not exist, it will be created. The reason this is called a “connection” is that sometimes the database is stored on a separate “database server” from the server on which we are running our application. In our simple examples the database will just be a local file in the same directory as the Python code we are running.

A `cursor` is like a file handle that we can use to perform operations on the data 12 stored in the database. Calling `cursor()` is very similar conceptually to calling `open()` when dealing with text files.

Once we have the cursor, we can begin to execute commands on the contents of 13 the database using the `execute()` method.

Database commands are expressed in a special language that has been standardized 14 across many different database vendors to allow us to learn a single database language. The database language is called *Structured Query Language* or *SQL* for short.

<http://en.wikipedia.org/wiki/SQL> 15

In our example, we are executing two SQL commands in our database. As a 16 convention, we will show the SQL keywords in uppercase and the parts of the

command that we are adding (such as the table and column names) will be shown 1 in lowercase.

The first SQL command removes the `Track` table from the database if it exists.2 This pattern is simply to allow us to run the same program to create the `Track` table over and over again without causing an error. Note that the `DROP TABLE` command deletes the table and all of its contents from the database (i.e., there is no “undo”).

```
cur.execute('DROP TABLE IF EXISTS Track')3
```

The second command creates a table named `Track` with a text column named 4 `title` and an integer column named `plays`.

```
cur.execute('CREATE TABLE Track (title TEXT, plays INTEGER)')5
```

Now that we have created a table named `Track`, we can put some data into that 6 table using the SQL `INSERT` operation. Again, we begin by making a connection to the database and obtaining the `cursor`. We can then execute SQL commands using the `cursor`.

The SQL `INSERT` command indicates which table we are using and then defines a 7 new row by listing the fields we want to include (`title`, `plays`) followed by the `VALUES` we want placed in the new row. We specify the values as question marks `(?, ?)` to indicate that the actual values are passed in as a tuple `('My Way', 15)` as the second parameter to the `execute()` call.

```
import sqlite38

conn = sqlite3.connect('music.sqlite')
cur = conn.cursor()

cur.execute('INSERT INTO Track (title, plays) VALUES (?, ?)', 
           ('Thunderstruck', 20))
cur.execute('INSERT INTO Track (title, plays) VALUES (?, ?)', 
           ('My Way', 15))
conn.commit()

print('Track:')
cur.execute('SELECT title, plays FROM Track')
for row in cur:
    print(row)

cur.execute('DELETE FROM Track WHERE plays < 100')
conn.commit()

cur.close()

# Code: https://www.py4e.com/code3/db2.py
```

Tracks¹

title	plays	²
Thunderstruck	20	
My Way	15	

Figure 15.3: Rows in a Table³

First we `INSERT` two rows into our table and use `commit()` to force the data to be⁴ written to the database file.

Then we use the `SELECT` command to retrieve the rows we just inserted from the⁵ table. On the `SELECT` command, we indicate which columns we would like (`title`, `plays`) and indicate which table we want to retrieve the data from. After we execute the `SELECT` statement, the cursor is something we can loop through in a `for` statement. For efficiency, the cursor does not read all of the data from the database when we execute the `SELECT` statement. Instead, the data is read on demand as we loop through the rows in the `for` statement.

The output of the program is as follows:⁶

```
Track:7
('Thunderstruck', 20)8
('My Way', 15)9
```

Our `for` loop finds two rows, and each row is a Python tuple with the first value¹⁰ as the `title` and the second value as the number of `plays`.

At the very end of the program, we execute an SQL command to `DELETE` the¹¹ rows we have just created so we can run the program over and over. The `DELETE` command shows the use of a `WHERE` clause that allows us to express a selection criterion so that we can ask the database to apply the command to only the rows that match the criterion. In this example the criterion happens to apply to all the rows so we empty the table out so we can run the program repeatedly. After the `DELETE` is performed, we also call `commit()` to force the data to be removed from the database.

15.5 Structured Query Language summary¹²

So far, we have been using the Structured Query Language in our Python examples¹³ and have covered many of the basics of the SQL commands. In this section, we look at the SQL language in particular and give an overview of SQL syntax.

Since there are so many different database vendors, the Structured Query Language¹⁴ (SQL) was standardized so we could communicate in a portable manner to database systems from multiple vendors.

A relational database is made up of tables, rows, and columns. The columns generally have a type such as text, numeric, or date data. When we create a table, we indicate the names and types of the columns:1

```
CREATE TABLE Track (title TEXT, plays INTEGER)2
```

To insert a row into a table, we use the SQL **INSERT** command:3

```
INSERT INTO Track (title, plays) VALUES ('My Way', 15)4
```

The **INSERT** statement specifies the table name, then a list of the fields/columns that you would like to set in the new row, and then the keyword **VALUES** and a list of corresponding values for each of the fields.5

The SQL **SELECT** command is used to retrieve rows and columns from a database.6
The **SELECT** statement lets you specify which columns you would like to retrieve as well as a **WHERE** clause to select which rows you would like to see. It also allows an optional **ORDER BY** clause to control the sorting of the returned rows.

```
SELECT * FROM Track WHERE title = 'My Way'7
```

Using ***** indicates that you want the database to return all of the columns for each row that matches the **WHERE** clause.8

Note, unlike in Python, in a SQL **WHERE** clause we use a single equal sign to indicate a test for equality rather than a double equal sign. Other logical operations allowed in a **WHERE** clause include **<**, **>**, **<=**, **>=**, **!=**, as well as **AND** and **OR** and parentheses to build your logical expressions.9

You can request that the returned rows be sorted by one of the fields as follows:10

```
SELECT title, plays FROM Track ORDER BY title11
```

It is possible to **UPDATE** a column or columns within one or more rows in a table12 using the SQL **UPDATE** statement as follows:

```
UPDATE Track SET plays = 16 WHERE title = 'My Way'13
```

The **UPDATE** statement specifies a table and then a list of fields and values to change14 after the **SET** keyword and then an optional **WHERE** clause to select the rows that are to be updated. A single **UPDATE** statement will change all of the rows that match the **WHERE** clause. If a **WHERE** clause is not specified, it performs the **UPDATE** on all of the rows in the table.

To remove a row, you need a **WHERE** clause on an SQL **DELETE** statement. The **WHERE** clause determines which rows are to be deleted:15

```
DELETE FROM Track WHERE title = 'My Way'16
```

These four basic SQL commands (INSERT, SELECT, UPDATE, and DELETE) allow the four basic operations needed to create and maintain data. We use “CRUD” (Create, Read, Update, and Delete) to capture all these concepts in a single term.²¹

15.6 Multiple tables and basic data modeling²

The real power of a relational database is when we create multiple tables and make links between those tables. The act of deciding how to break up your application data into multiple tables and establishing the relationships between the tables is called *data modeling*. The design document that shows the tables and their relationships is called a *data model*.³

Data modeling is a relatively sophisticated skill and we will only introduce the most basic concepts of relational data modeling in this section. For more detail on data modeling you can start with:⁴

http://en.wikipedia.org/wiki/Relational_model⁵

Lets say for our tracks database we wanted to track the name of the `artist` for each track in addition to the `title` and number of plays for each track. A simple approach might be to simply add another column to the database called `artist` and put the name of the artist in the column as follows:⁶

```
DROP TABLE IF EXISTS Track;7
CREATE TABLE Track (title TEXT, plays INTEGER, artist TEXT);
```

Then we could insert a few tracks into our table.⁸

```
INSERT INTO Track (title, plays, artist)9
VALUES ('My Way', 15, 'Frank Sinatra');
INSERT INTO Track (title, plays, artist)
VALUES ('New York', 25, 'Frank Sinatra');
```

If we were to look at our data with a `SELECT * FROM Track` statement, it looks like we have done a fine job.¹⁰

```
sqlite> SELECT * FROM Track;11
My Way|15|Frank Sinatra
New York|25|Frank Sinatra
sqlite>
```

We have made a *very bad error* in our data modeling. We have violated the rules of *database normalization*.¹²

https://en.wikipedia.org/wiki/Database_normalization¹³

²Yes there is a disconnect between “CRUD” term and the first letters of the four SQL statements that implement “CRUD”. A possible explanation might be to claim that “CRUD” is the “concept” and SQL is the implementation. Another possible explanation is that “CRUD” is more fun to say than “ISUD”.

While database normalization seems very complex on the surface and contains a lot of mathematical justifications, for now we can reduce it all into one simple rule that we will follow.

We should never put the same string data in a column more than once. If we need the data more than once, we create a numeric *key* for the data and reference the actual data using this key. Especially if the multiple entries refer to the same object.

To demonstrate the slippery slope we are going down by assigning string columns to our database model, think about how we would change the data model if we wanted to keep track of the eye color of our artists? Would we do this?

```
DROP TABLE IF EXISTS Track; 4
CREATE TABLE Track (title TEXT, plays INTEGER,
    artist TEXT, eyes TEXT);
INSERT INTO Track (title, plays, artist, eyes)
    VALUES ('My Way', 15, 'Frank Sinatra', 'Blue');
INSERT INTO Track (title, plays, artist, eyes)
    VALUES ('New York', 25, 'Frank Sinatra', 'Blue');
```

Since Frank Sinatra recorded over 1200 songs, are we really going to put the string 'Blue' in 1200 rows in our **Track** table. And what would happen if we decided his eye color was 'Light Blue'? Something just does not feel right.

The correct solution is to create a table for the each **Artist** and store all the data about the artist in that table. And then somehow we need to make a connection between a row in the **Track** table to a row in the **Artist** table. Perhaps we could call this "link" between two "tables" a "relationship" between two tables. And that is exactly what database experts decided to all these links.

Lets make an **Artist** table as follows:

```
DROP TABLE IF EXISTS Artist; 8
CREATE TABLE Artist (name TEXT, eyes TEXT);
INSERT INTO Artist (name, eyes)
    VALUES ('Frank Sinatra', 'blue');
```

Now we have two tables but we need a way to *link* rows in the two tables. To do this, we need what we call 'keys'. These keys will just be integer numbers that we can use to lookup a row in different table. If we are going to make links to rows inside of a table, we need to add a *primary key* to the rows in the table. By convention we usually name the primary key column 'id'. So our **Artist** table looks as follows:

```
DROP TABLE IF EXISTS Artist; 10
CREATE TABLE Artist (id INTEGER, name TEXT, eyes TEXT);
INSERT INTO Artist (id, name, eyes)
    VALUES (42, 'Frank Sinatra', 'blue');
```

Now we have a row in the table for 'Frank Sinatra' (and his eye color) and a primary key of '42' to use to link our tracks to him. So we alter our **Track** table as follows:

```
DROP TABLE IF EXISTS Track; 1
CREATE TABLE Track (title TEXT, plays INTEGER,
    artist_id INTEGER);
INSERT INTO Track (title, plays, artist_id)
    VALUES ('My Way', 15, 42);
INSERT INTO Track (title, plays, artist_id)
    VALUES ('New York', 25, 42);
```

The `artist_id` column is an integer, and by naming convention is a *foreign key* 2 pointing at a *primary* key in the `Artist` table. We call it a foreign key because it is pointing to a row in a different table.

Now we are following the rules of database normalization, but when we want to 3 get data out of our database, we don't want to see the 42, we want to see the name and eye color of the artist. To do this we use the `JOIN` keyword in our `SELECT` statement.

```
SELECT title, plays, name, eyes 4
FROM Track JOIN Artist
ON Track.artist_id = Artist.id;
```

The `JOIN` clause includes an `ON` condition that defines how the rows are to be connected. For each row in `Track` add the data from `Artist` from the row where `artist_id` `Track` table matches the `id` from the `Artist` table. 5

The output would be: 6



```
My Way|15|Frank Sinatra|blue 7
New York|25|Frank Sinatra|blue
```

While it might seem a little clunky and your instincts might tell you that it would 8 be faster just to keep the data in one table, it turns out the limit on database performance is how much data needs to be scanned when retrieving a query. While the details are very complex, integers are a lot smaller than strings (especially Unicode) and far quicker to move and compare.

15.7 Data model diagrams 9

While our `Track` and `Artist` database design is simple with just two tables and a single one-to-many relationship, these data models can get complicated quickly and are easier to understand if we can make a graphical representation of our data model. 10

While there are many graphical representations of data models, we will use one of the “classic” approaches, called “Crow’s Foot Diagrams” as shown in Figure 15.4. Each table is shown as a box with the name of the table and its columns. Then where there is a relationship between two tables a line is drawn connecting the tables with a notation added to the end of each line indicating the nature of the relationship. 11

Figure 15.4: A Verbose One-to-Many Data Model 2Figure 15.5: A Succinct One-to-Many Data Model 4

https://en.wikipedia.org/wiki/Entity-relationship_model 5

In this case, “many” tracks can be associated with each artist. So the track end is shown with the crow’s foot spread out indicating it is the “many” end. The artist end is shown with a vertical like that indicates “one”. There will be “many” artists in general, but the important aspect is that for each artist there will be many tracks. And each of those artists may be associated with multiple tracks. 6

You will note that the column that holds the *foreign key* like `artist_id` is on the “many” end and the *primary key* is at the “one” end. 7

Since the pattern of foreign and primary key placement is so consistent and follows the “many” and “one” ends of the lines, we never include either the primary or foreign key columns in our diagram of the data model as shown in the second diagram as shown in Figure 15.5. The columns are thought of as “implementation detail” to capture the nature of the relationship details and not an essential part of the data being modeled. 8

15.8 Automatically creating primary keys 9

In the above example, we arbitrarily assigned Frank the primary key of 42. However when we are inserting millions or rows, it is nice to have the database automatically generate the values for the `id` column. We do this by declaring the `id` column as a PRIMARY KEY and leave out the `id` value when inserting the row: 10

```

DROP TABLE IF EXISTS Artist; 11
CREATE TABLE Artist (id INTEGER PRIMARY KEY,
    name TEXT, eyes TEXT);
INSERT INTO Artist (name, eyes)
    VALUES ('Frank Sinatra', 'blue');
  
```

Now we have instructed the database to auto-assign us a unique value to the Frank Sinatra row. But we then need a way to have the database tell us the `id` value for the recently inserted row. One way is to use a `SELECT` statement to retrieve data from an SQLite built-in-function called `last_insert_rowid()`. 12

```
sqlite> DROP TABLE IF EXISTS Artist; 1
sqlite> CREATE TABLE Artist (id INTEGER PRIMARY KEY,
...>     name TEXT, eyes TEXT);
sqlite> INSERT INTO Artist (name, eyes)
...>     VALUES ('Frank Sinatra', 'blue');
sqlite> select last_insert_rowid();
1
sqlite> SELECT * FROM Artist;
1|Frank Sinatra|blue
sqlite>
```

Once we know the `id` of our ‘Frank Sinatra’ row, we can use it when we `INSERT` 2 the tracks into the `Track` table. As a general strategy, we add these `id` columns to any table we create:

```
sqlite> DROP TABLE IF EXISTS Track; 3
sqlite> CREATE TABLE Track (id INTEGER PRIMARY KEY,
...>     title TEXT, plays INTEGER, artist_id INTEGER);
```

Note that the `artist_id` value is the new auto-assigned row in the `Artist` table 4 and that while we added an `INTEGER PRIMARY KEY` to the `Track` table, we did not include `id` in the list of fields on the `INSERT` statements into the `Track` table. Again this tells the database to choose a unique value for us for the `id` column.

```
sqlite> INSERT INTO Track (title, plays, artist_id) 5
...>     VALUES ('My Way', 15, 1);
sqlite> select last_insert_rowid();
1
sqlite> INSERT INTO Track (title, plays, artist_id)
...>     VALUES ('New York', 25, 1);
sqlite> select last_insert_rowid();
2
sqlite>
```

You can call `SELECT last_insert_rowid();` after each of the inserts to retrieve 6 the value that the database assigned to the `id` of each newly created row. Later when we are coding in Python, we can ask for the `id` value in our code and store it in a variable for later use.

15.9 Logical keys for fast lookup⁷

If we had a table full of artists and a table full of tracks, each with a foreign key 8 link to a row in a table full of artists and we wanted to list all the tracks that were sung by ‘Frank Sinatra’ as follows:

```
SELECT title, plays, name, eyes 9
FROM Track JOIN Artist
ON Track.artist_id = Artist.id
WHERE Artist.name = 'Frank Sinatra';
```

Since we have two tables and a foreign key between the two tables, our data is **1** well-modeled, but if we are going to have millions of records in the `Artist` table and going to do a lot of lookups by artist name, we would benefit if we gave the database a hint about our intended use of the `name` column.

We do this by adding an “index” to a text column that we intend to use in `WHERE` **2** clauses:

```
CREATE INDEX artist_name ON Artist(name); 3
```

When the database has been told that an index is needed on a column in a table, it **4** stores extra information to make it possible to look up a row more quickly using the indexed field (`name` in this example). Once you request that an index be created, there is nothing special that is needed in the SQL to access the table. The database keeps the index up to date as data is inserted, deleted, and updated, and uses it automatically if it will increase the performance of a database query.

These text columns that are used to find rows based on some information in the **5** “real world” like the name of an artist are called *Logical keys*.

15.10 Adding constraints to the database^{**6**}

We can also use an index to enforce a constraint (i.e. rules) on our database **7** operations. The most common constraint is a *uniqueness constraint* which insists that all of the values in a column are unique. We can add the optional `UNIQUE` keyword, to the `CREATE INDEX` statement to tell the database that we would like it to enforce the constraint on our SQL. We can drop and re-create the `artist_name` index with a `UNIQUE` constraint as follows.

```
DROP INDEX artist_name; 8
CREATE UNIQUE INDEX artist_name ON Artist(name);
```

If we try to insert ‘Frank Sinatra’ a second time, it will fail with an error. **9**

```
sqlite> SELECT * FROM Artist; 10
1|Frank Sinatra|blue
sqlite> INSERT INTO Artist (name, eyes)
...>     VALUES ('Frank Sinatra', 'blue');
Runtime error: UNIQUE constraint failed: Artist.name (19)
sqlite>
```

We can tell the database to ignore any duplicate key errors by adding the `IGNORE` **11** keyword to the `INSERT` statement as follows:

```
sqlite> INSERT OR IGNORE INTO Artist (name, eyes) 12
...>     VALUES ('Frank Sinatra', 'blue');
sqlite> SELECT id FROM Artist WHERE name='Frank Sinatra';
1
sqlite>
```




Figure 15.6: Tracks, Albums, and Artists 2

By combining an `INSERT OR IGNORE` and a `SELECT` we can insert a new record if the name is not already there and whether or not the record is already there, retrieve the *primary key* of the record. 3

```

sqlite> INSERT OR IGNORE INTO Artist (name, eyes) 4
...>     VALUES ('Elvis', 'blue');
sqlite> SELECT id FROM Artist WHERE name='Elvis';
2
sqlite> SELECT * FROM Artist;
1|Frank Sinatra|blue
2|Elvis|blue
sqlite>
  
```

Since we have not added a uniqueness constraint to the eye color column, there is 5 no problem having multiple ‘Blue’ values in the `eye` column. 5

15.11 Sample multi-table application 6

A sample application called `tracks_csv.py` shows how these ideas can be combined 7 to parse textual data and load it into several tables using a proper data model with relational connections between the tables. 7

This application reads and parses a comma-separated file `tracks.csv` based on an 8 export from Dr. Chuck’s iTunes library. 8

```

Another One Bites The Dust,Queen,Greatest Hits,55,100,217103 9
Asche Zu Asche,Rammstein,Herzeleid,79,100,231810
Beauty School Dropout,Various,Grease,48,100,239960
Black Dog,Led Zeppelin,IV,109,100,296620
...
  
```

The columns in this file are: title, artist, album, number of plays, rating (0-100) 10 and length in milliseconds. 10

Our data model is shown in Figure 15.6 and described in SQL as follows: 11

```

DROP TABLE IF EXISTS Artist; 12
DROP TABLE IF EXISTS Album;
DROP TABLE IF EXISTS Track;
  
```

```

CREATE TABLE Artist (
    id INTEGER PRIMARY KEY,
    name TEXT UNIQUE
);

CREATE TABLE Album (
    id INTEGER PRIMARY KEY,
    artist_id INTEGER,
    title TEXT UNIQUE
);

CREATE TABLE Track (
    id INTEGER PRIMARY KEY,
    title TEXT UNIQUE,
    album_id INTEGER,
    len INTEGER, rating INTEGER, count INTEGER
);

```

1

We are adding the `UNIQUE` keyword to `TEXT` columns that we would like to have a uniqueness constraint that we will use in `INSERT IGNORE` statements. This is more succinct than separate `CREATE INDEX` statements but has the same effect.

With these tables in place, we write the following code `tracks_csv.py` to parse the data and insert it into the tables:

```

import sqlite3

conn = sqlite3.connect('trackdb.sqlite')
cur = conn.cursor()

handle = open('tracks.csv')

for line in handle:
    line = line.strip()
    pieces = line.split(',')
    if len(pieces) != 6 : continue

    name = pieces[0]
    artist = pieces[1]
    album = pieces[2]
    count = pieces[3]
    rating = pieces[4]
    length = pieces[5]

    print(name, artist, album, count, rating, length)

    cur.execute(''':> INSERT OR IGNORE INTO Artist (name)
VALUES ( ? )''', ( artist, ) )
    cur.execute('SELECT id FROM Artist WHERE name = ? ', (artist, ))
    artist_id = cur.fetchone()[0]

```

4

2

3

```

cur.execute(''INSERT OR IGNORE INTO Album (title, artist_id)
            VALUES ( ?, ? )'', ( album, artist_id ) )
cur.execute('SELECT id FROM Album WHERE title = ? ', (album, ))
album_id = cur.fetchone()[0]

cur.execute(''INSERT OR REPLACE INTO Track
            (title, album_id, len, rating, count)
            VALUES ( ?, ?, ?, ?, ? )'',
            ( name, album_id, length, rating, count ) )

conn.commit()

```

You can see that we are repeating the pattern of `INSERT OR IGNORE` followed by 1 a `SELECT` to get the appropriate `artist_id` and `album_id` for use in later `INSERT` statements. We start from `Artist` because we need `artist_id` to insert the `Album` and need the `album_id` to insert the `Track`.

If we look at the `Album` table, we can see that the entries were added and assigned 2 a *primary* key as necessary as the data was parsed. We can also see the *foreign key* pointing to a row in the `Artist` table for each `Album` row.

```

sqlite> .mode column
sqlite> SELECT * FROM Album LIMIT 5;
id  artist_id  title
--  -----
1   1           Greatest Hits
2   2           Herzeleid
3   3           Grease
4   4           IV
5   5           The Wall [Disc 2]

```

4


We can reconstruct all of the `Track` data, following all the relations using `JOIN` / 5 `ON` clauses. You can see both ends of each of the (2) relational connections in each row in the output below:

```

sqlite> .mode line
sqlite> SELECT * FROM Track
...> JOIN Album ON Track.album_id = Album.id
...> JOIN Artist ON Album.artist_id = Artist.id
...> LIMIT 2;
    id = 1
    title = Another One Bites The Dust
    album_id = 1
    len = 217103
    rating = 100
    count = 55
    id = 1
    artist_id = 1
    title = Greatest Hits
    id = 1
    name = Queen

```

6

Figure 15.7: A Many to Many Relationship 2

```

id = 2 3
title = Asche Zu Asche
album_id = 2
len = 231810
rating = 100
count = 79
id = 2
artist_id = 2
title = Herzeleid
id = 2
name = Rammstein
  
```

This example shows three tables and two *one-to-many* relationships between the 4 tables. It also shows how to use indexes and uniqueness constraints to programmatically construct the tables and their relationships.

[https://en.wikipedia.org/wiki/One-to-many_\(data_model\)](https://en.wikipedia.org/wiki/One-to-many_(data_model)) 5

Up next we will look at the many-to-many relationships in data models. 6

15.12 Many to many relationships in databases 7

Some data relationships cannot be modeled by a simple one-to-many relationship. 8 For example, lets say we are going to build a data model for a course management system. There will be courses, users, and rosters. A user can be on the roster for many courses and a course will have many users on its roster.

It is pretty simple to *draw* a many-to-many relationship as shown in Figure 15.7. 9 We simply draw two tables and connect them with a line that has the “many” indicator on both ends of the lines. The problem is how to *implement* the relationship using primary keys and foreign keys.

Before we explore how we implement many-to-many relationships, let’s see if we 10 could hack something up by extending a one-to many relationship.

If SQL supported the notion of arrays, we might try to define this: 11

```

CREATE TABLE Course (
  id      INTEGER PRIMARY KEY,
  title   TEXT UNIQUE
  student_ids ARRAY OF INTEGER;
);
12
  
```



Figure 15.8: A Many to Many Connector Table 2

Sadly, while this is a tempting idea, SQL does not support arrays.³³

Or we could just make long string and concatenate all the User primary keys into⁴⁴ a long string separated by commas.

```
CREATE TABLE Course (          5
    id      INTEGER PRIMARY KEY,
    title   TEXT UNIQUE
    student_ids ARRAY OF INTEGER;
);
```

```
INSERT INTO Course (title, student_ids) 6
VALUES( 'si311', '1,3,4,5,6,9,14');
```

This would be very inefficient because as the course roster grows in size and the⁷⁷ number of courses increases it becomes quite expensive to figure out which courses have student 14 on their roster.

Instead of either of these approaches, we model a many-to-many relationship using⁸⁸ an additional table that we call a “junction table”, “through table”, “connector table”, or “join table” as shown in Figure 15.8. The purpose of this table is to capture the *connection* between *a* course and *a* student.

In a sense the table sits between the Course and User table and has a one-to-many⁹⁹ relationship to both tables. By using an intermediate table we break a many-to-many relationship into two one-to-many relationships. Databases are very good at modeling and processing one-to-many relationships.

An example Member table would be as follows:¹⁰¹⁰

```
CREATE TABLE User (          11
    id      INTEGER PRIMARY KEY,
    name   TEXT UNIQUE
);

CREATE TABLE Course (
    id      INTEGER PRIMARY KEY,
    title   TEXT UNIQUE
);
```

```
CREATE TABLE Member ( 12
```

³Some SQL dialects support arrays but arrays do not scale well. NoSQL databases use arrays and data replication but at a cost of database integrity. NoSQL is a story for another course <https://www.pg4e.com/>

```

user_id      INTEGER,          1
course_id    INTEGER,
PRIMARY KEY (user_id, course_id)
);

```

Following our naming convention, Member.user_id and Member.course_id are 2 foreign keys pointing at the corresponding rows in the User and Course tables. Each entry in the member table links a row in the User table to a row in the Course table by going *through* the Member table.

We indicate that the *combination* of course_id and user_id is the PRIMARY KEY 3 for the Member table, also creating an uniqueness constraint for a course_id / user_id combination.

Now lets say we need to insert a number of students into the rosters of a number of 4 courses. Lets assume the data comes to us in a JSON-formatted file with records like this:

```

[           5
  [ "Charley", "si110"], 
  [ "Mea", "si110"], 
  [ "Hattie", "si110"], 
  [ "Keziah", "si110"], 
  [ "Rosa", "si106"], 
  [ "Mea", "si106"], 
  [ "Mairin", "si106"], 
  [ "Zendel", "si106"], 
  [ "Honie", "si106"], 
  [ "Rosa", "si106"], 
  ...
]

```

We could write code as follows to read the JSON file and insert the members of 6 each course roster into the database using the following code:

```

import json
import sqlite3

conn = sqlite3.connect('rosterdb.sqlite')
cur = conn.cursor()

str_data = open('roster_data_sample.json').read()
json_data = json.loads(str_data)

for entry in json_data:

    name = entry[0]
    title = entry[1]

    print((name, title))

    cur.execute(''INSERT OR IGNORE INTO User (name)

```

```

    VALUES ( ? )'''', ( name, ) )
cur.execute('SELECT id FROM User WHERE name = ? ', (name, ))
user_id = cur.fetchone()[0]

cur.execute('''INSERT OR IGNORE INTO Course (title)
    VALUES ( ? )''', ( title, ) )
cur.execute('SELECT id FROM Course WHERE title = ? ', (title, ))
course_id = cur.fetchone()[0]

cur.execute('''INSERT OR REPLACE INTO Member
    (user_id, course_id) VALUES ( ?, ? )''',
    ( user_id, course_id ) )

conn.commit()

```

Like in a previous example, we first make sure that we have an entry in the `User` 2 table and know the primary key of the entry as well as an entry in the `Course` table and know its primary key. We use the ‘`INSERT OR IGNORE`’ and ‘`SELECT`’ pattern so our code works regardless of whether the record is in the table or not.

Our insert into the `Member` table is simply inserting the two integers as a new or 3 existing row depending on the constraint to make sure we do not end up with duplicate entries in the `Member` table for a particular `user_id` / `course_id` combination.

To reconstruct our data across all three tables, we again use `JOIN / ON` to construct 4 a `SELECT` query;

```

sqlite> SELECT * FROM Course
...> JOIN Member ON Course.id = Member.course_id
...> JOIN User ON Member.user_id = User.id;
+-----+-----+-----+-----+
| id | title | user_id | course_id | id | name |
+-----+-----+-----+-----+
| 1 | si110 | 1 | 1 | 1 | Charley |
| 1 | si110 | 2 | 1 | 2 | Mea |
| 1 | si110 | 3 | 1 | 3 | Hattie |
| 1 | si110 | 4 | 1 | 4 | Lyena |
| 1 | si110 | 5 | 1 | 5 | Keziah |
| 1 | si110 | 6 | 1 | 6 | Ellyce |
| 1 | si110 | 7 | 1 | 7 | Thalia |
| 1 | si110 | 8 | 1 | 8 | Meabh |
| 2 | si106 | 2 | 2 | 2 | Mea |
| 2 | si106 | 10 | 2 | 10 | Mairin |
| 2 | si106 | 11 | 2 | 11 | Zendel |
| 2 | si106 | 12 | 2 | 12 | Honie |
| 2 | si106 | 9 | 2 | 9 | Rosa |
+-----+-----+-----+-----+
sqlite>

```

You can see the three tables from left to right - `Course`, `Member`, and `User` and you 6 can see the connections between the primary keys and foreign keys in each row of output.

15.13 Modeling data at the many-to-many connection¹

While we have presented the “join table” as having two foreign keys making a connection between rows in two tables, this is the simplest form of a join table. It is quite common to want to add some data to the connection itself.²

Continuing with our example of users, courses, and rosters to model a simple learning management system, we will also need to understand the *role* that each user is assigned in each course.³

If we first try to solve this by adding an “instructor” flag to the `User` table, we will find that this does not work because a user can be an instructor in one course and a student in another course. If we add an `instructor_id` to the `Course` table it will not work because a course can have multiple instructors. And there is no one-to-many hack that can deal with the fact that the number of roles will expand into roles like Teaching Assistant or Parent.⁴

But if we simply add a `role` column to the `Member` table - we can represent a wide range of roles, role combinations, etc.⁵

Lets change our member table as follows:⁶

```
DROP TABLE Member;7
```

```
CREATE TABLE Member (
    user_id      INTEGER,
    course_id    INTEGER,
    role         INTEGER,
    PRIMARY KEY (user_id, course_id)
);8
```

For simplicity, we will decide that zero in the role means “student” and one in the `role` means instructor. Lets assume our JSON data is augmented with the role as follows:⁹

```
[10
  [ "Charley", "si110", 1],
  [ "Mea", "si110", 0],
  [ "Hattie", "si110", 0],
  [ "Keziah", "si110", 0],
  [ "Rosa", "si106", 0],
  [ "Mea", "si106", 1],
  [ "Mairin", "si106", 0],
  [ "Zendel", "si106", 0],
  [ "Honie", "si106", 0],
  [ "Rosa", "si106", 0],
  ...
]
```

We could alter the `roster.py` program above to incorporate role as follows:¹¹

```

for entry in json_data: 1

    name = entry[0]
    title = entry[1]
    role = entry[2]

    ...

    cur.execute(''INSERT OR REPLACE INTO Member
        (user_id, course_id, role) VALUES ( ?, ?, ? )''',
        ( user_id, course_id, role ) )

```

In a real system, we would probably build a `Role` table and make the `role` column 2
in `Member` a foreign key into the `Role` table as follows:

```
DROP TABLE Member; 3
```

```

CREATE TABLE Member ( 4
    user_id      INTEGER,
    course_id    INTEGER,
    role_id      INTEGER,
    PRIMARY KEY (user_id, course_id, role_id)
);

```

```

CREATE TABLE Role ( 5
    id           INTEGER PRIMARY KEY,
    name         TEXT UNIQUE
);

```

```

INSERT INTO Role (id, name) VALUES (0, 'Student'); 6
INSERT INTO Role (id, name) VALUES (1, 'Instructor');

```

Notice that because we declared the `id` column in the `Role` table as a PRIMARY 7
KEY, we *could* omit it in the `INSERT` statement. But we can also choose the `id`
value as long as the value is not already in the `id` column and does not violate the
implied `UNIQUE` constraint on primary keys.

15.14 Summary⁸

This chapter has covered a lot of ground to give you an overview of the basics of 9
using a database in Python. It is more complicated to write the code to use a
database to store data than Python dictionaries or flat files so there is little reason
to use a database unless your application truly needs the capabilities of a database.
The situations where a database can be quite useful are: (1) when your application
needs to make many small random updates within a large data set, (2) when your
data is so large it cannot fit in a dictionary and you need to look up information
repeatedly, or (3) when you have a long-running process that you want to be able
to stop and restart and retain the data from one run to the next.

You can build a simple database with a single table to suit many application needs,¹ but most problems will require several tables and links/relationships between rows in different tables. When you start making links between tables, it is important to do some thoughtful design and follow the rules of database normalization to make the best use of the database's capabilities. Since the primary motivation for using a database is that you have a large amount of data to deal with, it is important to model your data efficiently so your programs run as fast as possible.

15.15 Debugging²

One common pattern when you are developing a Python program to connect to an SQLite database will be to run a Python program and check the results using the Database Browser for SQLite. The browser allows you to quickly check to see if your program is working properly.³

You must be careful because SQLite takes care to keep two programs from changing the same data at the same time. For example, if you open a database in the browser and make a change to the database and have not yet pressed the “save” button in the browser, the browser “locks” the database file and keeps any other program from accessing the file. In particular, your Python program will not be able to access the file if it is locked.⁴

So a solution is to make sure to either close the database browser or use the *File* menu to close the database in the browser before you attempt to access the database from Python to avoid the problem of your Python code failing because the database is locked.⁵

15.16 Glossary⁶

attribute One of the values within a tuple. More commonly called a “column” or “field”.⁷

constraint When we tell the database to enforce a rule on a field or a row in a table. A common constraint is to insist that there can be no duplicate values in a particular field (i.e., all the values must be unique).⁸

cursor A cursor allows you to execute SQL commands in a database and retrieve data from the database. A cursor is similar to a socket or file handle for network connections and files, respectively.⁹

database browser A piece of software that allows you to directly connect to a database and manipulate the database directly without writing a program.¹⁰

foreign key A numeric key that points to the primary key of a row in another table. Foreign keys establish relationships between rows stored in different tables.¹¹

index Additional data that the database software maintains as rows and inserts into a table to make lookups very fast.¹²

logical key A key that the “outside world” uses to look up a particular row. For example in a table of user accounts, a person’s email address might be a good candidate as the logical key for the user’s data.¹³

normalization	Designing a data model so that no data is replicated. We store each item of data at one place in the database and reference it elsewhere using a foreign key.	1
primary key	A numeric key assigned to each row that is used to refer to one row in a table from another table. Often the database is configured to automatically assign primary keys as rows are inserted.	2
relation	An area within a database that contains tuples and attributes. More typically called a “table”.	3
tuple	A single entry in a database table that is a set of attributes. More typically called “row”.	4

Chapter 16¹

Visualizing data²

So far we have been learning the Python language and then learning how to use³ Python, the network, and databases to manipulate data.

In this chapter, we take a look at three complete applications that bring all of these⁴ things together to manage and visualize data. You might use these applications as sample code to help get you started in solving a real-world problem.

Each of the applications is a ZIP file that you can download and extract onto your⁵ computer and execute.

16.1 Building a OpenStreetMap from geocoded⁶ data

In this project, we are using the OpenStreetMap geocoding API to clean up some⁷ user-entered geographic locations of university names and then placing the data on an actual OpenStreetMap.

To get started, download the application from:⁸

www.py4e.com/code3/opengeo.zip⁹

The first problem to solve is that these geocoding APIs are rate-limited to a certain¹⁰ number of requests per day. If you have a lot of data, you might need to stop and restart the lookup process several times. So we break the problem into two phases.

In the first phase we take our input “survey” data in the file *where.data* and read it¹¹ one line at a time, and retrieve the geocoded information from Google and store it in a database *geodata.sqlite*. Before we use the geocoding API for each user-entered location, we simply check to see if we already have the data for that particular line of input. The database is functioning as a local “cache” of our geocoding data to make sure we never ask Google for the same data twice.

You can restart the process at any time by removing the file *geodata.sqlite*.¹²

Run the *geoload.py* program. This program will read the input lines in *where.data*¹³ and for each line check to see if it is already in the database. If we don’t have the




Figure 16.1: An OpenStreetMap 2

data for the location, it will call the geocoding API to retrieve the data and store 3 it in the database.

Here is a sample run after there is already some data in the database: 4

Found in database AGH University of Science and Technology 5

Found in database Academy of Fine Arts Warsaw Poland 6

Found in database American University in Cairo 7

Found in database Arizona State University 8

Found in database Athens Information Technology 9

Retrieving <https://py4e-data.dr-chuck.net/> 10
opengeo?q=BITS+Pilani

Retrieved 794 characters {"type": "FeatureColl" 11

Retrieving <https://py4e-data.dr-chuck.net/> 12
opengeo?q=Babcock+University

Retrieved 760 characters {"type": "FeatureColl" 13

Retrieving <https://py4e-data.dr-chuck.net/> 14
opengeo?q=Banaras+Hindu+University

Retrieved 866 characters {"type": "FeatureColl" 15

... 16

The first five locations are already in the database and so they are skipped. The 17 program scans to the point where it finds new locations and starts retrieving them.

The *geoload.py* program can be stopped at any time, and there is a counter that you can use to limit the number of calls to the geocoding API for each run. Given that the *where.data* only has a few hundred data items, you should not run into the daily rate limit, but if you had more data it might take several runs over several days to get your database to have all of the geocoded data for your input.

Once you have some data loaded into *geodata.sqlite*, you can visualize the data using the *geodump.py* program. This program reads the database and writes the file *where.js* with the location, latitude, and longitude in the form of executable JavaScript code.

A run of the *geodump.py* program is as follows:

```
AGH University of Science and Technology, Czarnowiejska, 4
Czarna Wies, Krowodrza, Krakow, Lesser Poland
Voivodeship, 31-126, Poland 50.0657 19.91895
```

```
Academy of Fine Arts, Krakowskie Przedmiescie, 5
Northern Srodmiescie, Srodmiescie, Warsaw, Masovian
Voivodeship, 00-046, Poland 52.239 21.0155
```

```
...
260 lines were written to where.js
Open the where.html file in a web browser to view the data.
```

The file *where.html* consists of HTML and JavaScript to visualize a Google map.⁷ It reads the most recent data in *where.js* to get the data to be visualized. Here is the format of the *where.js* file:


```
myData = [
[50.0657,19.91895,
'AGH University of Science and Technology, Czarnowiejska,
Czarna Wies, Krowodrza, Krakow, Lesser Poland
Voivodeship, 31-126, Poland '],
[52.239,21.0155,
'Academy of Fine Arts, Krakowskie Przedmiesciee,
Srodmiescie Polnocne, Srodmiescie, Warsaw,
Masovian Voivodeship, 00-046, Poland'],
...
];
```

This is a JavaScript variable that contains a list of lists. The syntax for JavaScript list constants is very similar to Python, so the syntax should be familiar to you.

Simply open *where.html* in a browser to see the locations. You can hover over each map pin to find the location that the geocoding API returned for the user-entered input. If you cannot see any data when you open the *where.html* file, you might want to check the JavaScript or developer console for your browser.

16.2 Visualizing networks and interconnections¹¹

In this application, we will perform some of the functions of a search engine. We will first spider a small subset of the web and run a simplified version of the



1

Figure 16.2: A Page Ranking 2

Google page rank algorithm to determine which pages are most highly connected,
3 and then visualize the page rank and connectivity of our small corner of the web.
We will use the D3 JavaScript visualization library <http://d3js.org/> to produce the
visualization output.

You can download and extract this application from:
4

www.py4e.com/code3/pagerank.zip 5

The first program (*spider.py*) crawls a web site and pulls a series of pages into the
6 database (*spider.sqlite*), recording the links between pages. You can restart the
process at any time by removing the *spider.sqlite* file and rerunning *spider.py*.

```
Enter web url or enter: http://www.dr-chuck.com/ 7
['http://www.dr-chuck.com']
How many pages:2
1 http://www.dr-chuck.com/ 12
2 http://www.dr-chuck.com/csev-blog/ 57
How many pages:
```

In this sample run, we told it to crawl a website and retrieve two pages. If you
8 restart the program and tell it to crawl more pages, it will not re-crawl any pages
already in the database. Upon restart it goes to a random non-crawled page and
starts there. So each successive run of *spider.py* is additive.

```
Enter web url or enter: http://www.dr-chuck.com/ 9
['http://www.dr-chuck.com']
```

How many pages:3 1

```
3 http://www.dr-chuck.com/csev-blog 57 2
4 http://www.dr-chuck.com/dr-chuck/resume/speaking.htm 1
5 http://www.dr-chuck.com/dr-chuck/resume/index.htm 13
```

How many pages: 3

You can have multiple starting points in the same database—within the program,⁴ these are called “webs”. The spider chooses randomly amongst all non-visited links across all the webs as the next page to spider.

If you want to dump the contents of the *spider.sqlite* file, you can run *spdum.py*⁵ as follows:

```
(5, None, 1.0, 3, 'http://www.dr-chuck.com/csev-blog') 6
(3, None, 1.0, 4, 'http://www.dr-chuck.com/dr-chuck/resume/speaking.htm')
(1, None, 1.0, 2, 'http://www.dr-chuck.com/csev-blog/')
(1, None, 1.0, 5, 'http://www.dr-chuck.com/dr-chuck/resume/index.htm')
```

4 rows 7

This shows the number of incoming links, the old page rank, the new page rank,⁸ the id of the page, and the url of the page. The *spdum.py* program only shows pages that have at least one incoming link to them.

Once you have a few pages in the database, you can run page rank on the pages⁹ using the *sprank.py* program. You simply tell it how many page rank iterations to run.

How many iterations:2 10

```
1 0.546848992536 11
2 0.226714939664 12
```

[(1, 0.559), (2, 0.659), (3, 0.985), (4, 2.135), (5, 0.659)] 13

You can dump the database again to see that page rank has been updated:¹⁴

```
(5, 1.0, 0.985, 3, 'http://www.dr-chuck.com/csev-blog') 15
(3, 1.0, 2.135, 4, 'http://www.dr-chuck.com/dr-chuck/resume/speaking.htm')
(1, 1.0, 0.659, 2, 'http://www.dr-chuck.com/csev-blog/')
(1, 1.0, 0.659, 5, 'http://www.dr-chuck.com/dr-chuck/resume/index.htm')
4 rows
```

You can run *sprank.py* as many times as you like and it will simply refine the page rank each time you run it. You can even run *sprank.py* a few times and then go spider a few more pages with *spider.py* and then run *sprank.py* to reconverge the page rank values. A search engine usually runs both the crawling and ranking programs all the time.¹⁶

If you want to restart the page rank calculations without respidering the web pages,¹⁷ you can use *sprest.py* and then restart *sprank.py*.

How many iterations:50 18

```
1 0.546848992536 19
2 0.226714939664
```

```

3 0.0659516187242
4 0.0244199333
5 0.0102096489546
6 0.00610244329379
...
42 0.000109076928206
43 9.91987599002e-05
44 9.02151706798e-05
45 8.20451504471e-05
46 7.46150183837e-05
47 6.7857770908e-05
48 6.17124694224e-05
49 5.61236959327e-05
50 5.10410499467e-05
[(512, 0.0296), (1, 12.79), (2, 28.93), (3, 6.808), (4, 13.46)]

```

1

For each iteration of the page rank algorithm it prints the average change in page rank per page. The network initially is quite unbalanced and so the individual page rank values change wildly between iterations. But in a few short iterations, the page rank converges. You should run *sprank.py* long enough that the page rank values converge.

If you want to visualize the current top pages in terms of page rank, run *spjson.py* to read the database and write the data for the most highly linked pages in JSON format to be viewed in a web browser.

```

Creating JSON output on spider.json...
How many nodes? 30
Open force.html in a browser to view the visualization

```

4

You can view this data by opening the file *force.html* in your web browser. This shows an automatic layout of the nodes and links. You can click and drag any node and you can also double-click on a node to find the URL that is represented by the node.

If you rerun the other utilities, rerun *spjson.py* and press refresh in the browser to get the new data from *spider.json*.

16.3 Visualizing mail data⁷

Up to this point in the book, you have become quite familiar with our *mbox-short.txt* and *mbox.txt* data files. Now it is time to take our analysis of email data to the next level.

In the real world, sometimes you have to pull down mail data from servers. That might take quite some time and the data might be inconsistent, error-filled, and need a lot of cleanup or adjustment. In this section, we work with an application that is the most complex so far and pull down nearly a gigabyte of data and visualize it.

You can download this application from:¹⁰
<https://www.py4e.com/code3/gmane.zip>¹¹

2

3

5

6

8

9




Figure 16.3: A Word Cloud from the Sakai Developer List 2

We will be using data from a free email list archiving service that was called *gmane* - the service has since been shut down and for the purposes of this course, a partial archive has been kept at <http://mbox.dr-chuck.net>. The gmane service was very popular with open source projects because it provided a nice searchable archive of their email activity.

<http://mbox.dr-chuck.net/export.php> 4

When the Sakai email data was spidered using this software, it produced nearly a Gigabyte of data and took a number of runs on several days. The file *README.txt* in the above ZIP may have instructions as to how you can download a pre-spidered copy of the *content.sqlite* file for a majority of the Sakai email corpus so you don't have to spider for five days just to run the programs. If you download the pre-spidered content, you should still run the spidering process to catch up with more recent messages.

The first step is to spider the repository. The base URL is hard-coded in the *gmane.py* and is hard-coded to the Sakai developer list. You can spider another repository by changing that base url. Make sure to delete the *content.sqlite* file if you switch the base url.

The *gmane.py* file operates as a responsible caching spider in that it runs slowly and retrieves one mail message per second so as to avoid getting throttled. It stores all of its data in a database and can be interrupted and restarted as often as needed. It may take many hours to pull all the data down. So you may need to restart several times.

Here is a run of *gmane.py* retrieving the last five messages of the Sakai developer list:

```
How many messages:10
http://mbox.dr-chuck.net/sakai.devel/51410/51411 9460
    nealcaidin@sakaifoundation.org 2013-04-05 re: [building ...
http://mbox.dr-chuck.net/sakai.devel/51411/51412 3379
    samuelgutierrezjimenez@gmail.com 2013-04-06 re: [building ...
http://mbox.dr-chuck.net/sakai.devel/51412/51413 9903
    da1@vt.edu 2013-04-05 [building sakai] melete 2.9 oracle ...
http://mbox.dr-chuck.net/sakai.devel/51413/51414 349265
    m.shedid@elraed-it.com 2013-04-07 [building sakai] ...
http://mbox.dr-chuck.net/sakai.devel/51414/51415 3481
    samuelgutierrezjimenez@gmail.com 2013-04-07 re: ...
http://mbox.dr-chuck.net/sakai.devel/51415/51416 0
```

Does not start with From 2

The program scans *content.sqlite* from one up to the first message number not 3 already spidered and starts spidering at that message. It continues spidering until it has spidered the desired number of messages or it reaches a page that does not appear to be a properly formatted message.

Sometimes the repository is missing a message. Perhaps administrators can delete 4 messages or perhaps they get lost. If your spider stops, and it seems it has hit a missing message, go into the SQLite Manager and add a row with the missing id leaving all the other fields blank and restart *gmane.py*. This will unstick the spidering process and allow it to continue. These empty messages will be ignored in the next phase of the process.

One nice thing is that once you have spidered all of the messages and have them 5 in *content.sqlite*, you can run *gmane.py* again to get new messages as they are sent to the list.

The *content.sqlite* data is pretty raw, with an inefficient data model, and not 6 compressed. This is intentional as it allows you to look at *content.sqlite* in the SQLite Manager to debug problems with the spidering process. It would be a bad idea to run any queries against this database, as they would be quite slow.

The second process is to run the program *gmodel.py*. This program reads the raw 7 data from *content.sqlite* and produces a cleaned-up and well-modeled version of the data in the file *index.sqlite*. This file will be much smaller (often 10X smaller) than *content.sqlite* because it also compresses the header and body text.

Each time *gmodel.py* runs it deletes and rebuilds *index.sqlite*, allowing you to adjust 8 its parameters and edit the mapping tables in *content.sqlite* to tweak the data cleaning process. This is a sample run of *gmodel.py*. It prints a line out each time 250 mail messages are processed so you can see some progress happening, as this program may run for a while processing nearly a Gigabyte of mail data.

```
Loaded allsenders 1588 and mapping 28 dns mapping 1 9
1 2005-12-08T23:34:30-06:00 ggolden22@mac.com
251 2005-12-22T10:03:20-08:00 tpamsler@ucdavis.edu
501 2006-01-12T11:17:34-05:00 lance@indiana.edu
751 2006-01-24T11:13:28-08:00 vrajgopalan@ucmerced.edu
...
```

The *gmodel.py* program handles a number of data cleaning tasks.¹

Domain names are truncated to two levels for .com, .org, .edu, and .net. Other² domain names are truncated to three levels. So si.umich.edu becomes umich.edu and caret.cam.ac.uk becomes cam.ac.uk. Email addresses are also forced to lower case, and some of the @gmane.org address like the following

```
arwhyte-63aXycvo3TyHXe+LvDLADg@public.gmane.org3
```

are converted to the real address whenever there is a matching real email address⁴ elsewhere in the message corpus.

In the *mapping.sqlite* database there are two tables that allow you to map both⁵ domain names and individual email addresses that change over the lifetime of the email list. For example, Steve Githens used the following email addresses as he changed jobs over the life of the Sakai developer list:

```
s-githens@northwestern.edu6  
sgithens@cam.ac.uk  
swgithen@mtu.edu
```

We can add two entries to the Mapping table in *mapping.sqlite* so *gmodel.py* will⁷ map all three to one address:

```
s-githens@northwestern.edu -> swgithen@mtu.edu8  
sgithens@cam.ac.uk -> swgithen@mtu.edu
```

You can also make similar entries in the DNSMapping table if there are multiple⁹ DNS names you want mapped to a single DNS. The following mapping was added to the Sakai data:

```
iupui.edu -> indiana.edu10
```

so all the accounts from the various Indiana University campuses are tracked to-¹¹gether.

You can rerun the *gmodel.py* over and over as you look at the data, and add¹² mappings to make the data cleaner and cleaner. When you are done, you will have a nicely indexed version of the email in *index.sqlite*. This is the file to use to do data analysis. With this file, data analysis will be really quick.

The first, simplest data analysis is to determine “who sent the most mail?” and¹³ “which organization sent the most mail”? This is done using *gbasic.py*:

```
How many to dump? 514  
Loaded messages= 51330 subjects= 25033 senders= 1584
```

```
Top 5 Email list participants15  
steve.swinsburg@gmail.com 2657  
azeckoski@unicon.net 1742  
ieb@tfd.co.uk 1591  
csev@umich.edu 1304
```

```
david.horwitz@uct.ac.za 1184 1
```

```
Top 5 Email list organizations 2
gmail.com 7339
umich.edu 6243
uct.ac.za 2451
indiana.edu 2258
unicon.net 2055
```

Note how much more quickly *gbasic.py* runs compared to *gmane.py* or even *gmodel.py*. They are all working on the same data, but *gbasic.py* is using the compressed and normalized data in *index.sqlite*. If you have a lot of data to manage, a multistep process like the one in this application may take a little longer to develop, but will save you a lot of time when you really start to explore and visualize your data. 3

You can produce a simple visualization of the word frequency in the subject lines 4 in the file *gword.py*:

```
Range of counts: 33229 129 5
Output written to gword.js 6
```

This produces the file *gword.js* which you can visualize using *gword.htm* to produce 7 a word cloud similar to the one at the beginning of this section.

A second visualization is produced by *gline.py*. It computes email participation by 8 organizations over time.

```
Loaded messages= 51330 senders= 1584 9
Top 10 Organizations
['gmail.com', 'umich.edu', 'uct.ac.za', 'indiana.edu',
'unicon.net', 'tfd.co.uk', 'berkeley.edu', 'longsight.com',
'stanford.edu', 'ox.ac.uk']
Output written to gline.js
```

Its output is written to *gline.js* which is visualized using *gline.htm*. 10

This is a relatively complex and sophisticated application and has features to do 11 some real data retrieval, cleaning, and visualization.




Figure 16.4: Sakai Mail Activity by Organization 2

Appendix A¹

Contributions²

A.1 Translations³

This book has been translated into a number of languages which can be found at ⁴
<https://www.py4e.com/book>.

- Spanish - Python para todos: Explorando la información con Python 3 ⁵
Translated book, autograders, resources, and web site at <https://es.py4e.com>
Contributors: Juan Carlos Perez Castellanos, Juan Dougnac, Daniel Merino Echeverría, Jaime Bermeo Ramírez and Fernando Tardío.
- Italian Python per tutti: Esplorare dati con Python3 Contributors: Alessandro Rossetti and Vittore Zen
- Portuguese - Python Para Todos: Explorando Dados com Python 3 Translation: Yuri Loia de Medeiros
- Polish - Python dla wszystkich: Odkrywanie danych z Python 3 Translated book, autograders, resources, and web site at <https://py4e.pl> Translation: Andrzej Wójtowicz (Adam Mickiewicz University in Poznań, Poland)
- Greek - Translated book, autograders, resources, and web site at <https://gr.py4e.com> Translation: Konstantia Kiourtidou
- Arabic - Translation: Electronics Go

Anyone who is given permission to translate the book agrees to provide a free copy ⁶
of the book online. Some of the translations are available in print and available
for purchase.

A.2 Contributor List for Python for Everybody⁷

The source for the book is maintained in [github](#) and you can find numerous ⁸
contributors for the book who find and suggest fixes using pull requests.

<https://github.com/csev/py4e/> ¹

Look under **Insights** → **Contributors** to see all the folks who have contributed ² through **github**.

Juan Carlos Perez Castellanos, Juan Dougnac, Daniel Merino Echeverría, Jaime Bermeo Ramírez, Fernando Tardío, Alessandro Rossetti, Vittore Zen, Yuri Loia de Medeiros, Konstantia Kiourtidou, Andrzej Wójtowicz, Elliott Hauser, Stephen Catto, Sue Blumenberg, Tamara Brunnock, Mihaela Mack, Chris Kolosiwsky, Dustin Farley, Jens Leerssen, Naveen KT, Mirza Ibrahimovic, Naveen (@togarnk), Zhou Fangyi, Alistair Walsh, Erica Brody, Jih-Sheng Huang, Louis Luangkesorn, and Michael Fudge ³

You can see contribution details at: ⁴

<https://github.com/csev/py4e/graphs/contributors> ⁵

A.3 Contributor List for Python for Informatics ⁶

Bruce Shields for copy editing early drafts, Sarah Hegge, Steven Cherry, Sarah Kathleen Barbarow, Andrea Parker, Radaphat Chongthammakun, Megan Hixon, Kirby Urner, Sarah Kathleen Barbrow, Katie Kujala, Noah Botimer, Emily Alinder, Mark Thompson-Kular, James Perry, Eric Hofer, Eytan Adar, Peter Robinson, Deborah J. Nelson, Jonathan C. Anthony, Eden Rassette, Jeannette Schroeder, Justin Feezell, Chuanqi Li, Gerald Gordinier, Gavin Thomas Strassel, Ryan Clement, Alissa Talley, Caitlin Holman, Yong-Mi Kim, Karen Stover, Cherie Edmonds, Maria Seiferle, Romer Kristi D. Aranas (RK), Grant Boyer, Hedemarie Dussan, ⁷

A.4 Preface for “Think Python” ⁸

A.4.1 The strange history of “Think Python” ⁹

(Allen B. Downey) ¹⁰

In January 1999 I was preparing to teach an introductory programming class in Java. I had taught it three times and I was getting frustrated. The failure rate in the class was too high and, even for students who succeeded, the overall level of achievement was too low. ¹¹

One of the problems I saw was the books. They were too big, with too much unnecessary detail about Java, and not enough high-level guidance about how to program. And they all suffered from the trap door effect: they would start out easy, proceed gradually, and then somewhere around Chapter 5 the bottom would fall out. The students would get too much new material, too fast, and I would spend the rest of the semester picking up the pieces. ¹²

Two weeks before the first day of classes, I decided to write my own book. My goals were: ¹³

- Keep it short. It is better for students to read 10 pages than not read 50 pages. ¹⁴

- Be careful with vocabulary. I tried to minimize the jargon and define each term at first use.
- Build gradually. To avoid trap doors, I took the most difficult topics and split them into a series of small steps.
- Focus on programming, not the programming language. I included the minimum useful subset of Java and left out the rest.

I needed a title, so on a whim I chose *How to Think Like a Computer Scientist*.²

My first version was rough, but it worked. Students did the reading, and they³ understood enough that I could spend class time on the hard topics, the interesting topics and (most important) letting the students practice.

I released the book under the GNU Free Documentation License, which allows⁴ users to copy, modify, and distribute the book.

What happened next is the cool part. Jeff Elkner, a high school teacher in Virginia,⁵ adopted my book and translated it into Python. He sent me a copy of his translation, and I had the unusual experience of learning Python by reading my own book.

Jeff and I revised the book, incorporated a case study by Chris Meyers, and in 2001⁶ we released *How to Think Like a Computer Scientist: Learning with Python*, also under the GNU Free Documentation License. As Green Tea Press, I published the book and started selling hard copies through Amazon.com and college book stores. Other books from Green Tea Press are available at greenteapress.com.

In 2003 I started teaching at Olin College and I got to teach Python for the first⁷ time. The contrast with Java was striking. Students struggled less, learned more, worked on more interesting projects, and generally had a lot more fun.

Over the last five years I have continued to develop the book, correcting errors,⁸ improving some of the examples and adding material, especially exercises. In 2008 I started work on a major revision—at the same time, I was contacted by an editor at Cambridge University Press who was interested in publishing the next edition. Good timing!

I hope you enjoy working with this book, and that it helps you learn to program⁹ and think, at least a little bit, like a computer scientist.

A.4.2 Acknowledgements for “Think Python”¹⁰

(Allen B. Downey)¹¹

First and most importantly, I thank Jeff Elkner, who translated my Java book into¹² Python, which got this project started and introduced me to what has turned out to be my favorite language.

I also thank Chris Meyers, who contributed several sections to *How to Think Like a Computer Scientist*.¹³

And I thank the Free Software Foundation for developing the GNU Free Documenta-¹⁴ tion License, which helped make my collaboration with Jeff and Chris possible.

I also thank the editors at Lulu who worked on *How to Think Like a Computer Scientist*.¹

I thank all the students who worked with earlier versions of this book and all the² contributors (listed in an Appendix) who sent in corrections and suggestions.

And I thank my wife, Lisa, for her work on this book, and Green Tea Press, and³ everything else, too.

Allen B. Downey⁴
Needham MA

Allen Downey is an Associate Professor of Computer Science at the Franklin W.⁵ Olin College of Engineering.

A.5 Contributor List for “Think Python”⁶

(Allen B. Downey)⁷

More than 100 sharp-eyed and thoughtful readers have sent in suggestions and⁸ corrections over the past few years. Their contributions, and enthusiasm for this project, have been a huge help.

For the detail on the nature of each of the contributions from these individuals,⁹ see the “Think Python” text.

Lloyd Hugh Allen, Yvon Boulianne, Fred Bremmer, Jonah Cohen, Michael Conlon, Benoit Girard, Courtney Gleason and Katherine Smith, Lee Harr, James Kaylin, David Kershaw, Eddie Lam, Man-Yong Lee, David Mayo, Chris McAloon, Matthew J. Moelter, Simon Dicon Montford, John Ouzts, Kevin Parks, David Pool, Michael Schmitt, Robin Shaw, Paul Sleigh, Craig T. Snydal, Ian Thomas, Keith Verheyden, Peter Winstanley, Chris Wrobel, Moshe Zadka, Christoph Zwierschke, James Mayer, Hayden McAfee, Angel Arnal, Tauhidul Hoque and Lex Berezhny, Dr. Michele Alzetta, Andy Mitchell, Kalin Harvey, Christopher P. Smith, David Hutchins, Gregor Lingl, Julie Peters, Florin Oprina, D. J. Webre, Ken, Ivo Wever, Curtis Yanko, Ben Logan, Jason Armstrong, Louis Cordier, Brian Cain, Rob Black, Jean-Philippe Rey at Ecole Centrale Paris, Jason Mader at George Washington University made a number Jan Gundtofte-Bruun, Abel David and Alexis Dinno, Charles Thayer, Roger Sperberg, Sam Bull, Andrew Cheung, C. Corey Capel, Alessandra, Wim Champagne, Douglas Wright, Jared Spindor, Lin Peiheng, Ray Hagtvedt, Torsten Hübsch, Inga Petuhhov, Arne Babenhauserheide, Mark E. Casida, Scott Tyler, Gordon Shephard, Andrew Turner, Adam Hobart, Daryl Hammond and Sarah Zimmerman, George Sass, Brian Bingham, Leah Engelbert-Fenton, Joe Funke, Chao-chao Chen, Jeff Paine, Lubos Pintes, Gregg Lind and Abigail Heithoff, Max Hailperin, Chotipat Pornavalai, Stanislaw Antol, Eric Pashman, Miguel Azevedo, Jianhua Liu, Nick King, Martin Zuther, Adam Zimmerman, Ratnakar Tiwari, Anurag Goel, Kelli Kratzer, Mark Griffiths, Roydan Ongie, Patryk Wolowiec, Mark Chonofsky, Russell Coleman, Wei Huang, Karen Barber, Nam Nguyen, Stéphane Morin, Fernando Tardio, and Paul Stoop.¹⁰

Appendix B¹

Copyright Detail²

This work is licensed under a Creative Common Attribution-NonCommercial-ShareAlike 3.0 Unported License. This license is available at

creativecommons.org/licenses/by-nc-sa/3.0/.⁴

I would have preferred to license the book under the less restrictive CC-BY-SA license. But unfortunately there are a few unscrupulous organizations who search for and find freely licensed books, and then publish and sell virtually unchanged copies of the books on a print on demand service such as LuLu or KDP. KDP has (thankfully) added a policy that gives the wishes of the actual copyright holder preference over a non-copyright holder attempting to publish a freely licensed work. Unfortunately there are many print-on-demand services and very few have as well-considered a policy as KDP.

Regretfully, I added the NC element to the license this book to give me recourse in case someone tries to clone this book and sell it commercially. Unfortunately, adding NC limits uses of this material that I would like to permit. So I have added this section of the document to describe specific situations where I am giving my permission in advance to use the material in this book in situations that some might consider commercial.

- If you are printing a limited number of copies of all or part of this book for use in a course (e.g., like a coursepack), then you are granted CC-BY license to these materials for that purpose.⁷
- If you are a teacher at a university and you translate this book into a language other than English and teach using the translated book, then you can contact me and I will grant you a CC-BY-SA license to these materials with respect to the publication of your translation. In particular, you will be permitted to sell the resulting translated book commercially.

If you are intending to translate the book, you may want to contact me so we can make sure that you have all of the related course materials so you can translate them as well.⁸

Of course, you are welcome to contact me and ask for permission if these clauses are not sufficient. In all cases, permission to reuse and remix this material will be

granted as long as there is clear added value or benefit to students or teachers that **1** will accrue as a result of the new work.

Charles Severance **2**

www.dr-chuck.com **3**

Ann Arbor, MI, USA **4**

September 9, 2013 **5**

Index¹

access, 92 2
accumulator, 64 3
 sum, 62
algorithm, 53 4
aliasing, 99, 100, 105 5
 copying to avoid, 102 6
alternative execution, 34 7
and operator, 32 8
API, 166 9
 key, 165 10
append method, 94, 101 11
argument, 43, 47, 49, 50, 53, 100 12
 keyword, 121 13
 list, 100 14
 optional, 73, 97 15
arithmetic operator, 22 16
assignment, 29, 91 17
 item, 70, 92, 120 18
 tuple, 122, 128 19
assignment statement, 20 20
attribute, 179, 202 21
Auto increment, 190 22

BeautifulSoup, 154, 157, 169 23
binary file, 151 24
bisection, debugging by, 64 25
body, 40, 47, 53, 58 26
bool type, 31 27
boolean expression, 31, 40 28
boolean operator, 71 29
bracket 30
 squiggly, 109 31
bracket operator, 67, 92, 120 32
branch, 34, 40 33
break statement, 58 34
bug, 15 35
BY-SA, iv 36

cache, 205 37
case-sensitivity, variable names, 29 38
catch, 89 39
CC-BY-SA, iv 40

celsius, 37 41
central processing unit, 15 42
chained conditional, 34, 40 43
character, 67 44
child class, 179 45
choice function, 47 46
class, 173, 179 47
 float, 19 48
 int, 19 49
 str, 19 50
class keyword, 172 51
close method, 88 52
colon, 47 53
comment, 26, 29 54
comparable, 119, 128 55
comparison 56
 string, 71 57
 tuple, 120 58
comparison operator, 31 59
compile, 16 60
composition, 50, 53 61
compound statement, 32, 40 62
concatenation, 24, 29, 70, 98 63
 list, 93, 101 64
condition, 32, 40, 58 65
conditional 66
 chained, 34, 40 67
 nested, 35, 40 68
conditional executions, 32 69
conditional statement, 32, 40 70
connect function, 183 71
Connector table, 197 72
consistency check, 117 73
Constraint, 192 74
constraint, 202 75
construct, 173 76
constructor, 175, 179 77
continue statement, 59 78
contributors, 220 79
conversion 80
 type, 44 81
copy 82

slice, 69, 94 1
 to avoid aliasing, 102 2
 count method, 73 3
 counter, 64, 70, 76, 82, 111 4
 counting and looping, 70 5
 CPU, 15 6
CREATE INDEX, 192 7
 Creative Commons License, iv 8
 Crow's foot diagram, 196 9
 Crow's Foot diagrams, 189 10
CRUD, 185, 187 11
 curl, 157 12
 cursor, 202 13
 cursor function, 183 14

 Data model, 196 15
 Data model diagrams, 189 16
 Data Modelling, 187 17
 Data Normalization, 187 18
 Data Replication, 188 19
 data structure, 128 20
 database, 181 21
 indexes, 181 22
 database browser, 202 23
 Database Normalization, 187 24
 database normalization, 203 25
 debugging, 14, 28, 39, 52, 75, 88, 102, 26
 116, 128
 by bisection, 64 27
 decorate-sort-undecorate pattern, 121 28
 decrement, 57, 64 29
 def keyword, 47 30
 definition 31
 function, 47 32
 del statement, 95 33
 deletion, element of list, 95 34
 delimiter, 97, 105 35
 destructor, 175, 179 36
 deterministic, 46, 53 37
 development plan 38
 random walk programming, 14 39
 dict function, 109 40
 dictionary, 109, 117, 123 41
 looping with, 114 42
 traversal, 124 43
 dir, 174 44
 divisibility, 24 45
 division 46
 floating-point, 23 47
 dot notation, 45, 53, 72 48
 DSU pattern, 121, 128 49

 element, 91, 105 50
 element deletion, 95 51
 ElementTree, 160, 166 52
 find, 160 53
 findall, 161 54
 fromstring, 160 55
 get, 161 56
 elif keyword, 35 57
 ellipses, 47 58
 else keyword, 34 59
 email address, 123 60
 empty list, 91 61
 empty string, 76, 98 62
 encapsulation, 70 63
 end of line character, 89 64
 Entity-Relationship diagrams, 189 65
 equivalence, 99 66
 equivalent, 105 67
 error 68
 runtime, 28 69
 semantic, 20, 29 70
 shape, 128 71
 syntax, 28 72
 error message, 20, 28 73
 evaluate, 23 74
 exception, 28 75
 IndexError, 68, 92 76
 IOError, 86 77
 KeyError, 110 78
 TypeError, 67, 70, 120 79
 ValueError, 26, 123 80
 experimental debugging, 14 81
 expression, 22, 23, 29 82
 boolean, 31, 40 83
 extend method, 94 84
 eXtensible Markup Language, 166 85

 fahrenheit, 37 86
 False special value, 31 87
 file, 79 88
 open, 79 89
 reading, 82 90
 writing, 87 91
 file handle, 80 92
 filter pattern, 83 93
 findall, 134 94
 flag, 76 95
 float function, 44 96
 float type, 19 97
 floating-point, 29 98
 floating-point division, 23 99

- flow control, 150 1
- flow of execution, 49, 53, 58 2
- for loop, 68, 92 3
- for statement, 60 4
- Foreign key, 189, 195 5
- foreign key, 190, 202 6
- formatted string literals, 74 7
- Free Documentation License, GNU, 219 8
- frequency, 111 9
 - letter, 129 10
- fruitful function, 51, 53 11
- function, 47, 53 12
 - choice, 47 13
 - connect, 183 14
 - cursor, 183 15
 - dict, 109 16
 - float, 44 17
 - int, 44 18
 - len, 68, 110 19
 - list, 97 20
 - log, 45 21
 - open, 79, 86 22
 - print, 16 23
 - randint, 46 24
 - random, 46 25
 - repr, 88 26
 - reversed, 127 27
 - sorted, 127 28
 - sqrt, 45 29
 - str, 44 30
 - tuple, 119 31
- function argument, 49 32
- function call, 43, 53 33
- function definition, 47, 48, 53 34
- function object, 48, 53 35
- function parameter, 49 36
- function, fruitful, 51 37
- function, math, 45 38
- function, reasons for, 52 39
- function, trigonometric, 45 40
- function, void, 51 41
- gather, 128 42
- get method, 112 43
- GNU Free Documentation License, 219 44
- Google 45
 - map, 205 46
 - page rank, 207 47
- greedy, 133, 142, 153 48
- greedy matching, 142 49
- grep, 141, 142 50
- guardian pattern, 38, 40, 76 51
- hardware, 3 52
 - architecture, 3 53
- hash function, 117 54
- hash table, 110 55
- hashable, 119, 126, 128 56
- hashtable, 117 57
- header, 47, 53 58
- high-level language, 16 59
- histogram, 111, 117 60
- HTML, 154, 169 61
- identical, 105 62
- identity, 99 63
- idiom, 102, 112, 114 64
- if statement, 32 65
- image 66
 - jpg, 148 67
- immutability, 70, 76, 100, 119, 127 68
- implementation, 111, 117 69
- import statement, 53 70
- in operator, 71, 92, 110 71
- increment, 57, 64 72
- indentation, 47 73
- index, 67, 76, 92, 105, 109, 202 74
 - looping with, 93 75
 - negative, 68 76
 - slice, 69, 94 77
 - starting at zero, 67, 92 78
- IndexError, 68, 92 79
- infinite loop, 58, 64 80
- inheritance, 179 81
- initialization (before update), 57 82
- instance, 173 83
- int function, 44 84
- int type, 19 85
- integer, 29 86
- interactive mode, 6, 16, 21, 51 87
- interpret, 16 88
- invocation, 72, 76 89
- IOError, 86 90
- is operator, 99 91
- item, 76, 91 92
 - dictionary, 117 93
- item assignment, 70, 92, 120 94
- item update, 93 95
- items method, 123 96
- iteration, 57, 64 97
- JavaScript Object Notation, 162, 166 98

join method, 98 1
 Join table, 197 2
 jpg, 148 3
 JSON, 162, 166 4
 parse, 198 5
 Junction table, 197 6
 key, 109, 117 7
 key-value pair, 109, 117, 123 8
 keyboard input, 25 9
 KeyError, 110 10
 keys method, 114 11
 keyword, 21, 29 12
 def, 47 13
 elif, 35 14
 else, 34 15
 keyword argument, 121 16
 keywords, 5 17
 language 18
 programming, 5 19
 len function, 68, 110 20
 letter frequency, 129 21
 list, 91, 97, 105, 127 22
 as argument, 100 23
 concatenation, 93, 101 24
 copy, 94 25
 element, 92 26
 empty, 91 27
 function, 97 28
 index, 92 29
 membership, 92 30
 method, 94 31
 nested, 91, 93 32
 operation, 93 33
 repetition, 93 34
 slice, 94 35
 traversal, 92, 105 36
 list comprehension, 127 37
 list object, 168 38
 log function, 45 39
 Logical key, 191 40
 logical key, 202 41
 logical operator, 31, 32 42
 lookup, 117 43
 loop, 58 44
 for, 68, 92 45
 infinite, 58 46
 maximum, 62 47
 minimum, 62 48
 nested, 112, 117 49
 traversal, 68 50
 while, 57 51
 looping 52
 with dictionaries, 114 53
 with indices, 93 54
 with strings, 70 55
 looping and counting, 70 56
 low-level language, 16 57
 machine code, 16 58
 main memory, 16 59
 Many to many relationship, 196 60
 math function, 45 61
 membership 62
 dictionary, 110 63
 list, 92 64
 set, 110 65
 message, 179 66
 method, 72, 76, 179 67
 append, 94, 101 68
 close, 88 69
 count, 73 70
 extend, 94 71
 get, 112 72
 items, 123 73
 join, 98 74
 keys, 114 75
 pop, 95 76
 remove, 95 77
 sort, 94, 102, 120 78
 split, 97, 123 79
 string, 76 80
 values, 110 81
 void, 95 82
 method, list, 94 83
 mnemonic, 27, 29 84
 module, 45, 53 85
 random, 46 86
 sqlite3, 183 87
 module object, 45 88
 modulus operator, 24, 29 89
 mutability, 70, 92, 94, 100, 119, 127 90
 negative index, 68 91
 nested conditional, 35, 40 92
 nested list, 91, 93, 105 93
 nested loops, 112, 117 94
 newline, 25, 81, 88, 89 95
 non-greedy, 153 96
 None special value, 51, 62, 95 97
 normalization, 203 98

- not operator, 32 1
- number, random, 46 2
- OAuth, 165 3
- object, 70, 76, 99, 105, 173, 179 4
 - function, 48 5
 - inheritance, 177 6
- object lifecycle, 175 7
- object-oriented, 167 8
- open function, 79, 86 9
- OpenStreetMap, 205 10
- operand, 22, 29 11
- operator, 29 12
 - and, 32 13
 - boolean, 71 14
 - bracket, 67, 92, 120 15
 - comparison, 31 16
 - in, 71, 92, 110 17
 - is, 99 18
 - logical, 31, 32 19
 - modulus, 24, 29 20
 - not, 32 21
 - or, 32 22
 - slice, 69, 94, 101, 120 23
 - string, 24 24
- operator, arithmetic, 22 25
- optional argument, 73, 97 26
- or operator, 32 27
- order of operations, 23, 29 28
- parameter, 49, 53, 100 29
- parent class, 180 30
- parentheses 31
 - argument in, 43 32
 - empty, 47, 72 33
 - overriding precedence, 23 34
 - parameters in, 50 35
 - regular expression, 137, 153 36
 - tuples in, 119 37
- parse, 16 38
- parsing 39
 - HTML, 154, 169 40
- parsing HTML, 152 41
- pass statement, 33 42
- pattern 43
 - decorate-sort-undecorate, 121 44
 - DSU, 121 45
 - filter, 83 46
 - guardian, 38, 40, 76 47
 - search, 76 48
 - swap, 122 49
- PEMDAS, 23 50
- persistence, 79 51
- pi, 45 52
- pop method, 95 53
- port, 157 54
- portability, 16 55
- precedence, 29 56
- Primary key, 190, 195 57
- primary key, 188, 203 58
- Primary key retrieval, 191 59
- print function, 16 60
- problem solving, 4, 16 61
- program, 11, 16 62
- programming language, 5 63
- prompt, 16, 25 64
- pseudorandom, 46, 53 65
- Python 2, 23, 25 66
- Python 3, 23 67
- Pythonic, 87, 89 68
- QA, 86, 89 69
- Quality Assurance, 86, 89 70
- quotation mark, 19, 69 71
- radian, 45 72
- randint function, 46 73
- random function, 46 74
- random module, 46 75
- random number, 46 76
- random walk programming, 14 77
- re module, 131 78
- reference, 100, 105 79
 - aliasing, 100 80
- regex, 131 81
 - character sets(brackets), 135 82
 - findall, 134 83
 - parentheses, 137, 153 84
 - search, 131 85
 - wild card, 132 86
- regular expressions, 131 87
- relation, 203 88
- Relational Model, 187 89
- remove method, 95 90
- repetition 91
 - list, 93 92
- repr function, 88 93
- reserved words, 5 94
- return value, 43, 53 95
- reversed function, 127 96
- Romeo and Juliet, 105, 113, 115, 121, 97 125

rules of precedence, 23, 29 1
 runtime error, 28 2
 sanity check, 117 3
 scaffolding, 117 4
 scatter, 128 5
 script, 10 6
 script mode, 21, 51 7
 search pattern, 76 8
 secondary memory, 16, 79 9
 semantic error, 16, 20, 29 10
 semantics, 16 11
 sequence, 67, 76, 91, 97, 119, 126 12
 Service Oriented Architecture, 166 13
 set membership, 110 14
 shape, 128 15
 shape error, 128 16
 short circuit, 38, 40 17
 sine function, 45 18
 singleton, 119, 128 19
 slice, 76 20
 copy, 69, 94 21
 list, 94 22
 string, 69 23
 tuple, 120 24
 update, 94 25
 slice operator, 69, 94, 101, 120 26
 SOA, 166 27
 socket, 158 28
 sort method, 94, 102, 120 29
 sorted function, 127 30
 source code, 16 31
 special value 32
 False, 31 33
 None, 51, 62, 95 34
 True, 31 35
 spider, 158 36
 split method, 97, 123 37
 SQL 38
 Constraint, 198 39
 CREATE, 187 40
 CRUD, 185, 187 41
 IGNORE, 192 42
 INDEX, 191 43
 JOIN, 189, 195, 199 44
 ON, 189, 195 45
 UNIQUE, 190, 192 46
 sqlite3 module, 183 47
 sqrt function, 45 48
 squiggly bracket, 109 49
 statement, 21, 29 50

assignment, 20 51
 break, 58 52
 compound, 32 53
 conditional, 32, 40 54
 continue, 59 55
 del, 95 56
 for, 60, 68, 92 57
 if, 32 58
 import, 53 59
 pass, 33 60
 try, 86 61
 while, 57 62
 str function, 44 63
 string, 19, 29, 97, 127 64
 comparison, 71 65
 empty, 98 66
 find, 132 67
 immutable, 70 68
 method, 72 69
 operation, 24 70
 slice, 69 71
 split, 137 72
 startswith, 132 73
 string method, 76 74
 string representation, 88 75
 string type, 19 76
 super class, 178 77
 swap pattern, 122 78
 syntax error, 28 79
 temperature conversion, 37 80
 text file, 89 81
 Through table, 197 82
 time, 149 83
 time.sleep, 149 84
 traceback, 37, 39, 40 85
 traversing, 68, 76, 111, 114, 121 86
 list, 92 87
 traverse 88
 dictionary, 124 89
 trigonometric function, 45 90
 True special value, 31 91
 try statement, 86 92
 tuple, 119, 127, 128, 203 93
 as key in dictionary, 126 94
 assignment, 122 95
 comparison, 120 96
 in brackets, 126 97
 singleton, 119 98
 slice, 120 99
 tuple assignment, 128 100

tuple function, 119 1
type, 19, 29, 174 2
 bool, 31 3
 dict, 109 4
 file, 79 5
 list, 91 6
 tuple, 119 7
type conversion, 44 8
TypeError, 67, 70, 120 9
typographical error, 14 10

underscore character, 21 11
Unicode, 185 12
update, 57 13
 item, 93 14
 slice, 94 15
urllib 16
 image, 148 17
use before def, 28, 49 18

value, 19, 29, 99, 117 19
ValueError, 26, 123 20
values method, 110 21
variable, 20, 29 22
 updating, 57 23
Visualization 24
 map, 205 25
 networks, 207 26
 page rank, 207 27
void function, 51, 53 28
void method, 95 29

web 30
 scraping, 152 31
wget, 157 32
while loop, 57 33
whitespace, 39, 52, 88 34
wild card, 132, 143 35

XML, 166 36

zero, index starting at, 67, 92 37